

Automate Infrastructure



Logistics



- Class Hours:
- Instructor will provide class start and end times.
- Breaks throughout class

- Lunch:
- 1 hour 15 minutes



- Telecommunication:
- Turn off or set electronic devices to vibrate
- Reading or attending to devices can be distracting to other students

- Miscellaneous
 - Courseware
 - Bathroom

Course Objectives



- Explain the benefits of DevSecOps
- Use Git for version control
- Describe what Configuration Drift is and how to avoid it using Infrastructure as Code
- Use Terraform to automate infrastructure management.
- Build Ansible playbooks to install and configure applications.
- Build pipelines using GitHub Actions to automate infrastructure and application deployment.

Meet your Instructor: Antoine Victor



- Technical Trainer | DevOps Consultant | Content Creator
- Instructor with 20+ years of Technical Training experience
- Passionate about Teaching by Doing
- Favorite Tools: VS Code, Terraform, Ansible, Azure, AWS, Jenkins, Jira, and more
- Known for hands-on labs, real-world demos, and interactive workshops
- MCSD, MCSE, MCSA CSA, MCT
- Certified Scrum Developer (CSD)
- Certified Scrum Master / Product Owner

- Based in Seattle, WA | Loves Snowboarding, Evs and Renewable Energy
- Drives a Tesla Model S | Enjoys making travel content with his dog Winston

Expertise

- Agile
- DevOps
- Cloud
- Automation
- CICD
- Full Stack Development



antoine@innovationinsoftware.com



<https://github.com/ProDataMan>



<https://www.linkedin.com/in/antoinievictor/>

Introductions

Hello!

- Name
- Job Role
- Your experience with (scale 1 – 5):
 - Ansible
 - Terraform
 - Git
 - GitHub Actions
- Expectations for course (please be specific)

<https://jruels.github.io/gh-auto-infra>

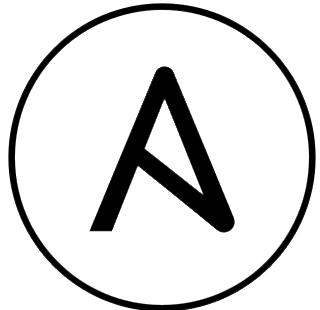


Automation:

Repeatability by automating activities.



Automation tools



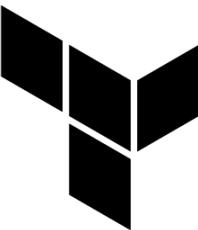
ANSIBLE



CHEF



puppet



HashiCorp
Terraform

Infrastructure as Code



What is IaC?

Infrastructure as code (IaC) is the process of managing and provisioning computer data centers through machine-readable definition files. Used with bare-metal as well as virtual machines and many other resources. Normally a declarative approach

Infrastructure as Code



- Programmatically provision and configure components
- Treat like any other code base
 - Version control
 - Automated testing
 - data backup

Infrastructure as Code



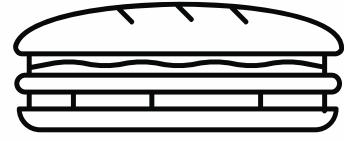
Provisioning infrastructure through software to achieve consistent and predictable environments.

Core Concepts



- Defined in code
- Stored in source control
- Declarative or imperative

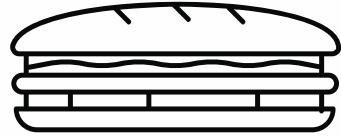
Imperative



```
# Make a sandwich  
get bread  
get mayo  
get turkey  
get lettuce
```

```
spread mayo on bread  
put lettuce in between bread  
put turkey in between bread on  
top of turkey
```

Declarative

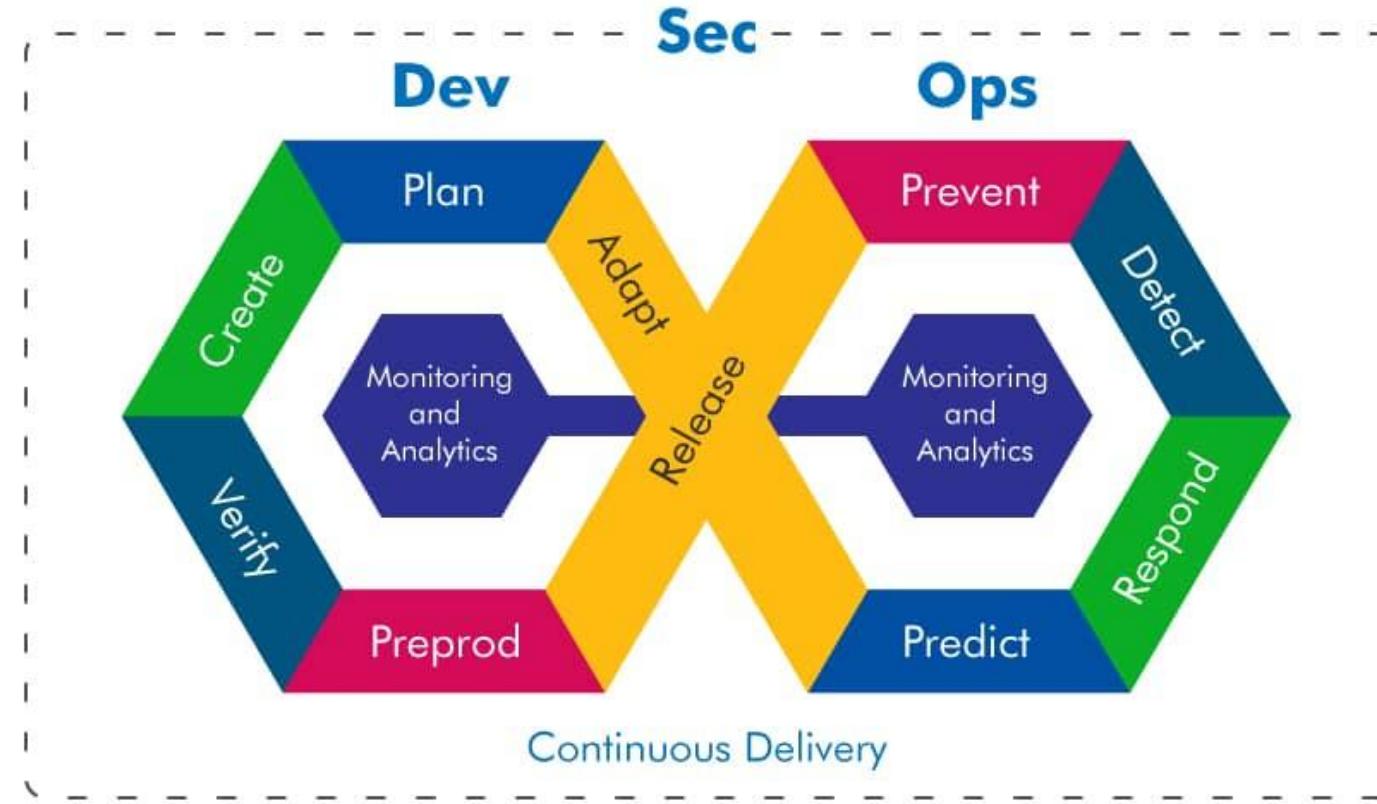


```
# Make a sandwich
food sandwich "turkey" {
    ingredients = [
        "bread", "turkey", "mayo", "lettuce"
    ]
}
```

DevSecOps and Policy-as-Code



Introduction to DevSecOps



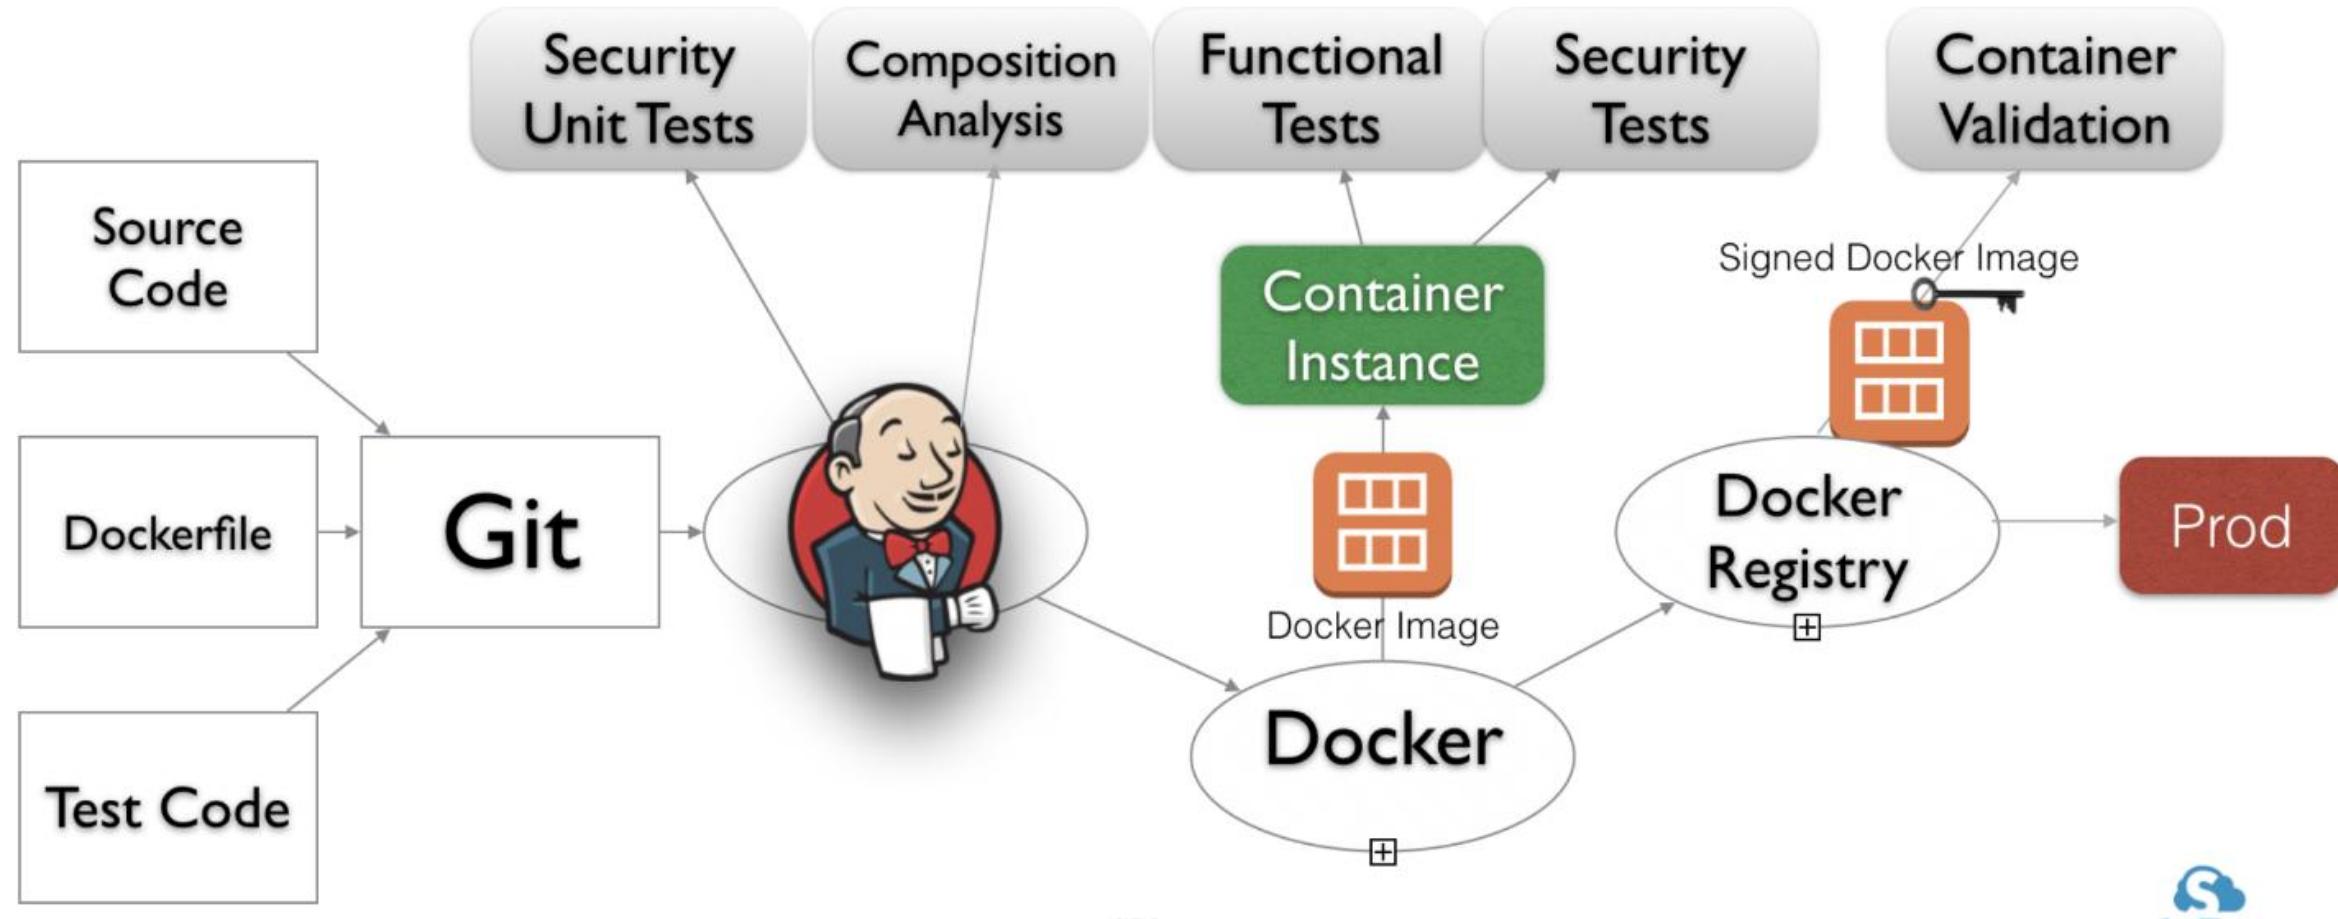
Continuously assessing security measures within the development cycle to gauge their effectiveness and foster ongoing improvement.

What is DevSecOps?



- **DevSecOps** integrates security practices into DevOps, making security a shared responsibility throughout the software development lifecycle.
- **Security** is embedded at every stage, from code development to deployment.
- **Goal:** Catch vulnerabilities early and ensure continuous protection.

Example DevSecOps Architecture



Key principles of DevSecOps



- **Continuous security integration**, automating security checks for faster, safer releases.
- **Close collaboration** among development, security, and operations teams for proactive threat mitigation.
- **Real-time monitoring and feedback loops**, allowing teams to assess and improve security measures based on actual performance.

Objective: To enable fast, secure software delivery by making security an integral part of the development pipeline.

Why DevSecOps Matters



- **Increasing Cyber Threats:** As cyber threats grow more sophisticated, integrating security early is essential to protect data and systems.
- **Faster Development Cycles:** Traditional security practices can slow down the pace; DevSecOps enables rapid, secure releases.
- **Cost Efficiency:** Catching vulnerabilities early reduces the cost of fixing issues post-deployment.
- **Compliance & Trust:** Automated checks ensure that applications meet regulatory standards, increasing reliability and trust.

Benefits of DevSecOps



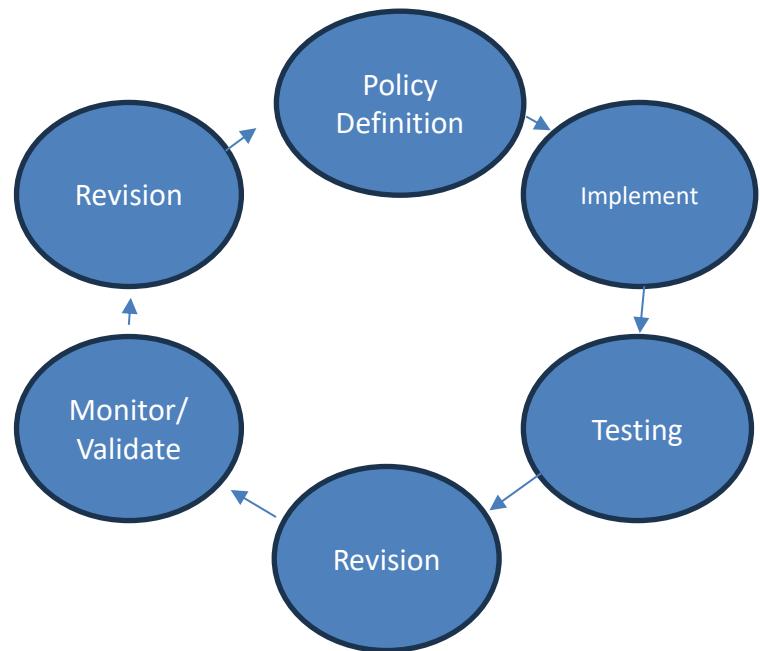
- **Improved Security Posture:** Continuous integration of security practices reduces risks and improves resilience.
- **Higher Quality Software:** Early detection of issues leads to more stable, reliable software.
- **Faster Time-to-Market:** Automating security allows for rapid deployment without compromising safety.
- **Enhanced Team Efficiency:** Cross-functional collaboration minimizes bottlenecks and empowers teams to work together on security goals.

Introduction to Policy-As-Code



Implementing and continuously evaluating security policies as code within the development process to ensure they remain effective, adaptable, and drive consistent compliance.

What is Policy as Code?



Policy as Code is the practice of defining and managing security and compliance policies as version-controlled code.

By **codifying policies**, organizations ensure that security and compliance requirements are automatically enforced within the development pipeline.

This approach enables **consistency**, **traceability**, and **automation**, making it easier to detect and correct policy violations early

Key principles of Policy as Code



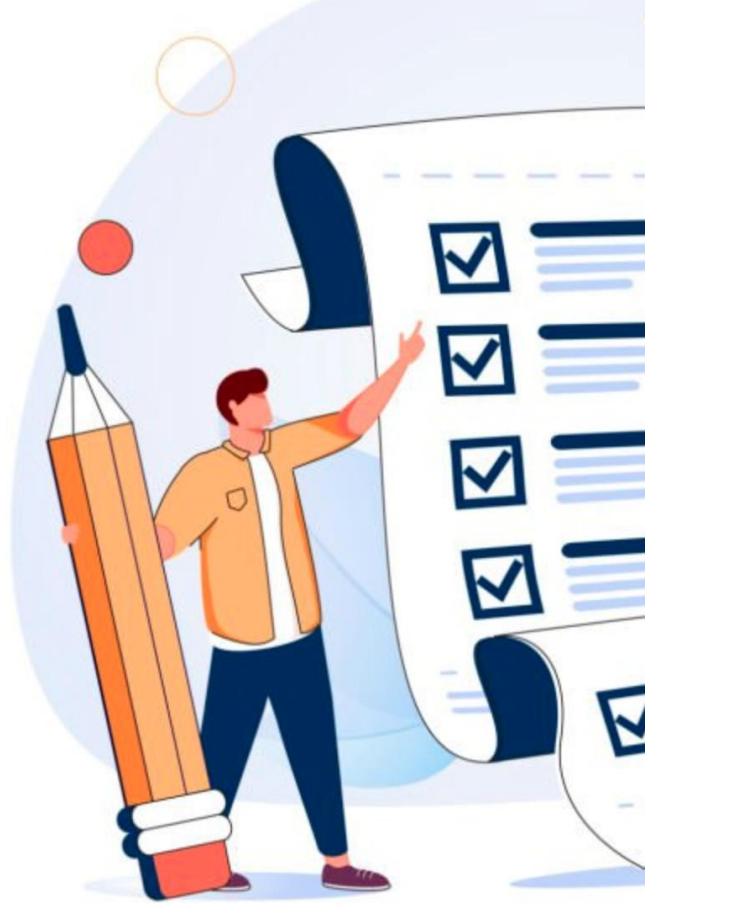
Declarative Policy Definitions: Policies are written in a clear, human-readable format, defining the “what” rather than the “how.”

Embedded in CI/CD: Policies are integrated within the CI/CD pipeline to ensure they are validated at every stage of development.

Automation and Enforcement: Policy violations trigger automated alerts or actions, reducing reliance on manual intervention.

Version Control and Collaboration: Policies are treated as code, allowing teams to review, update, and manage them collaboratively.

Why Policy as Code Matters



Automation of Compliance: Ensures that compliance checks happen automatically, reducing manual work and human error.

Consistency Across Environments: Policy as Code guarantees that security standards are enforced consistently across all stages and environments.

Adaptability and Agility: Codified policies can be quickly updated to respond to new threats or regulatory changes.

Auditability and Traceability: Version-controlled policies provide an auditable history, helping teams track policy changes over time.

Benefits of Policy as Code



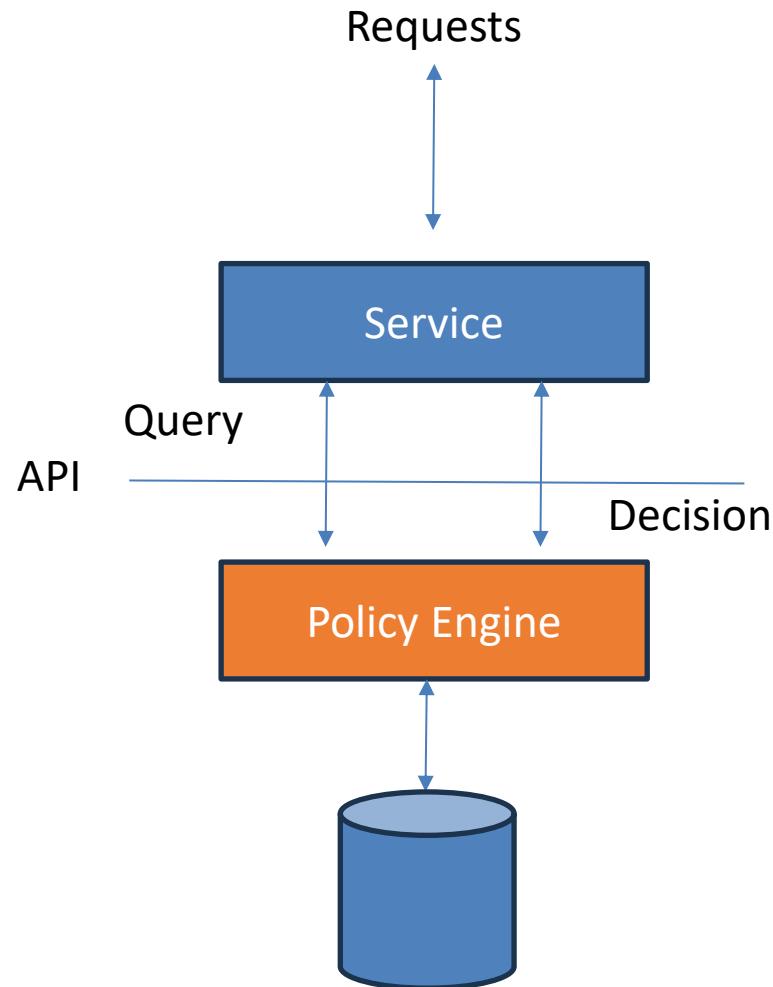
Enhanced Security Posture: By enforcing policies automatically, teams can detect and resolve non-compliance issues sooner.

Reduced Operational Overhead: Automation reduces the need for repetitive manual compliance checks.

Greater Transparency and Accountability: Codified policies provide a clear record of compliance standards and enforcement.

Improved Adaptability: Policies can be updated as code, allowing organizations to quickly respond to new security requirements.

Policy as Code Components



Policy Engine: Interprets policies and returns decisions.

Service: Requests policy checks for specific actions or access.

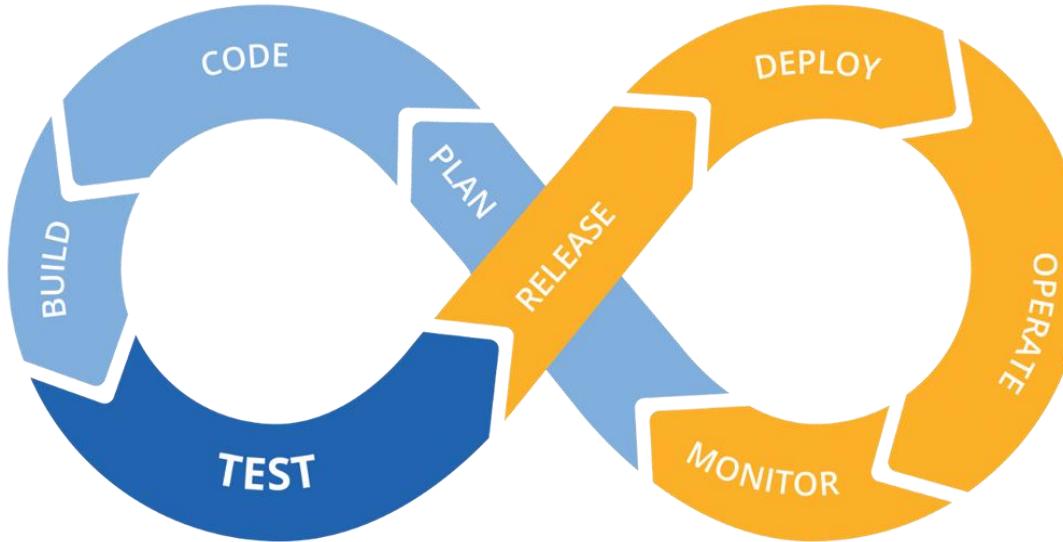
Requests: Data sent to the policy engine for evaluation.

API: Connects services to the policy engine for seamless policy checks.

Query: The policy engine evaluates input data against policies.

Decision: The result returned to the service, enabling rule enforcement.

Security Integration in the CI/CD Pipeline



Continuously embedding and assessing security checks throughout the CI/CD pipeline to uphold robust, adaptable safeguards and maintain consistent compliance from development to deployment.

Introduction to Security in CI/CD



Integrating Security Throughout: Embedding security at every stage of the CI/CD pipeline enables teams to catch and address vulnerabilities as they arise, preventing issues from reaching production.

Ensuring Compliance and Readiness: By weaving security checks into the pipeline, code remains secure, compliant, and ready for production with each deployment.

Why Integrate Security in CI/CD?



Proactive Threat Mitigation: Identifies vulnerabilities as soon as code is committed.

Cost-Effective: Fixing issues early reduces the cost and complexity of remediation.

Continuous Protection: Ensures secure code delivery by integrating automated security checks.

Key Security Stages in the CI/CD Pipeline



1. **Code Scanning:** Analyzes source code for vulnerabilities.
2. **Dependency Checks:** Validates that third-party libraries are free from known vulnerabilities.
3. **Container Security:** Scans container images for security risks.
4. **Infrastructure as Code (IaC) Scanning:** Checks for misconfigurations in IaC templates.

Benefits of Security Integration in CI/CD



Early Detection: Vulnerabilities are identified before reaching production.

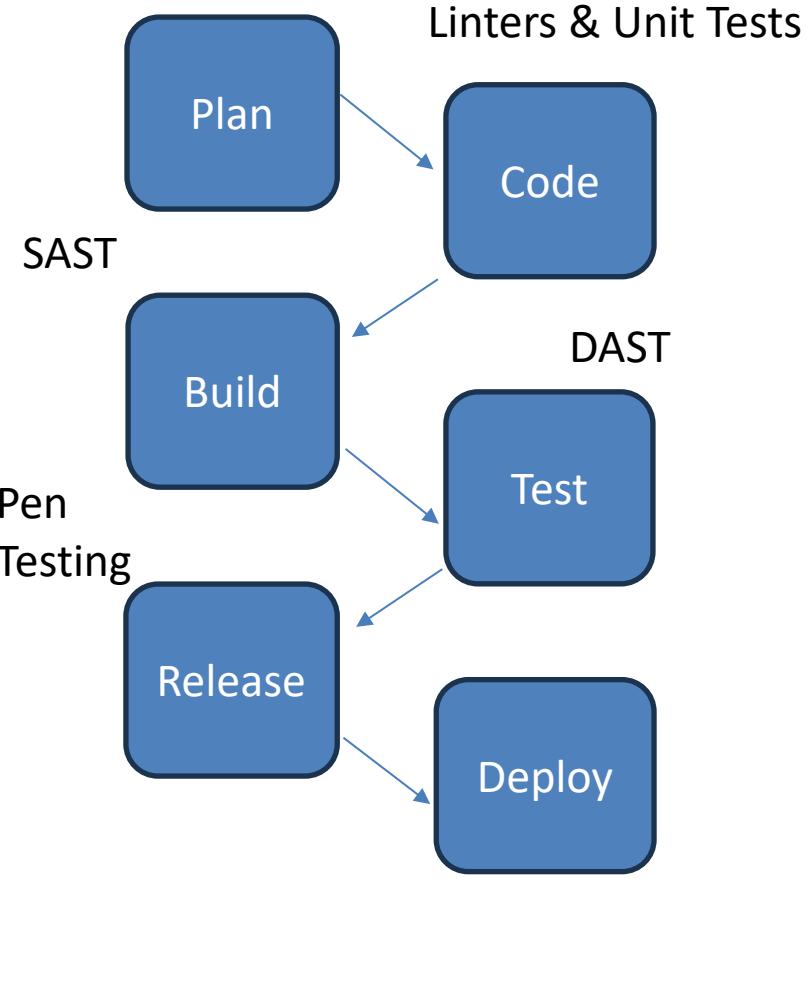
Automated Compliance: Security policies and compliance are enforced automatically.

Enhanced Team Efficiency: Developers receive instant feedback, allowing faster and safer code iterations.

Reduced Risk: Minimizes exposure to potential threats by ensuring secure code deployment.

Overview of CI/CD Security Integrations

Security Analysis



Code Scanning: Identifies vulnerabilities directly in the source code before merging.

Dependency Scanning: Checks third-party libraries for known vulnerabilities to ensure secure dependencies.

Container Security: Scans container images for configuration issues and vulnerabilities.

Infrastructure as Code (IaC) Scanning: Validates cloud and infrastructure templates to catch misconfigurations early.

Dynamic Application Security Testing (DAST): Tests the running application for security weaknesses from an attacker's perspective.

GitHub Actions for Code Scanning

```
name: Code Scanning

on:
  pull_request:
    branches: [main]

jobs:
  code-scan:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout Code
        uses: actions/checkout@v2

      - name: Run CodeQL Analysis
        uses: github/codeql-action/analyze@v1
        with:
          category: "security"
```

What It Does: Automatically scans the code with CodeQL whenever a pull request is opened to catch vulnerabilities early.

Why It's Useful: Gives developers immediate feedback on security issues, helping to fix them before they reach the main branch.

Impact on Security: Maintains a secure codebase by proactively identifying and resolving security flaws in the source code.

Dependency Scanning with GitHub Actions

```
name: Dependency Scan

on:
  schedule:
    - cron: '0 2 * * 1'

jobs:
  dependabot:
    runs-on: ubuntu-latest
    steps:
      - name: Dependency Scan
        uses: dependabot/fetch-metadata@v1
```

What It Does: Regularly scans all third-party libraries and dependencies to spot vulnerabilities.

Why It's Useful: Ensures that all dependencies remain secure by flagging outdated or insecure libraries.

Reduces risks associated with third-party code, keeping the project safer with up-to-date dependencies.

GitHub Actions for Container Security

```
name: Container Security

on:
  push:
    branches: [main]

jobs:
  container-scan:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout Code
        uses: actions/checkout@v2
      - name: Scan Docker Image
        uses: anchore/scan-action@v2
        with:
          image: 'my-image:latest'
```

What It Does: Scans Docker images to identify vulnerabilities in container layers, configurations, and dependencies.

Why It's Useful: Detects security risks within containers before they reach production, ensuring containers are safe to deploy.

Impact on Security: Prevents compromised or misconfigured containers from entering the deployment pipeline, securing the application environment.

IaC Security Checks with GitHub Actions

```
name: IaC Security Scan

on:
  pull_request:
    branches: [main]

jobs:
  iac-scan:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout Code
        uses: actions/checkout@v2

      - name: Run tfsec
        uses: aquasecurity/tfsec-action@v1
```

What It Does: Analyzes Infrastructure as Code (IaC) templates to check for misconfigurations or potential security issues.

Why It's Useful: Identifies insecure configurations before deploying infrastructure, helping to secure cloud and on-premises resources.

- Ensures that infrastructure is set up securely from the start, reducing the risk of exposure from insecure configurations.

Automated Security Testing Workflow Example

```
name: Security Workflow
on: [push, pull_request]

jobs:
  security-checks:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout Code
        uses: actions/checkout@v2

      - name: CodeQL Analysis
        uses: github/codeql-action/analyze@v1

      - name: Dependency Scan
        uses: dependabot/fetch-metadata@v1

      - name: IaC Scan
        uses: aquasecurity/tfsec-action@v1
```

What It Does: Runs code scanning, dependency checks, and IaC scans in a single, streamlined workflow.

Why It's Useful: Simplifies security by consolidating multiple checks, making it easier to manage and monitor security in one place.

Provides comprehensive security coverage throughout the pipeline, reducing the risk of missing critical vulnerabilities.

Dynamic Application Security Testing (DAST)

```
name: DAST Security Testing
on:
  push:
    branches: [main]
jobs:
  dast-scan:
    runs-on: ubuntu-latest
    steps:
      - name: Start Application
        run: |
          docker-compose up -d
      - name: OWASP ZAP Scan
        uses: zaproxy/action-full-scan@v0.4.0
        with:
          target: 'http://localhost:8080'
          cmdOptions: '-config api.disablekey=true'
      - name: Stop Application
        run: |
          docker-compose down
```

What It Does: Runs a dynamic security scan against a live instance of the application, identifying vulnerabilities in how it behaves under simulated attack conditions.

Why It's Useful: Tests the application as an attacker would, helping identify security issues that may not be evident in the code alone.

Impact on Security: Provides a final layer of defense by catching vulnerabilities from an external perspective, enhancing overall application resilience.

Introduction to Open Policy Agent (OPA)



Open Policy Agent

Assessing and enforcing policies across systems to maintain effective, adaptable security and compliance throughout the development lifecycle.

What is Open Policy Agent (OPA)



Open Policy Agent

OPA is an open-source, general-purpose policy engine that unifies policy enforcement across different systems.

- It enables organizations to define and enforce policies consistently across applications, microservices, Kubernetes, CI/CD pipelines, and more.
- By centralizing policy management, OPA makes it easier to update, audit, and apply security and compliance policies.

Key Benefits of OPA



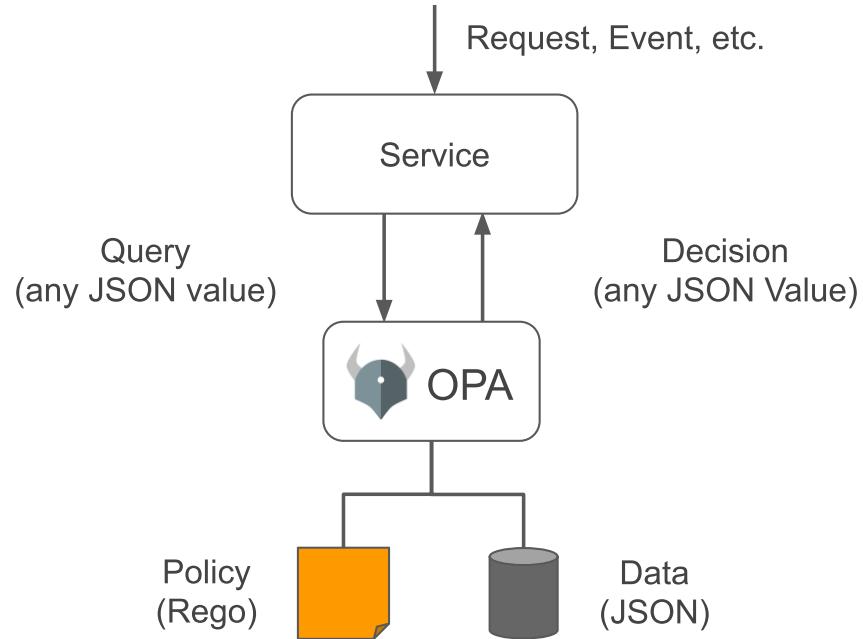
Open Policy Agent

Consistency: Standardizes policy enforcement across various environments and applications.

Flexibility: Allows policies to be written in a declarative language, Rego, adaptable to different use cases.

Auditability: Policies are version-controlled and auditable, supporting compliance requirements.

OPA Architecture Overview



Policy Engine: The OPA engine interprets policies written in Rego and makes real-time policy decisions.

Data Input: OPA evaluates policies against structured data (JSON) provided by applications or services.

Policy Decision API: Applications call OPA's API to receive authorization or policy decisions.

Writing a Basic OPA Policy

```
rego  
  
package authz  
  
allow {  
    input.user == "admin"  
}
```

This example shows a simple policy that checks if a user has access to a resource.

The policy allows access if the user in the input data is “admin.”

Applications can call OPA with JSON input like {"user": "admin"}, and OPA returns whether access is allowed.

Policy Enforcement in Microservices

```
rego

package authz

default allow = false

allow {
    input.user.role == "manager"
}
```

Enforce policies directly within microservices, ensuring consistent rules across the application.

- In this example, access is allowed only if the user's role is “manager.”
- This simplifies role-based access control in microservices by centralizing the logic in OPA.

Kubernetes Admission Control with OPA

```
rego

package kubernetes.admission

deny[msg] {
    input.request.kind.kind == "Pod"
    input.request.object.spec.containers[_]
        .securityContext.runAsUser == 0
    msg = "Running as root is not allowed"
}
```

OPA can be used to enforce security policies on Kubernetes clusters, such as blocking deployments with specific configurations.

- This example policy prevents deployment if a container runs as root, improving security.

Implementing CI/CD Pipeline Policies

```
rego  
  
package ci.policy  
  
deny[msg] {  
    input.vulnerabilities[_].severity == "critical"  
    msg = "Build blocked due to critical vulnerability"  
}
```

Enforce policies directly in CI/CD pipelines to ensure compliance before deployment.

- This example policy blocks builds if a critical vulnerability is detected, ensuring safer deployments.

Integrating OPA with REST APIs

```
rego  
  
package api.access  
  
allow {  
    input.user.type == "employee"  
    input.resource.accessible == true  
}
```

OPA policies can be tested with unit tests to ensure they work as expected.

- The example tests confirm that only “admin” users are allowed, ensuring policy accuracy.

Testing and Validating OPA Policies

```
rego

package authz

test_allow_admin {
    allow with input as {"user": "admin"}
}

test_deny_non_admin {
    not allow with input as {"user": "guest"}
}
```

OPA policies can be tested with unit tests to ensure they work as expected.

- The example tests confirm that only “admin” users are allowed, ensuring policy accuracy.

OPA Policy Evaluation in Action



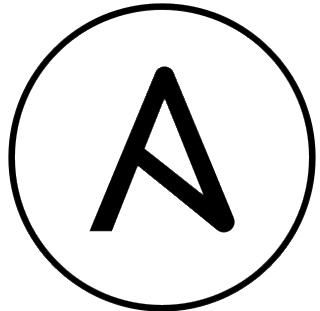
Open Policy Agent

Running OPA Locally: Use OPA run to load policies and data locally, allowing you to test policies before deploying.

Deploying in Production: OPA can be deployed as a sidecar or standalone service, integrating directly into application infrastructure.

OPA Decision Logging: Logs all policy decisions, providing visibility into enforcement for auditing and troubleshooting.

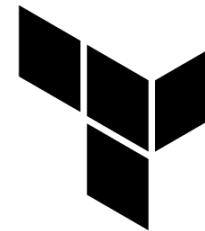
Git



A N S I B L E



git



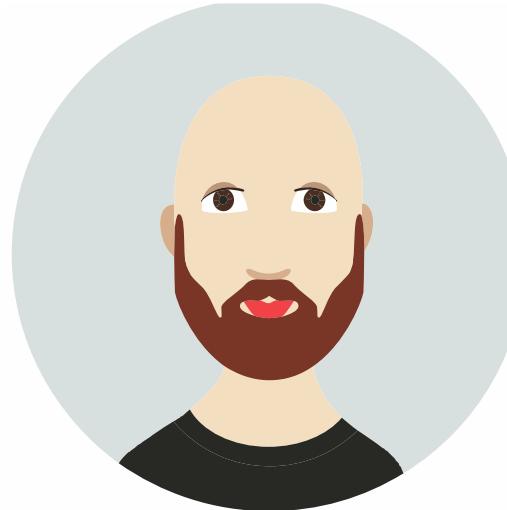
HashiCorp
Terraform



Git

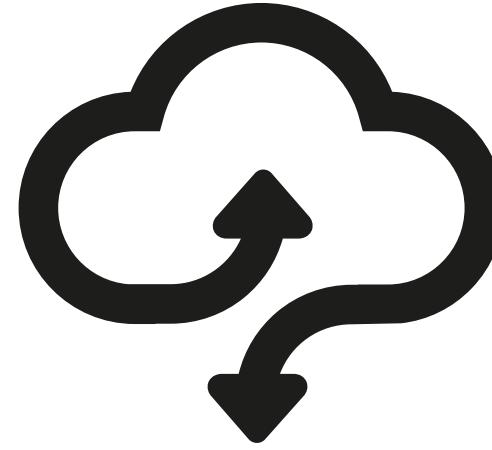


git



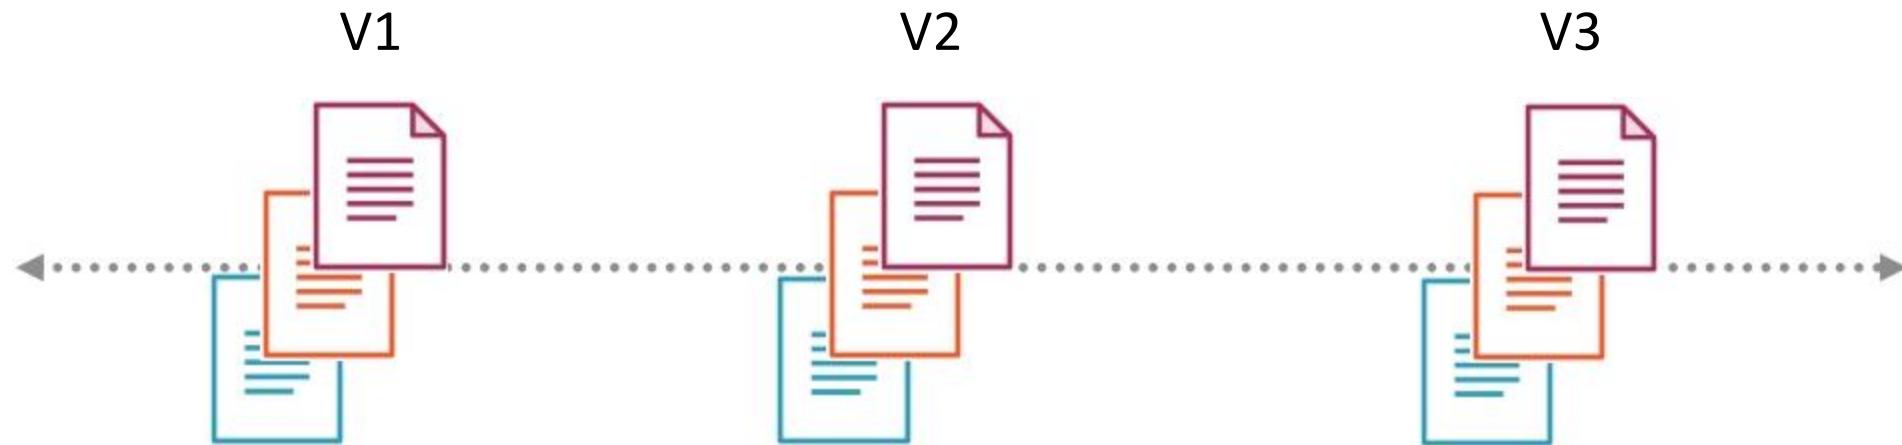
Git

git



What is Git?

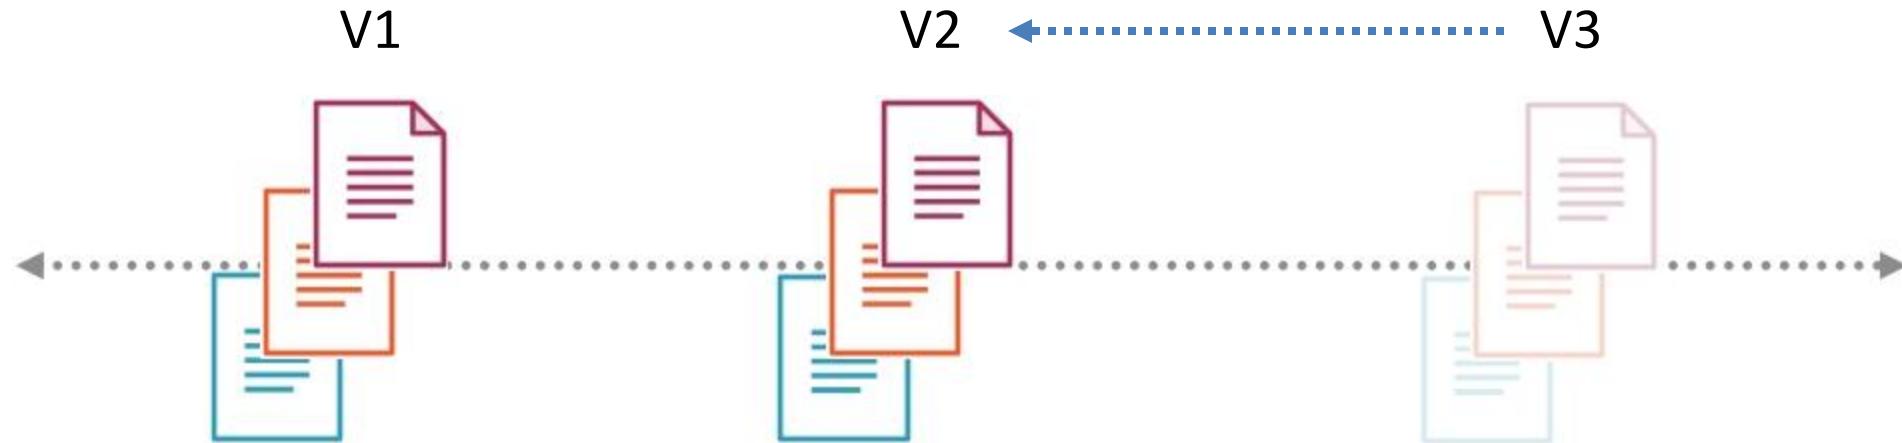
Version Control System



What is Git?

Version Control System

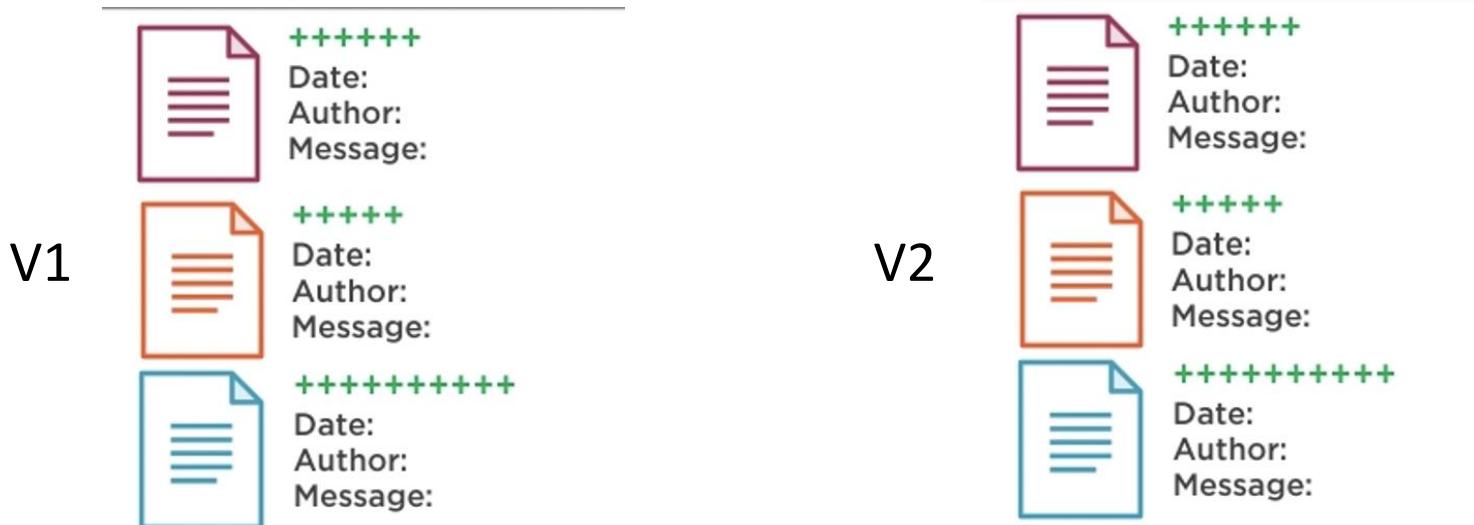
- Software designed to record changes to files made over time
- Ability to revert back to a previous file version or project version



What is Git?

Version Control System

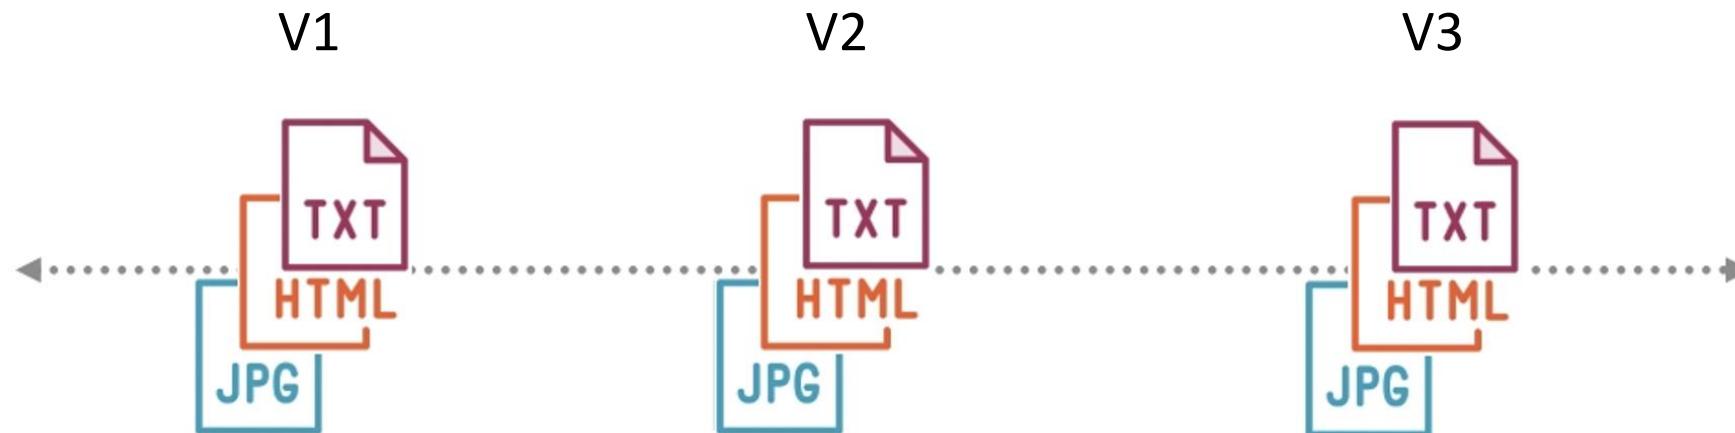
- Software designed to record changes to files made over time
- Ability to revert back to a previous file version or project version
- Compare changes made to files from one version to another



What is Git?

Version Control System

- Software designed to record changes to files made over time.
- Ability to revert back to a previous file version or project version.
- Compare changes made to files from one version to another.
- Ability to version control any plain text file, not just source code.



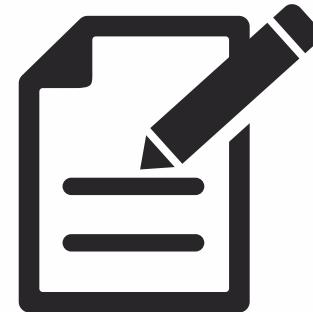
Git use-case



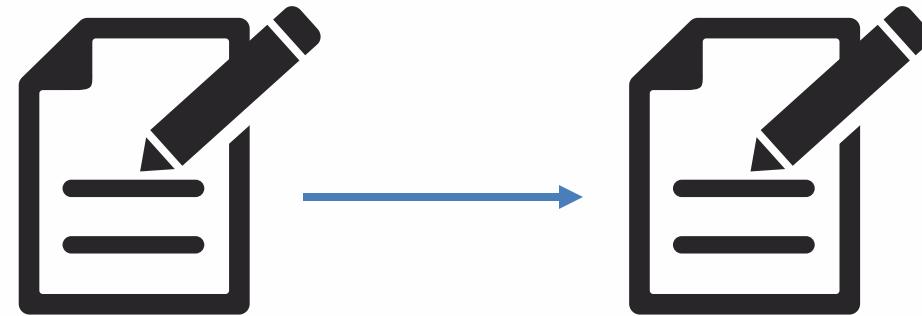
Git use-case



Latte



Git use-case



POP QUIZ: Dangers of copying files



What are some dangers of Steve's current process?

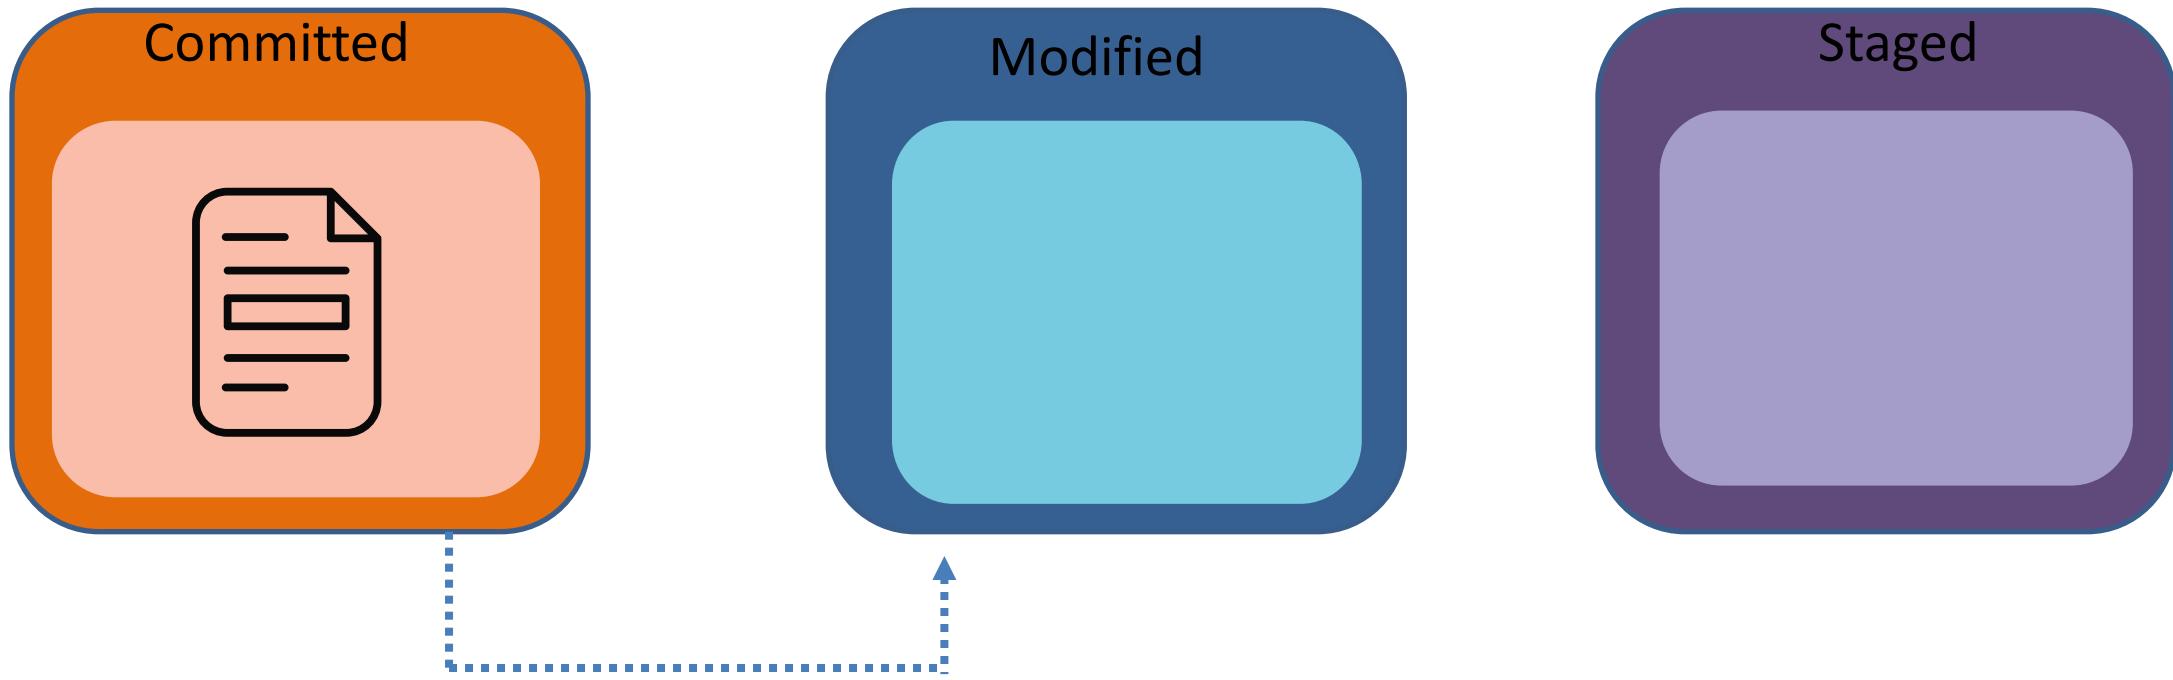
POP QUIZ: Dangers of copying files



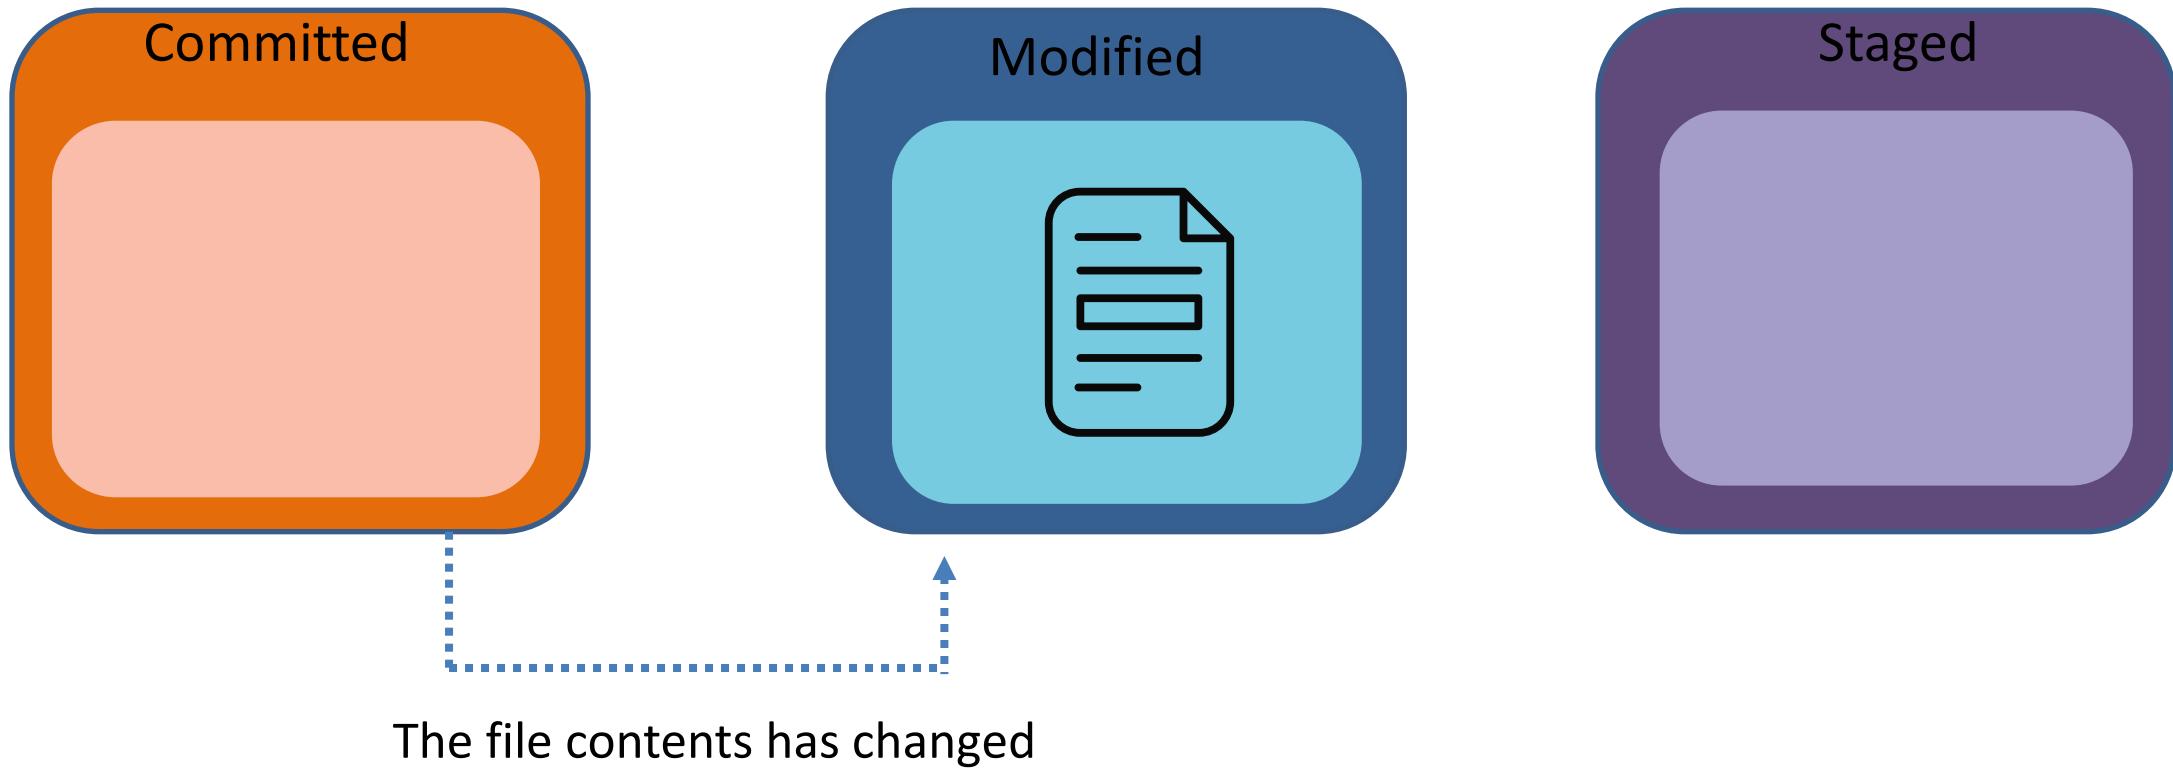
What are some dangers of Steve's current process?

- Files live in one location
- Accidentally overwrite existing recipes

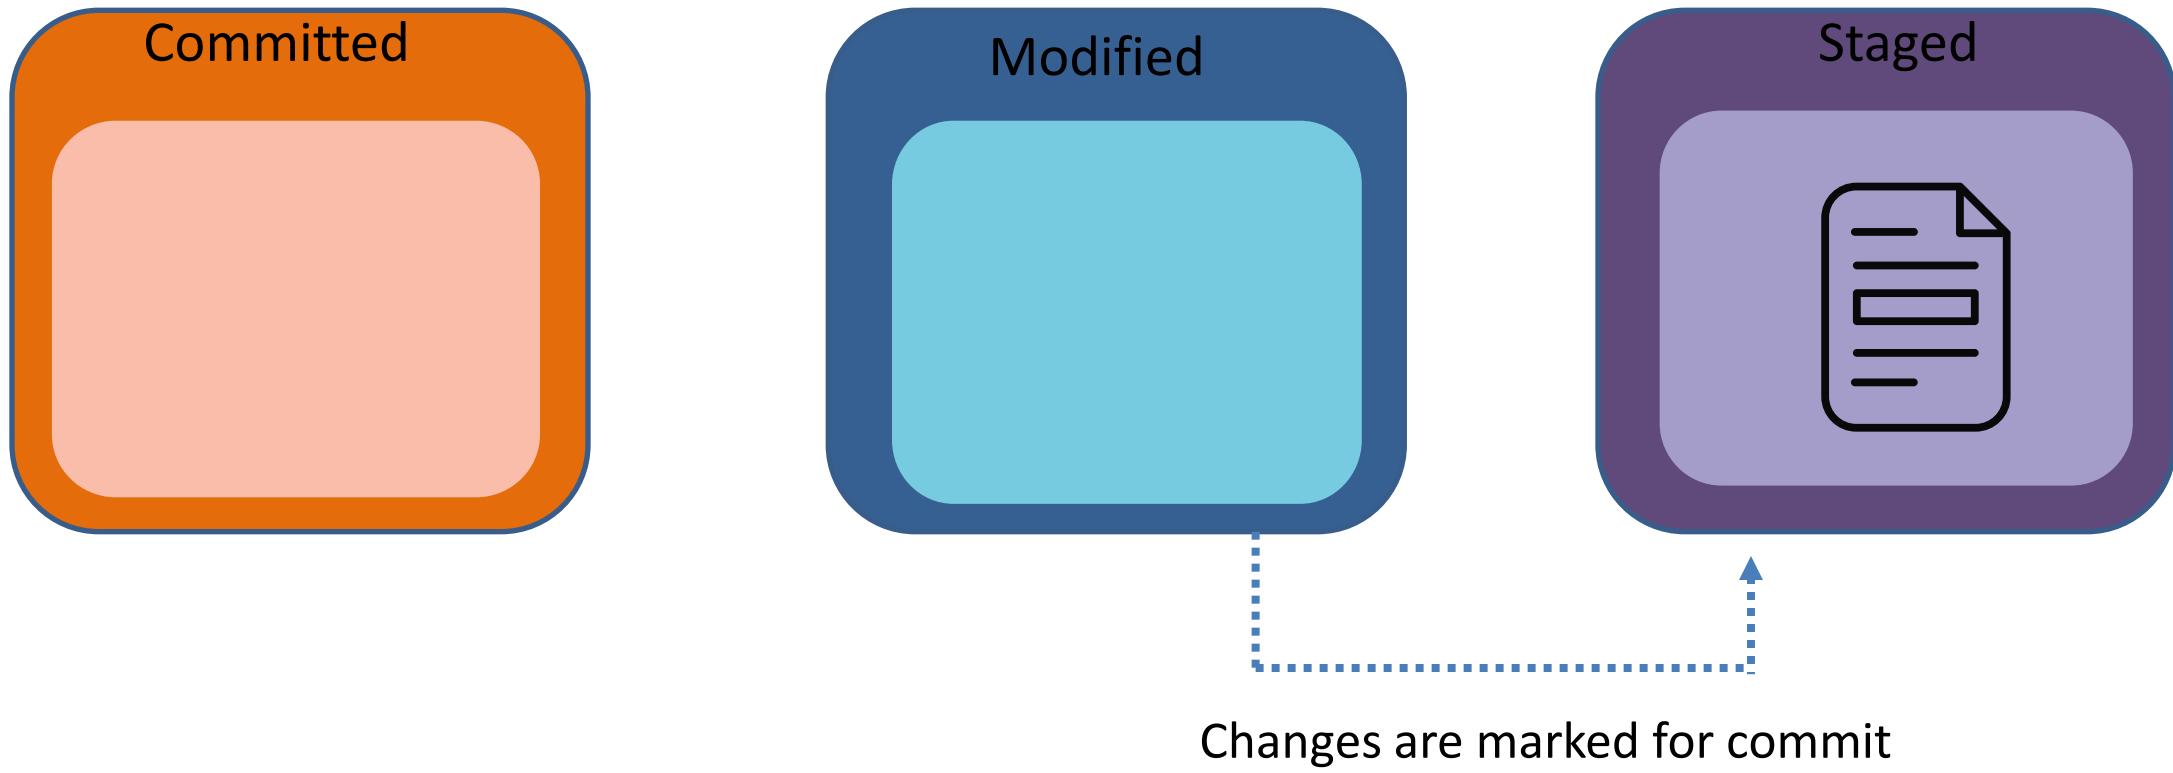
Three stages of a file



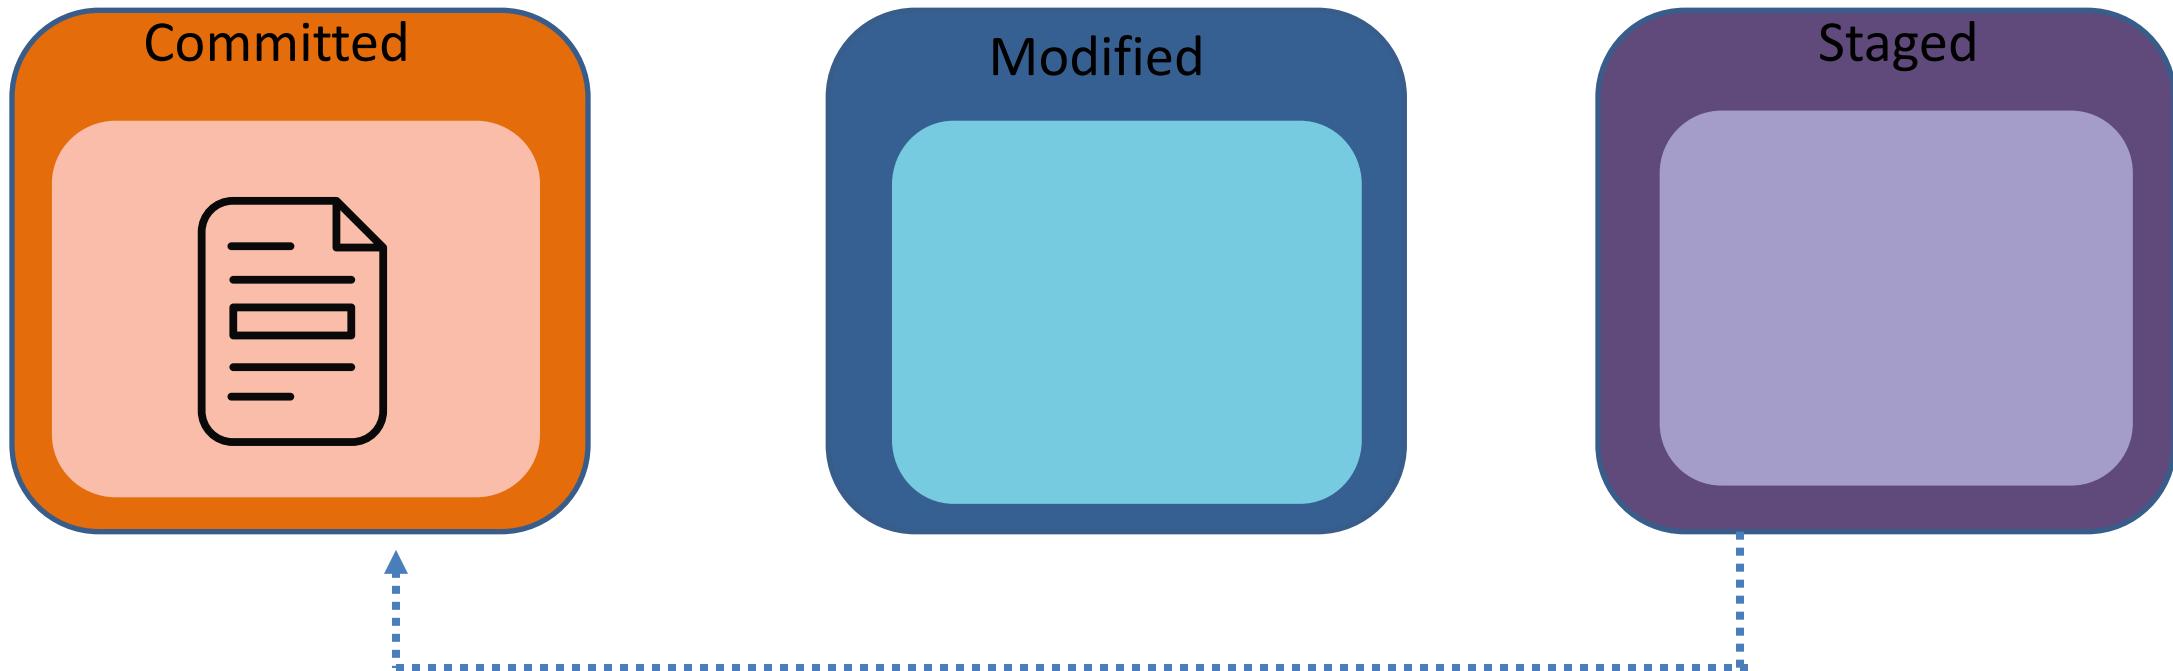
Three stages of a file



Three stages of a file



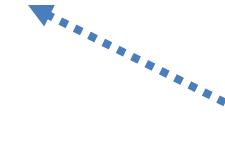
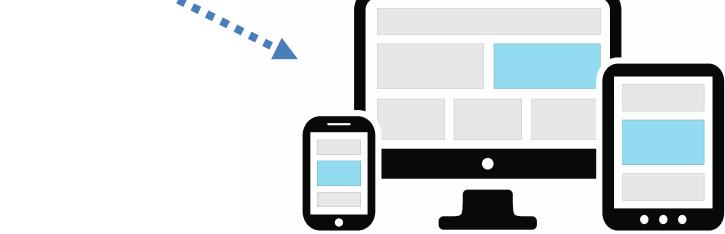
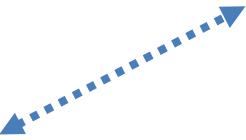
Three stages of a file



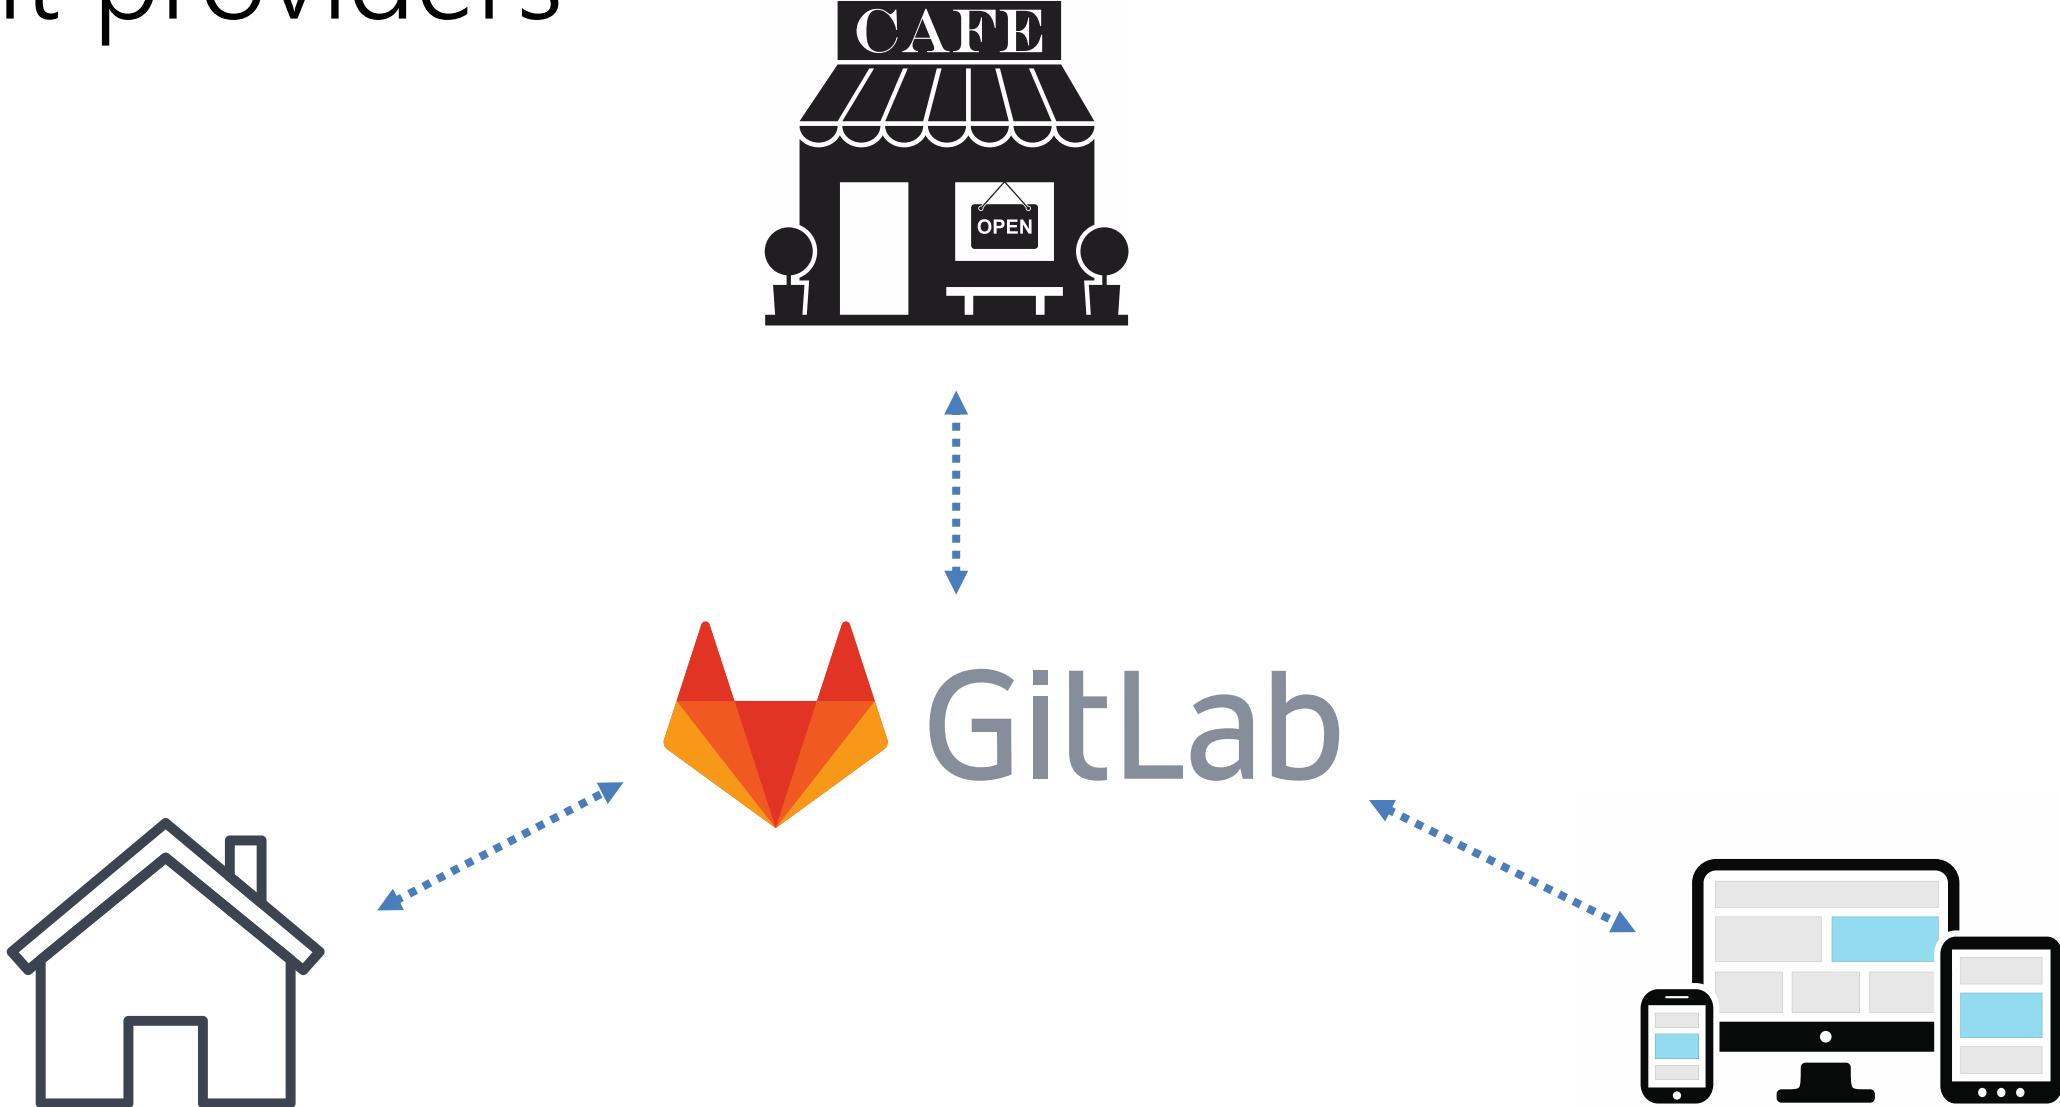
Git providers



GitHub



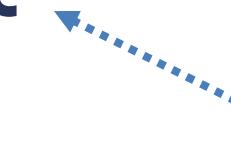
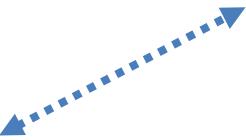
Git providers



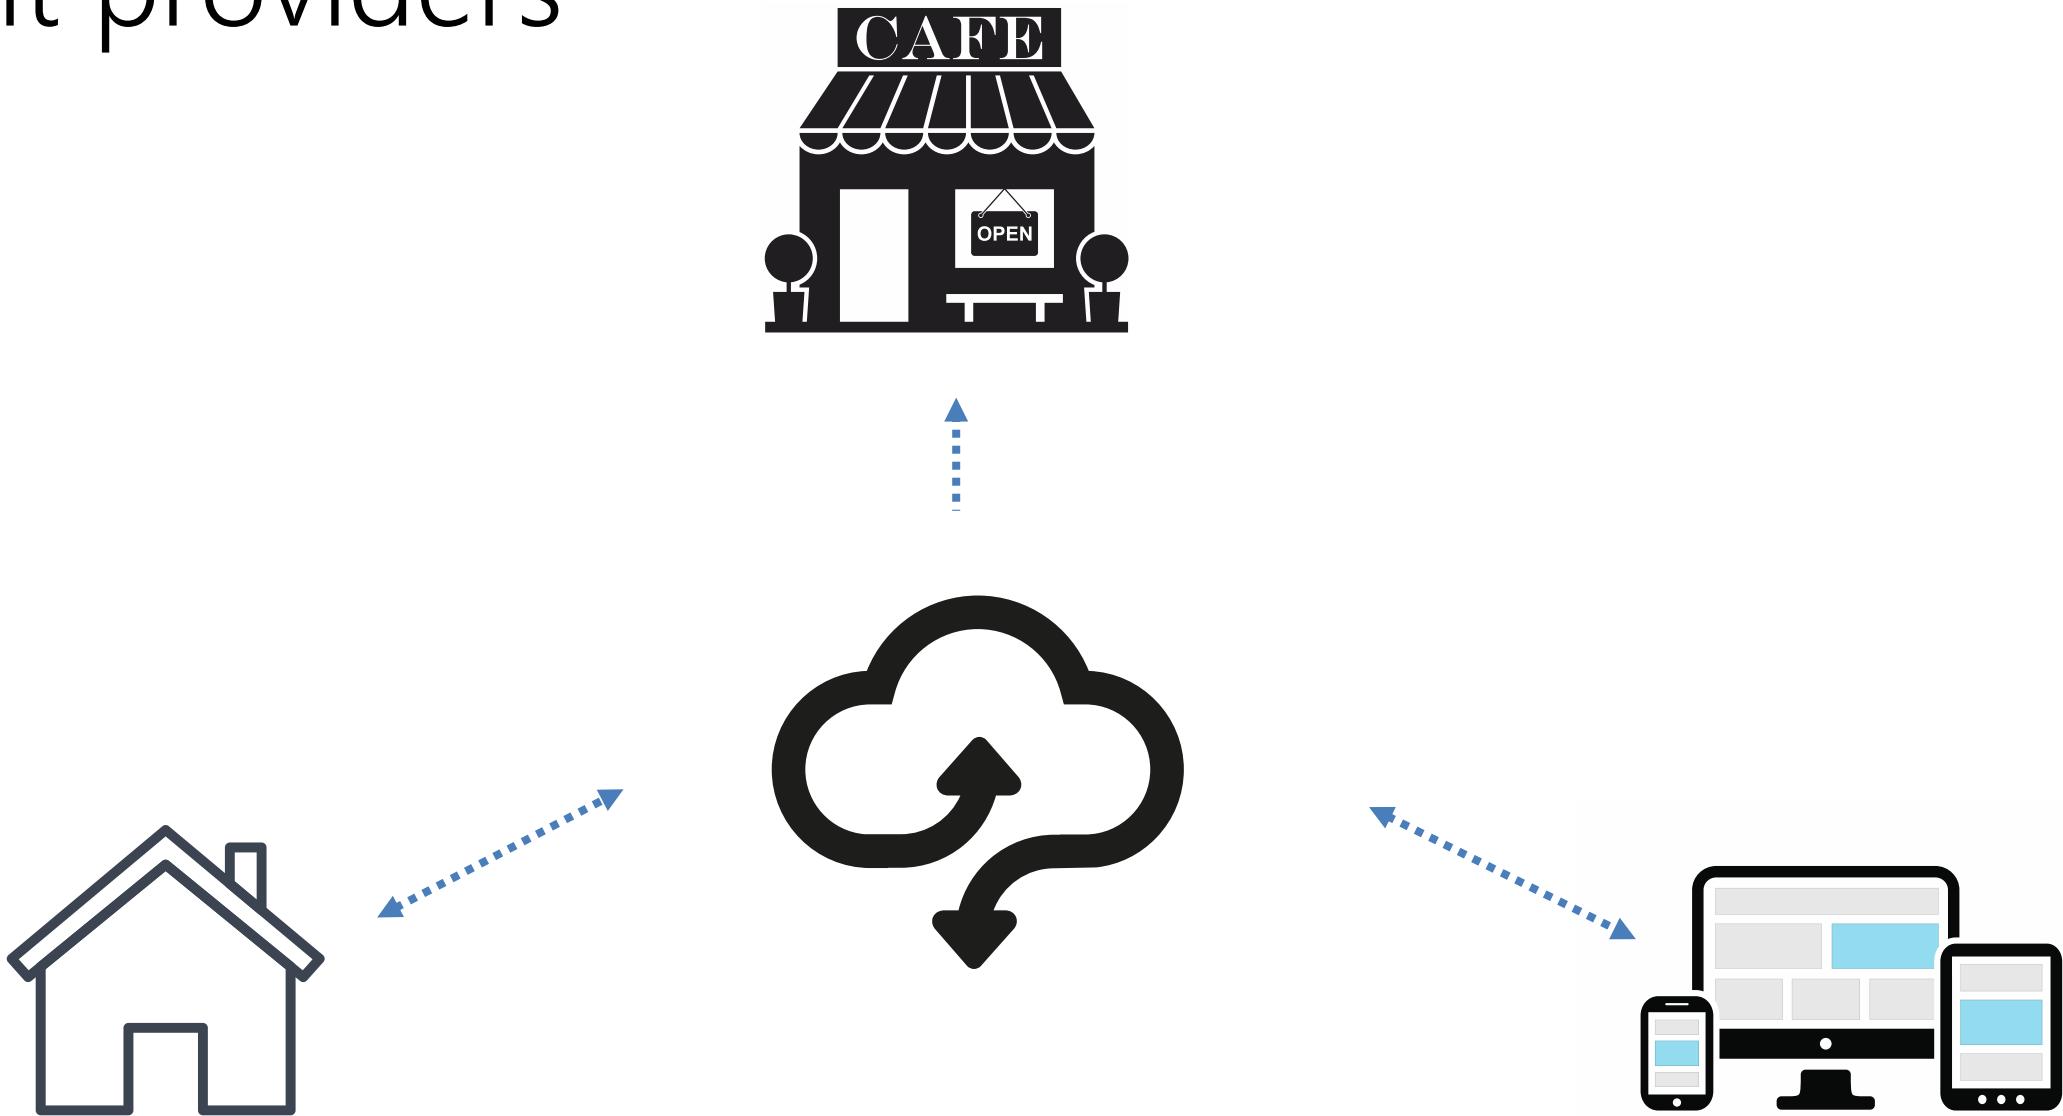
Git providers



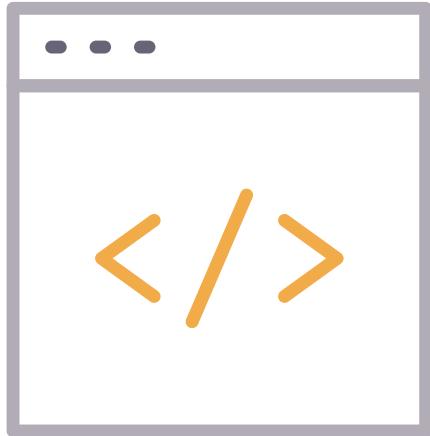
 Bitbucket



Git providers



Git basic commands

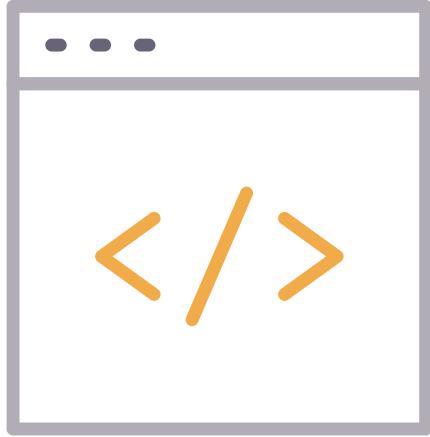


Clone from upstream:

```
$ git clone git://git.kernel.org/pub/scm/.../linux.git my-linux  
$ cd my-linux
```

- **git clone** is used to copy a repository
 - You can also clone a local repository
 - `git clone /path/to/repository`

Git basic commands



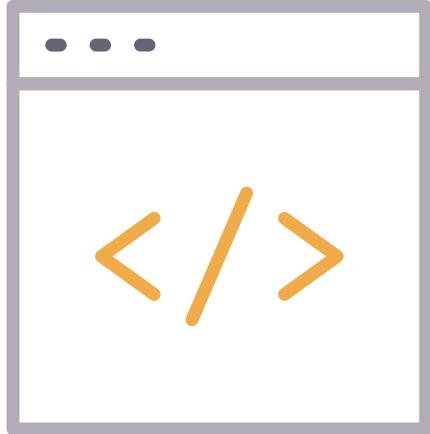
Start a new Git repository for an existing code base

```
$ cd /path/to/my/codebase  
$ git init      (1)  
$ git add .     (2)  
$ git commit    (3)
```

1. Create a `/path/to/my/codebase/.git` directory.
2. Add all existing files to the index.
3. Record the pristine state as the first commit in the history.

- Use `git init` to create a new local repository.
- Alternatively, you can create a repository within a new directory by specifying the project name: `git init [project name]`

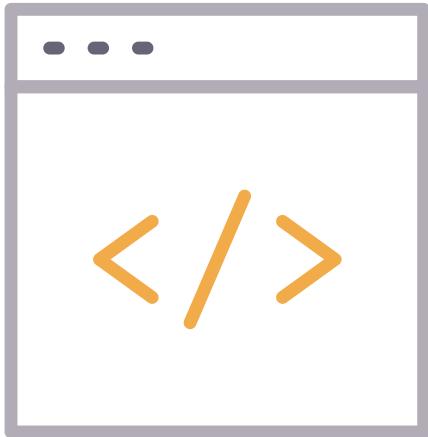
Git basic commands



```
$ git add git-*.*sh
```

- Use `git add <files/directories>` to add files to staging area.

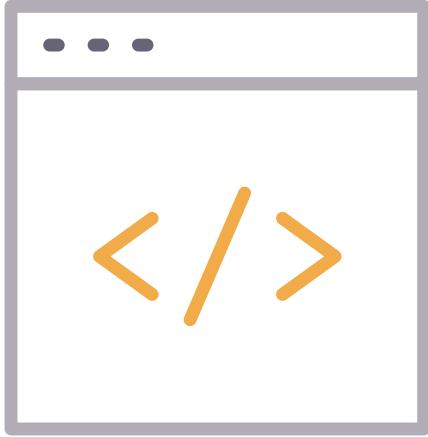
Git basic commands



```
$ git commit
```

- `git commit` will create a snapshot of the changes and save it to the git directory.
- `-m` will pass a commit message: `git commit -m "commit message"`

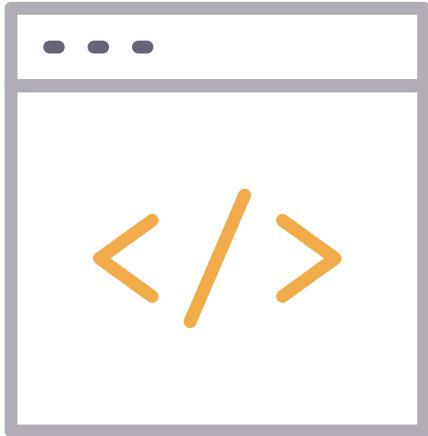
Git basic commands



```
git config --add core.gitproxy '"proxy-command" for example.com'
```

- `git config` can be used to set user-specific configuration values like email, username, file format, and so on.
- `git config --global user.email youremail@example.com`
- `--global` sets value for all repositories

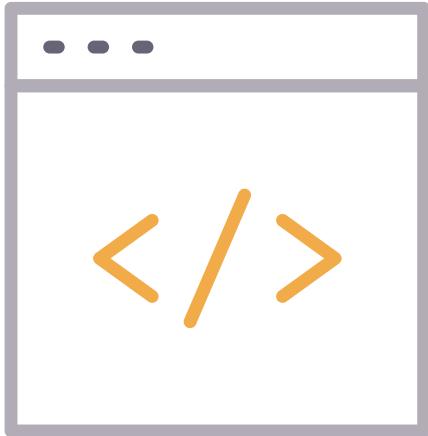
Git basic commands



```
git status [<options>...] [--] [<pathspec>...]
```

- `git status` displays the list of changed files together with the files that are yet to be staged or committed.

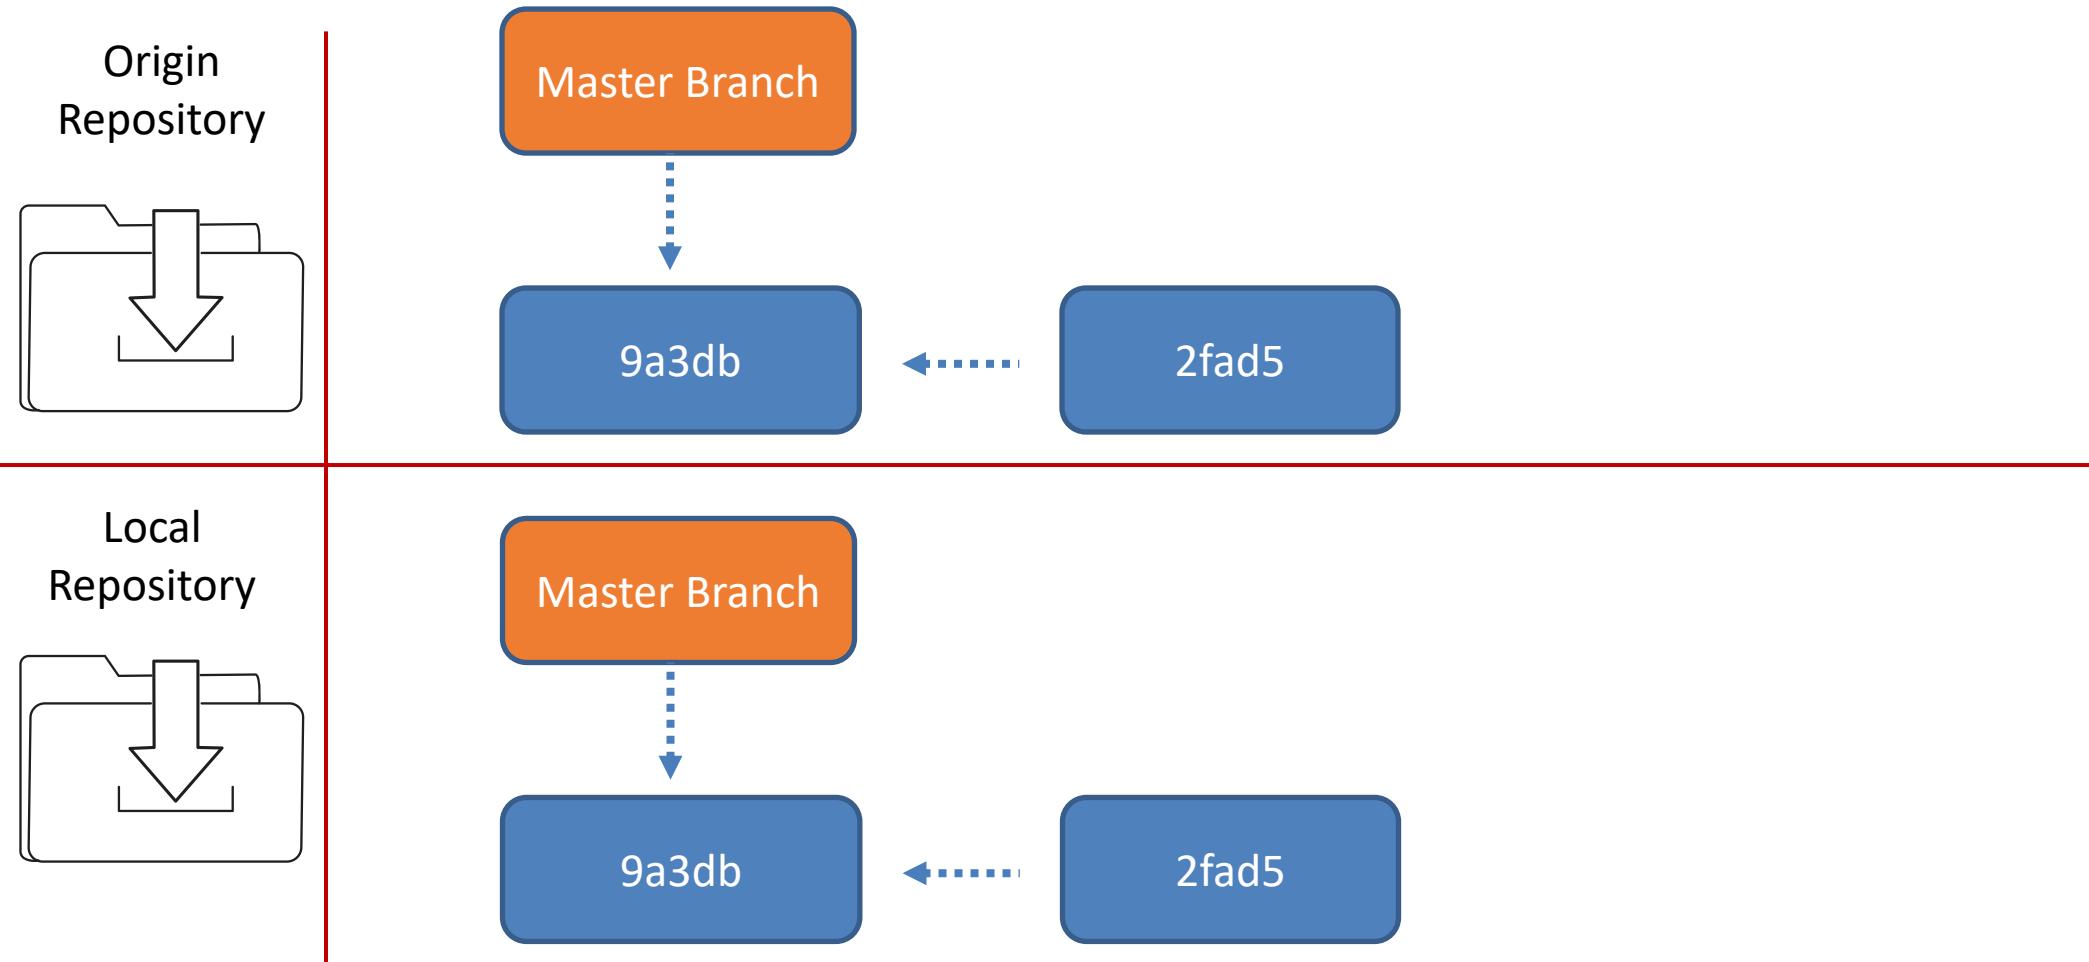
Git basic commands



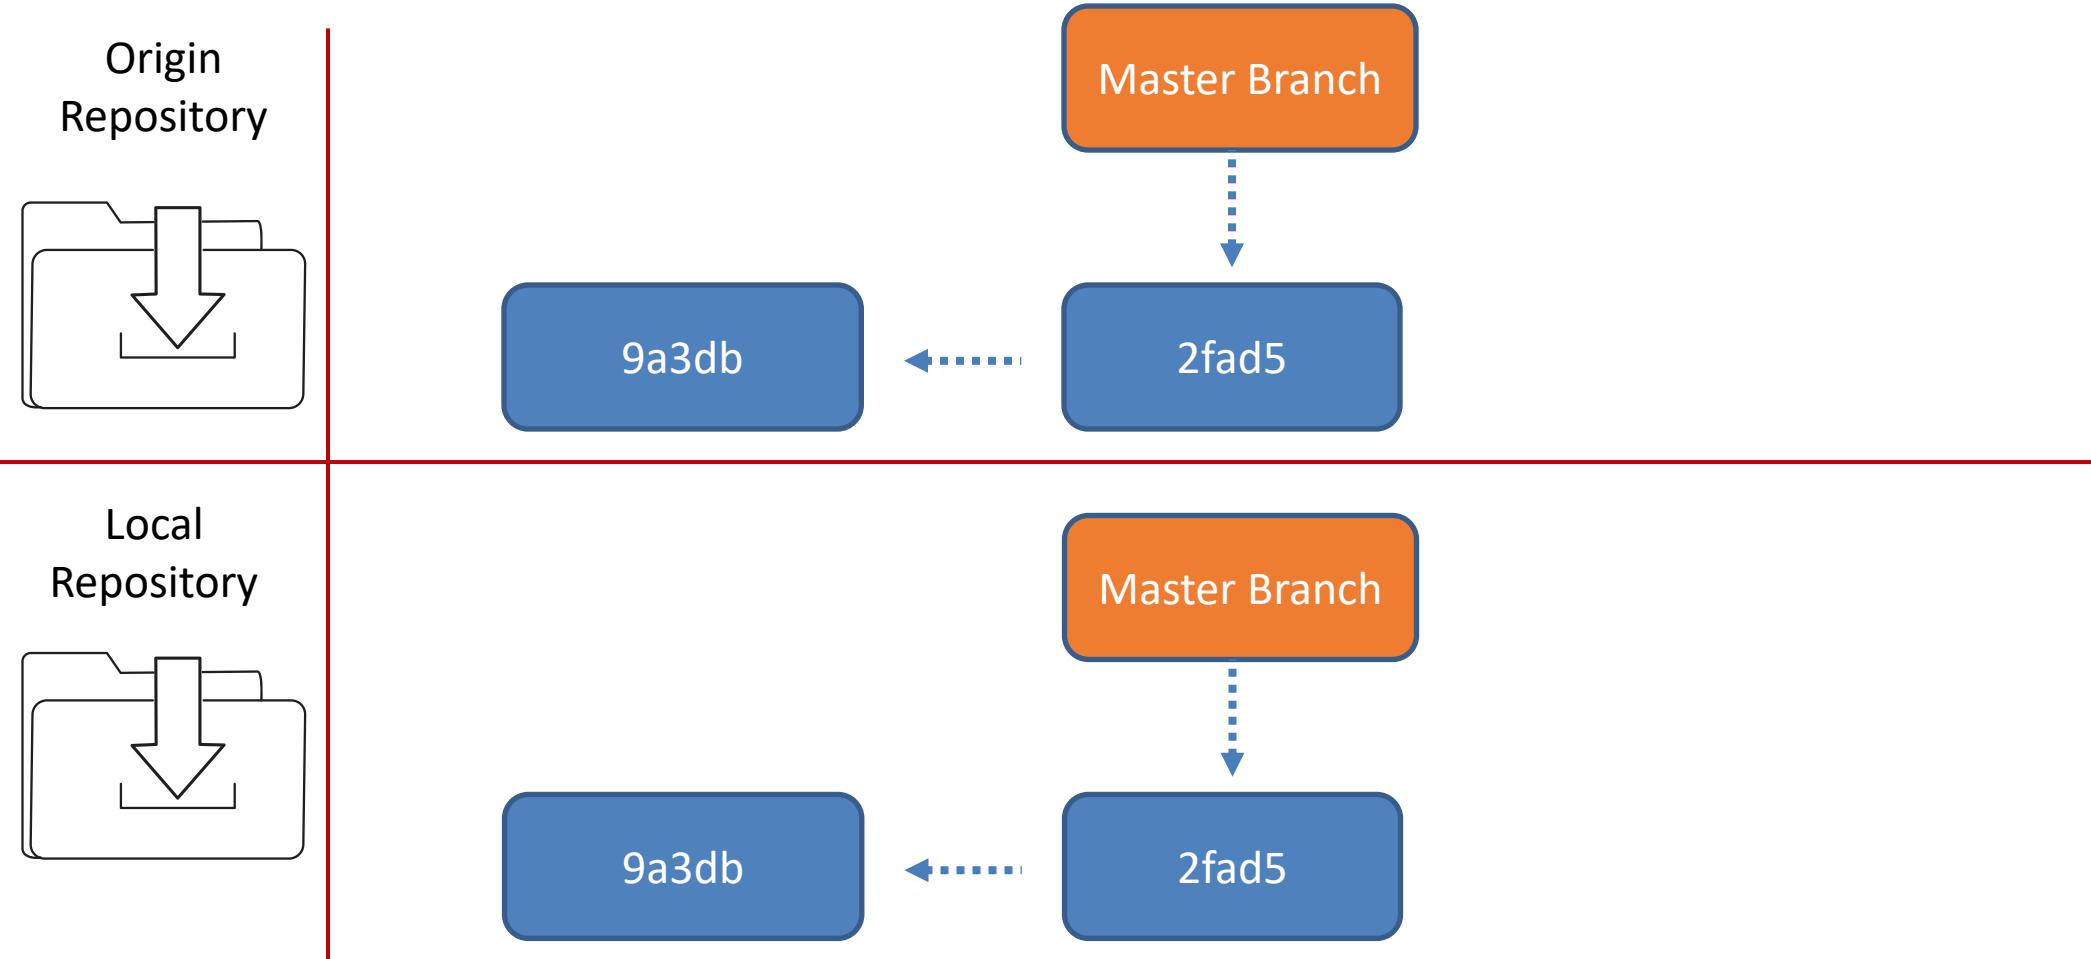
```
git push origin master
```

- **git push** is used to send local commits to the master branch of the remote repository.
- `git push origin <branch>`

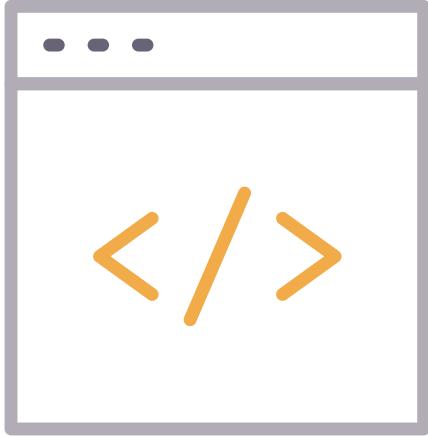
Git commit process



Git commit process



Git diff

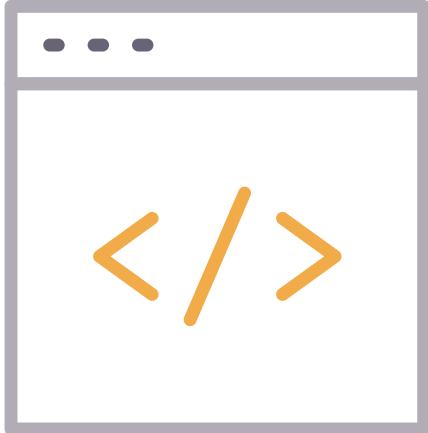


Various ways to check your working tree

```
$ git diff          (1)  
$ git diff --cached (2)  
$ git diff HEAD     (3)
```

- `git diff` lists down conflicts.
- Use it to show changes that have been made but have not been committed.

Git log

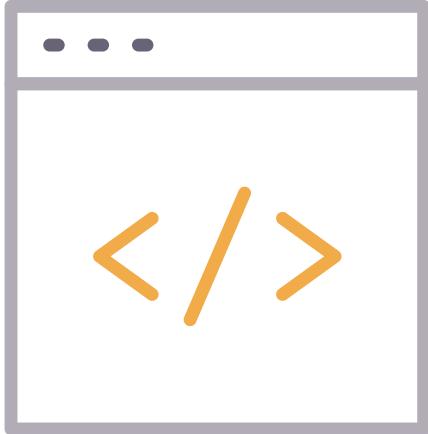


```
git log --no-merges  
Show the whole commit history, but skip any merges
```

```
commit a68112e12f896ca028f25a46e58aaaf46360cb847  
Author: Jason Smith <jason@smithss.org>  
Date:   Sun Jan 6 21:12:35 2019 +0000  
  
    initial
```

- `git log` is used to see the repositories history by listing certain commit details.

Git blame

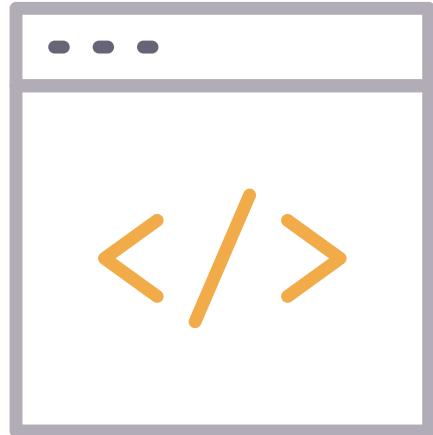


```
git blame CNAME
```

```
a76f1de5 (Jason Smith 2020-04-19 22:52:29 -0700 1) kickstart.innovationin.software  
5948ef48 (Jason Smith 2020-04-19 22:51:19 -0700 2)
```

- `git blame` is used to see which users changed a file.

Git blame



```
$ git commit ...
$ git reset --soft HEAD^      (1)
$ edit                      (2)
$ git commit -a -c ORIG_HEAD (3)
```

Undo commits permanently

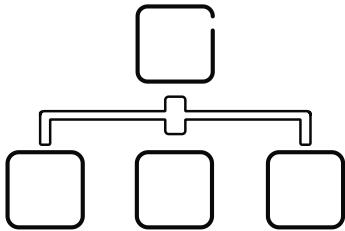
```
$ git commit ...
$ git reset --hard HEAD~3    (1)
```

- `git reset` **undo changes**
- `--soft` `--mixed` or `--hard` **depending on situation.**
- `--soft` Does NOT change local working directory. Index still points to latest commit
- `--mixed` Does NOT change local working directory. Index points to reverted commit
- `--hard` Changes local working directory and index to point to reverted commit. Deletes all comments and files from original commit.

Lab: Git

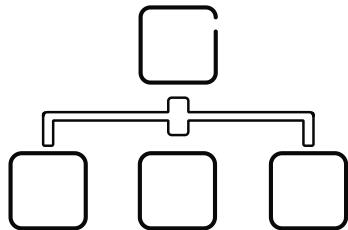


Software Development flows



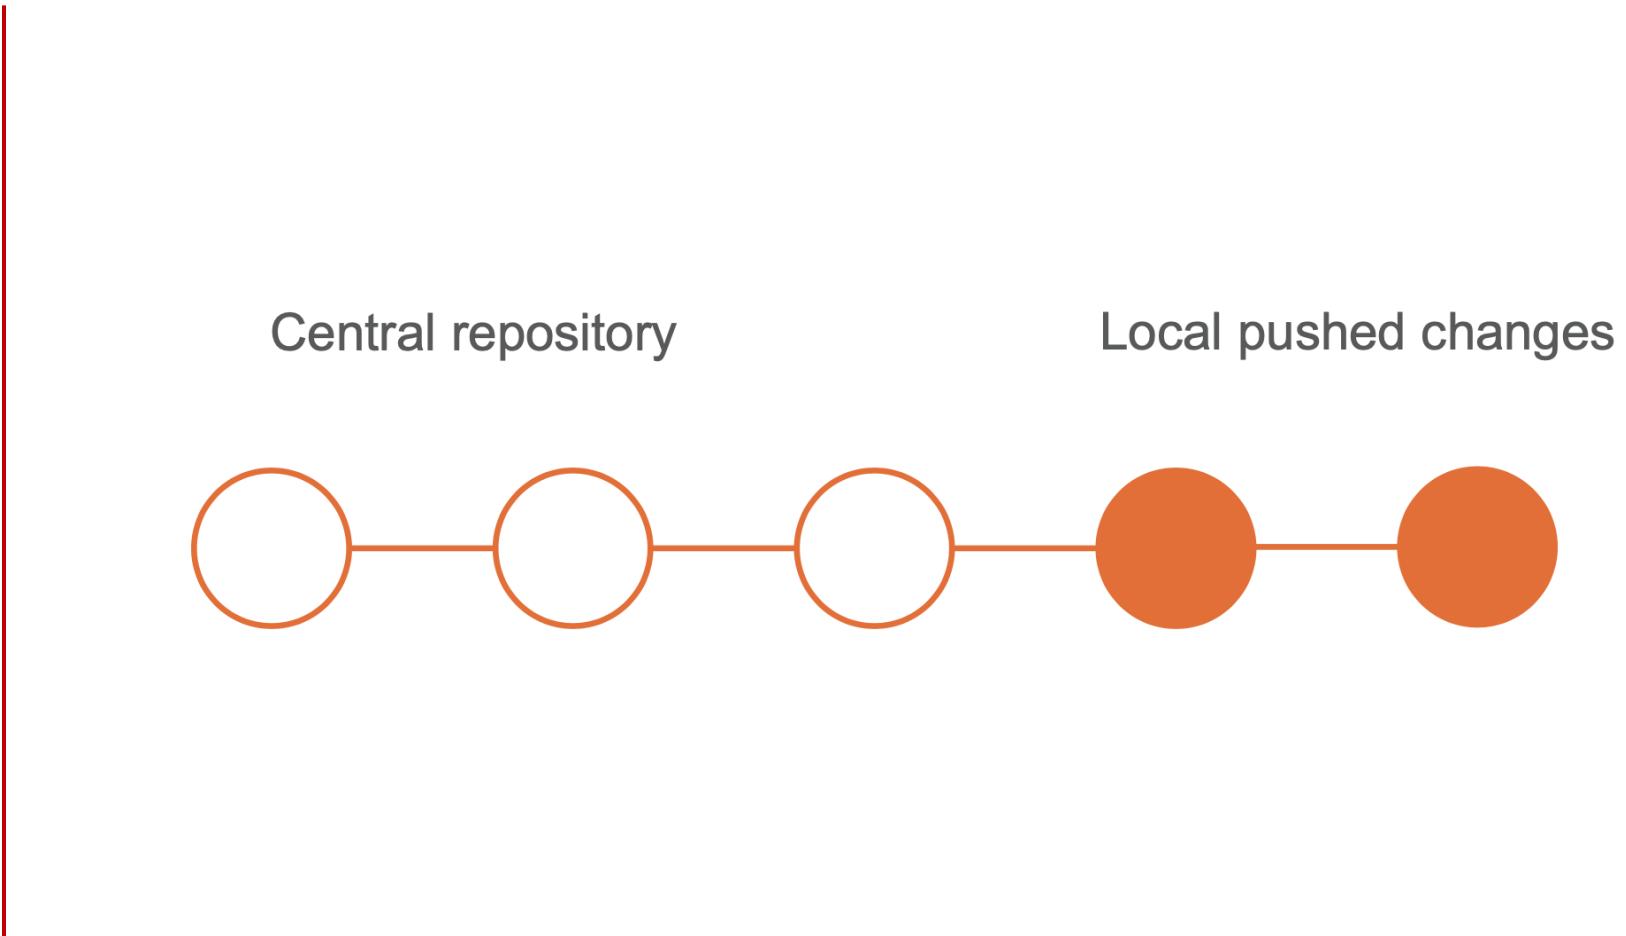
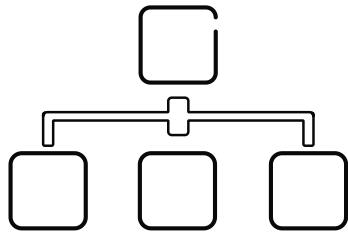
- Centralized flow
- GitHub
- Featured branch

Centralized flow

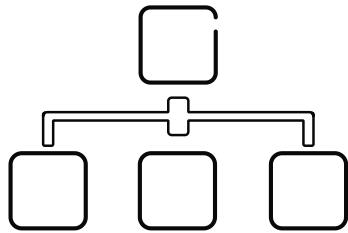


- Uses a centralized repository
- End users 'clone' repository
- Pull changes from central repository to local
- Make changes
- Push local changes back to centralized repository
- Similar to SVN

Centralized flow

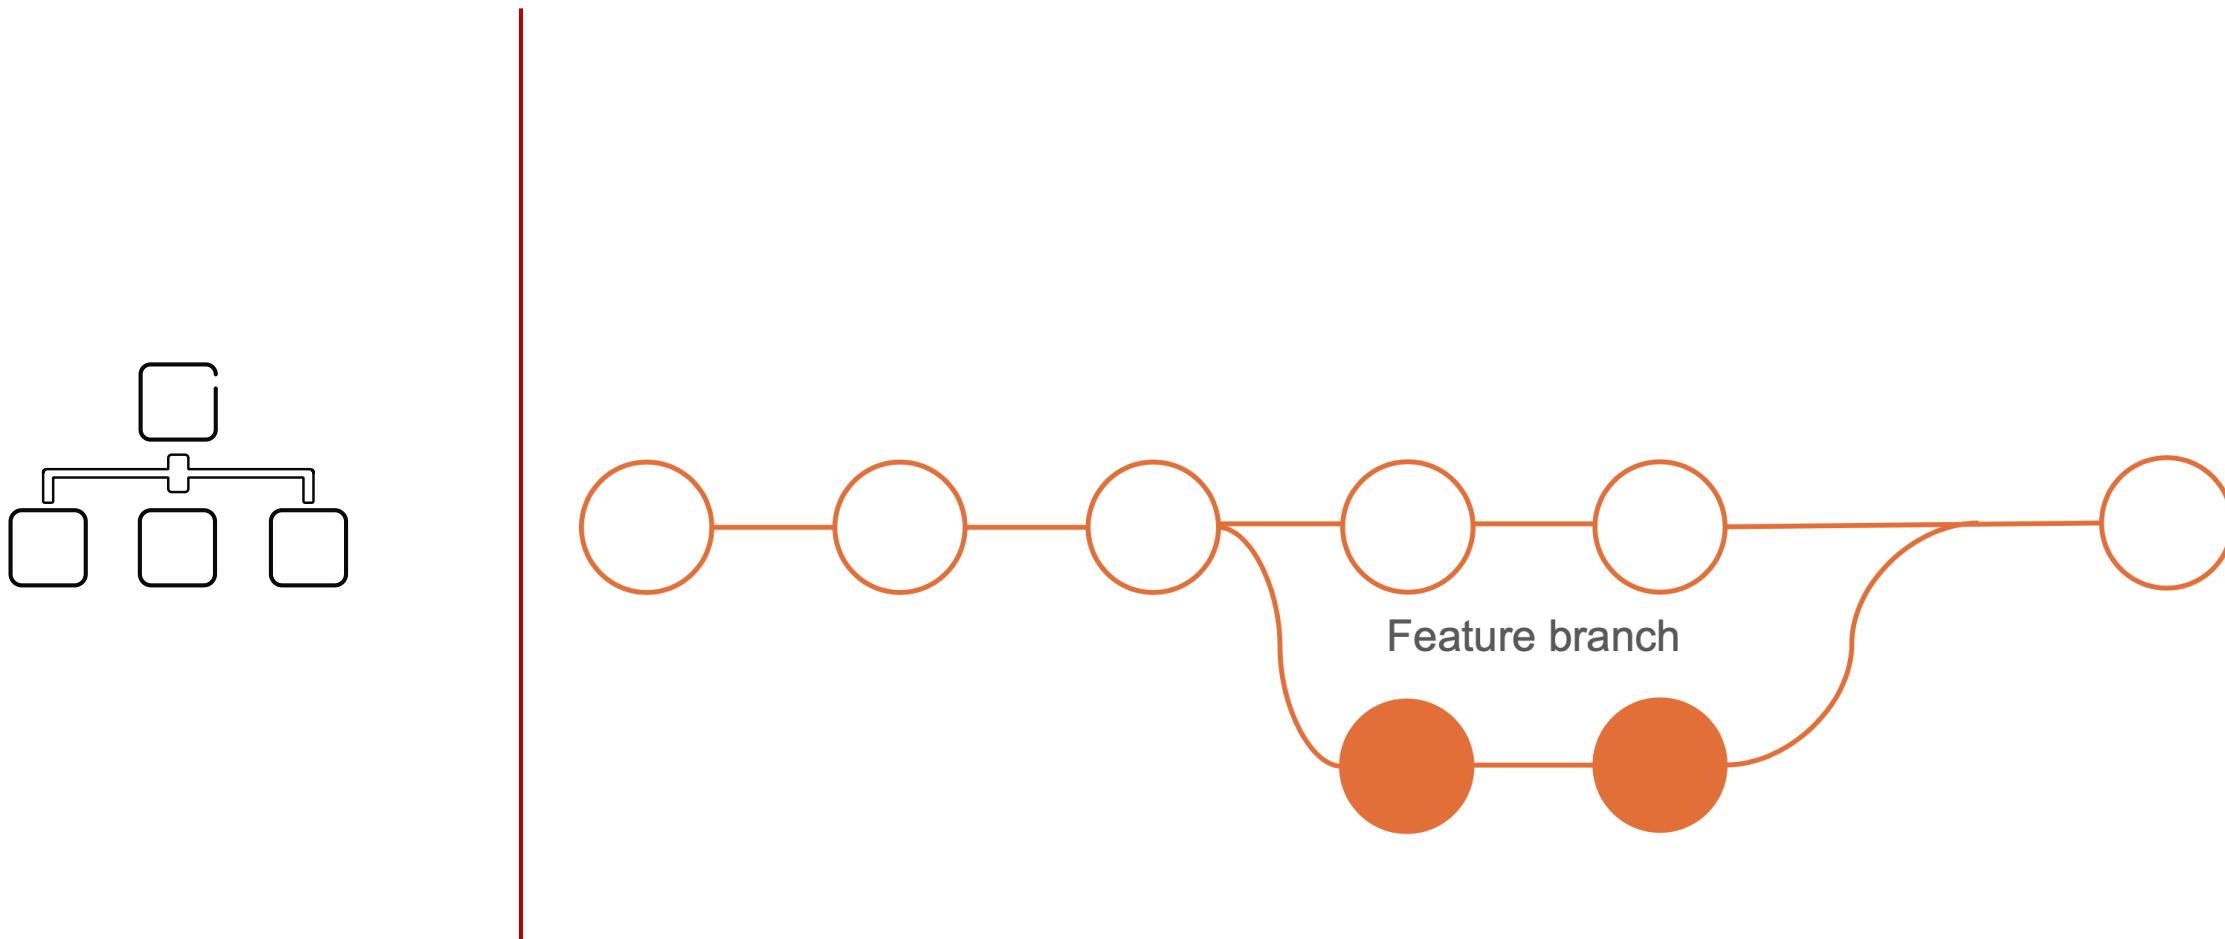


Feature branch flow

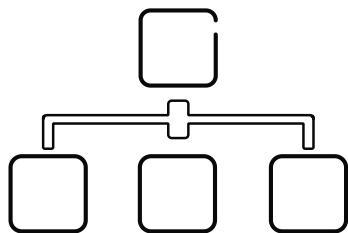


- Create a feature branch
- Add new code
- Open a pull request
- Merge code back into Master

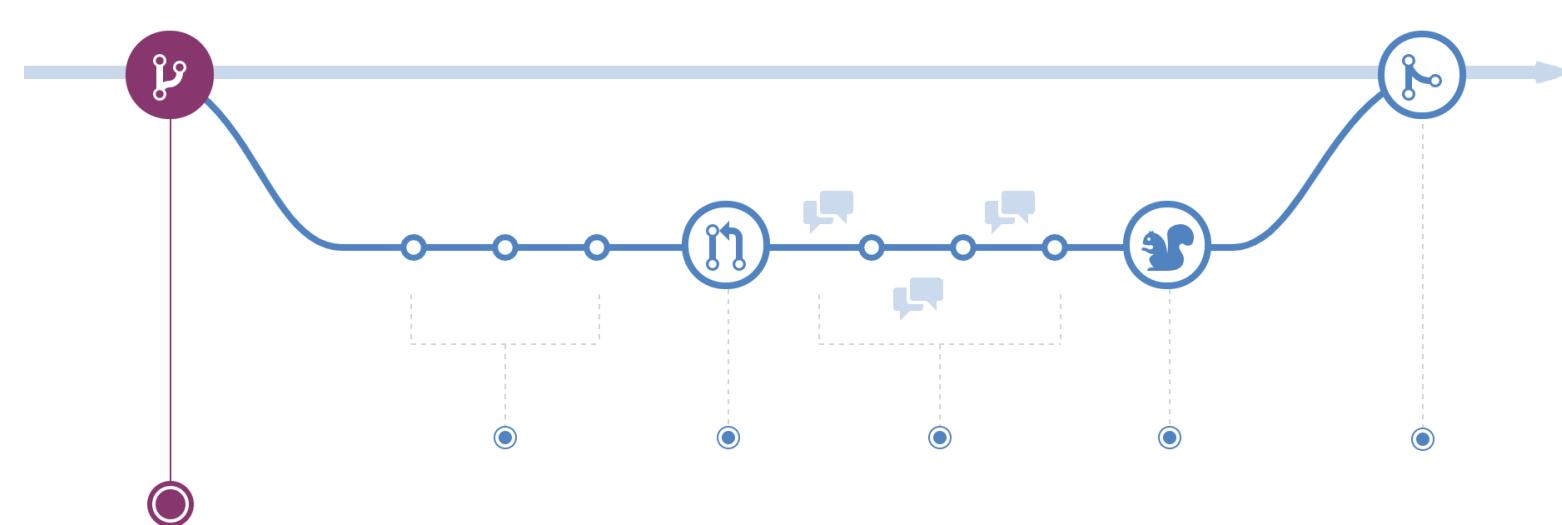
Feature branch flow



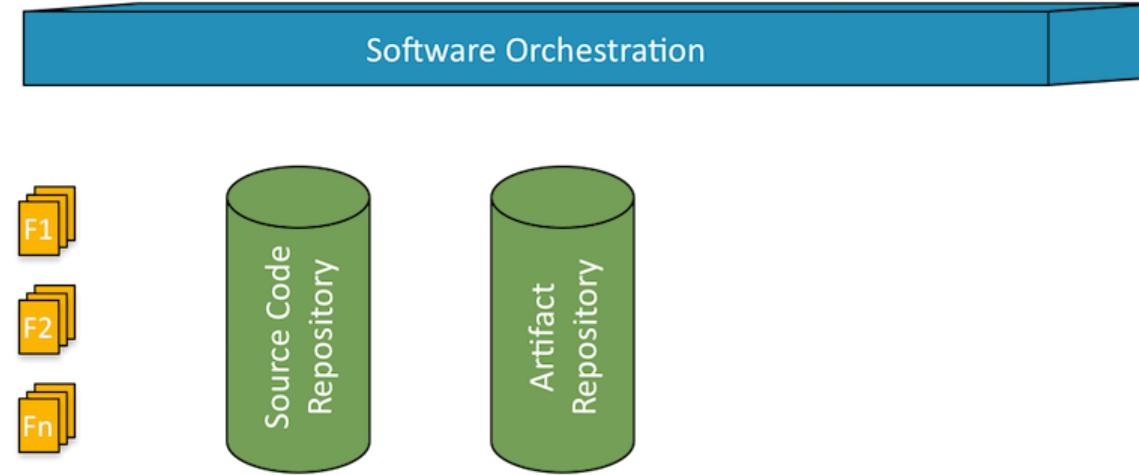
GitHub flow



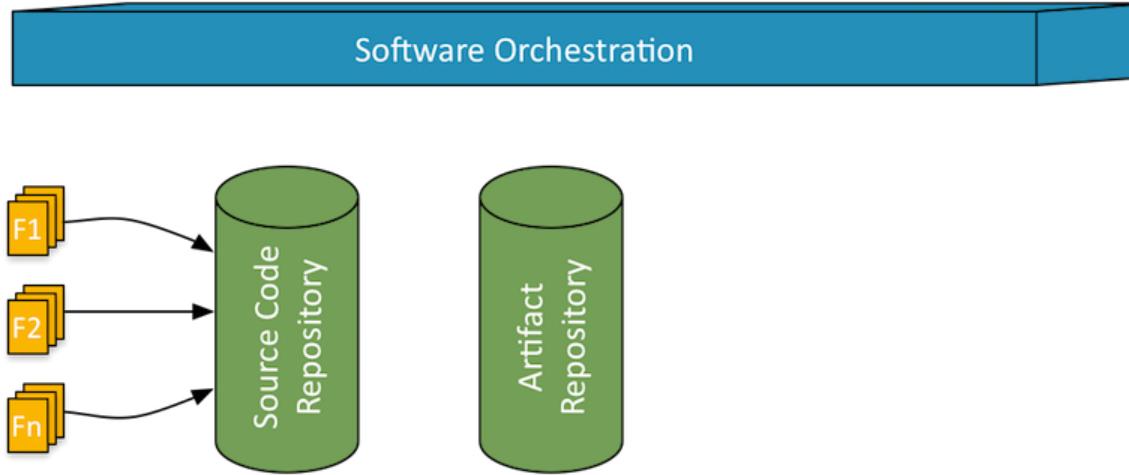
1. Create a branch



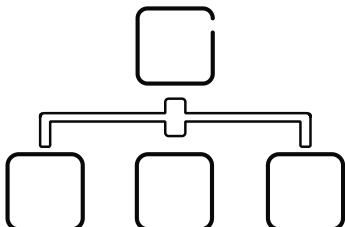
GitHub flow



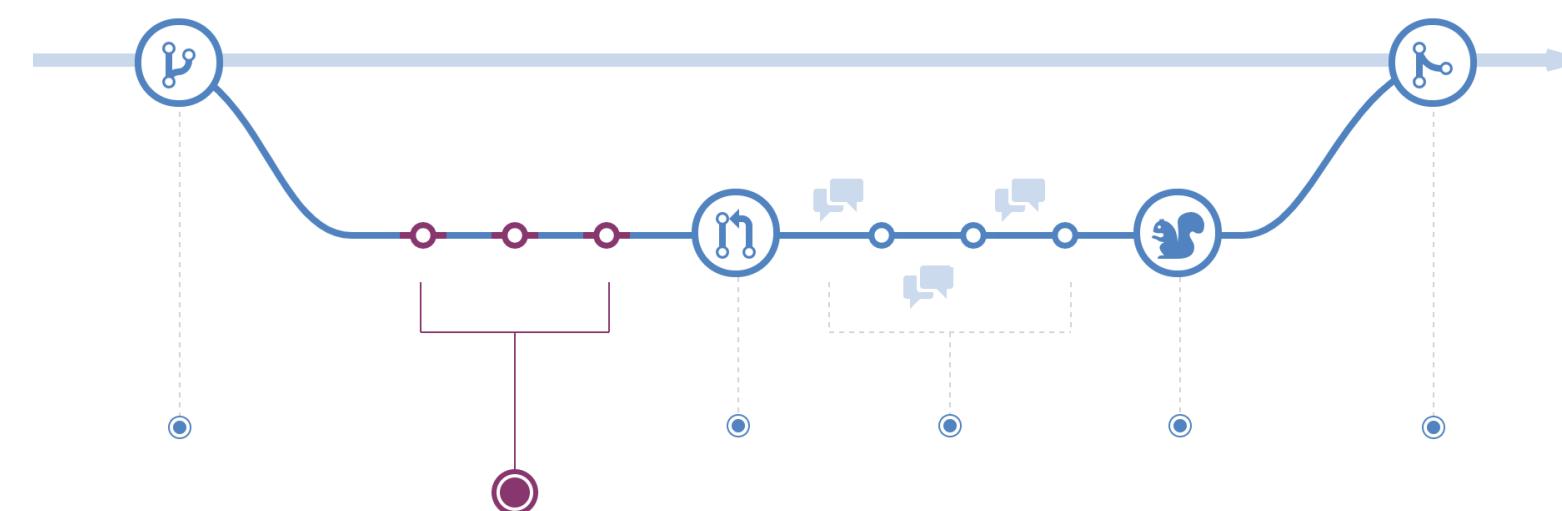
GitHub flow



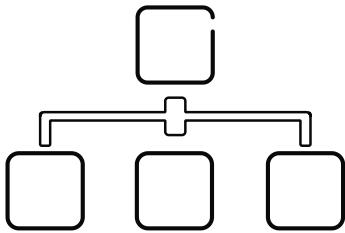
GitHub flow



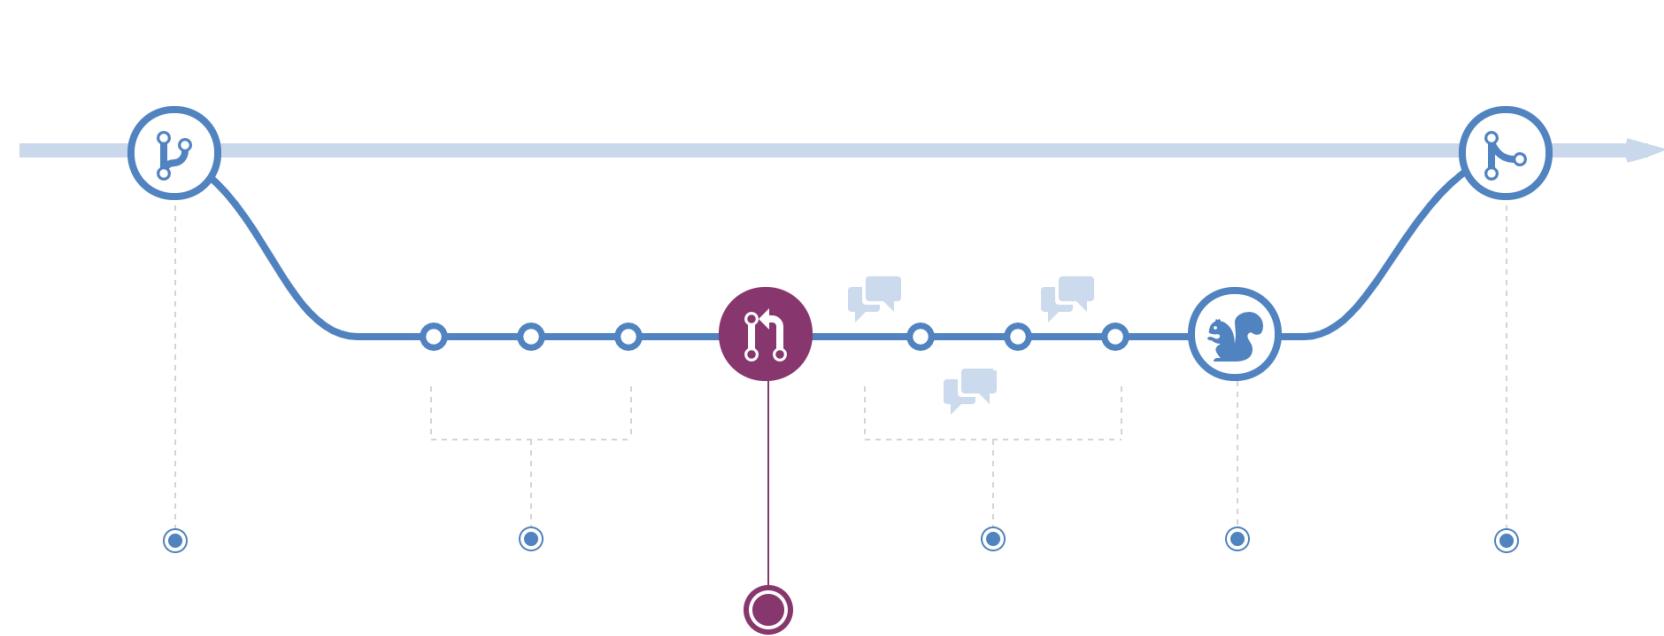
1. Create a branch
2. Checkout branch and make changes



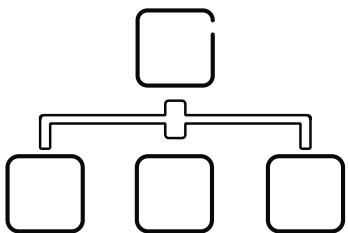
GitHub flow



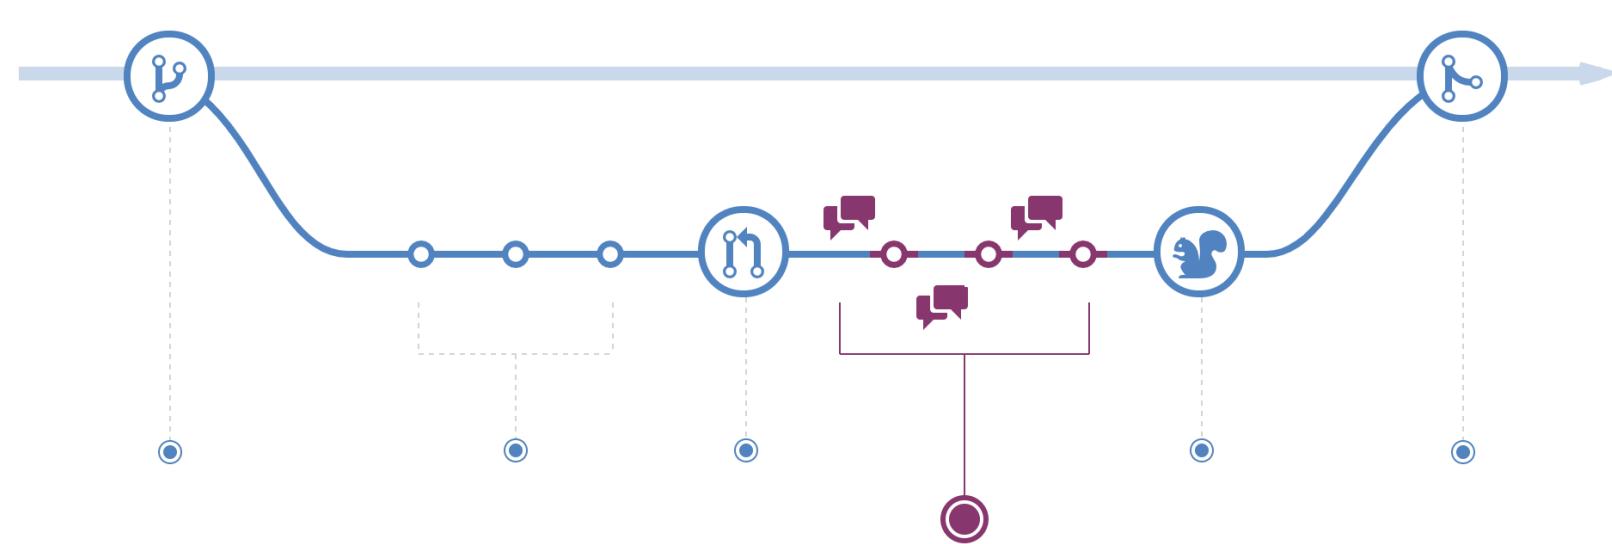
1. Create a branch
2. Checkout branch and make changes
3. Open pull request



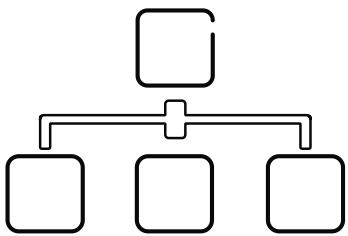
GitHub flow



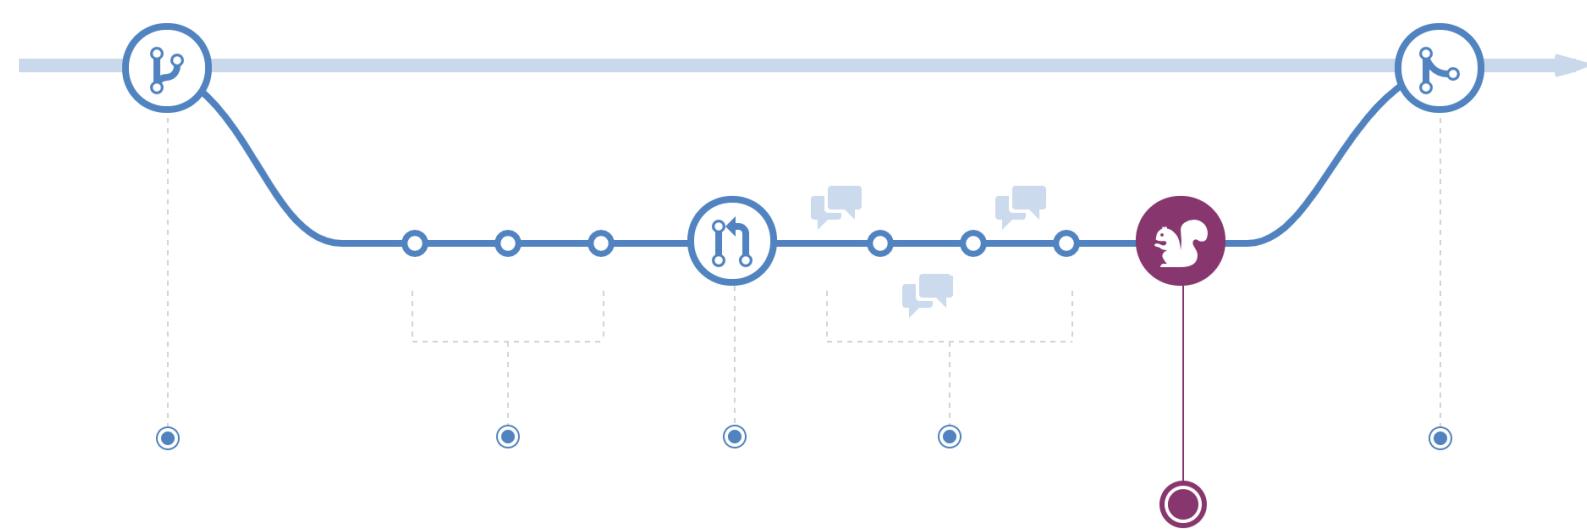
1. Create a branch
2. Checkout branch and make changes
3. Open pull request
4. Peer review



GitHub flow



1. Create a branch
2. Checkout branch and make changes
3. Open pull request
4. Peer review
5. Deploy



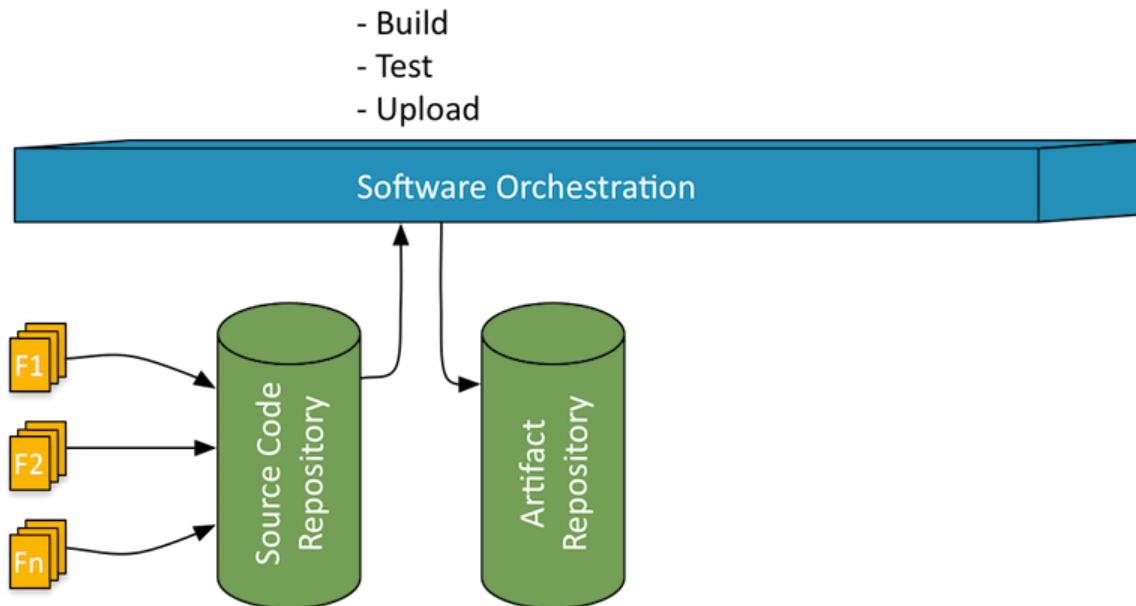
POP QUIZ: Your development flow



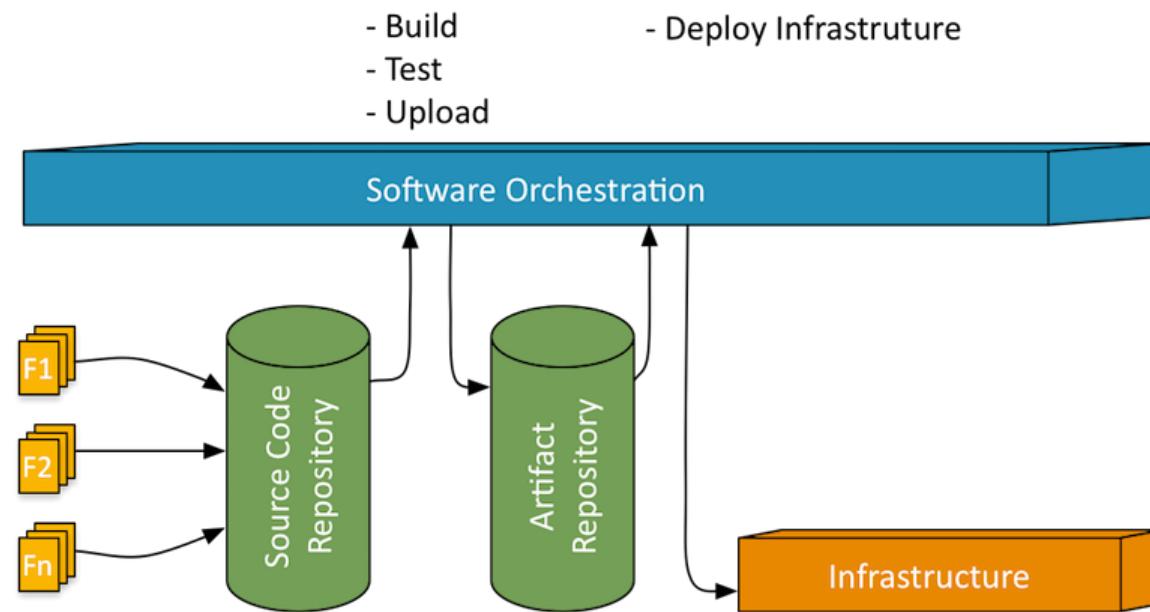
How does your team add new features to your applications?

- Feature branches?
- All changes to master?

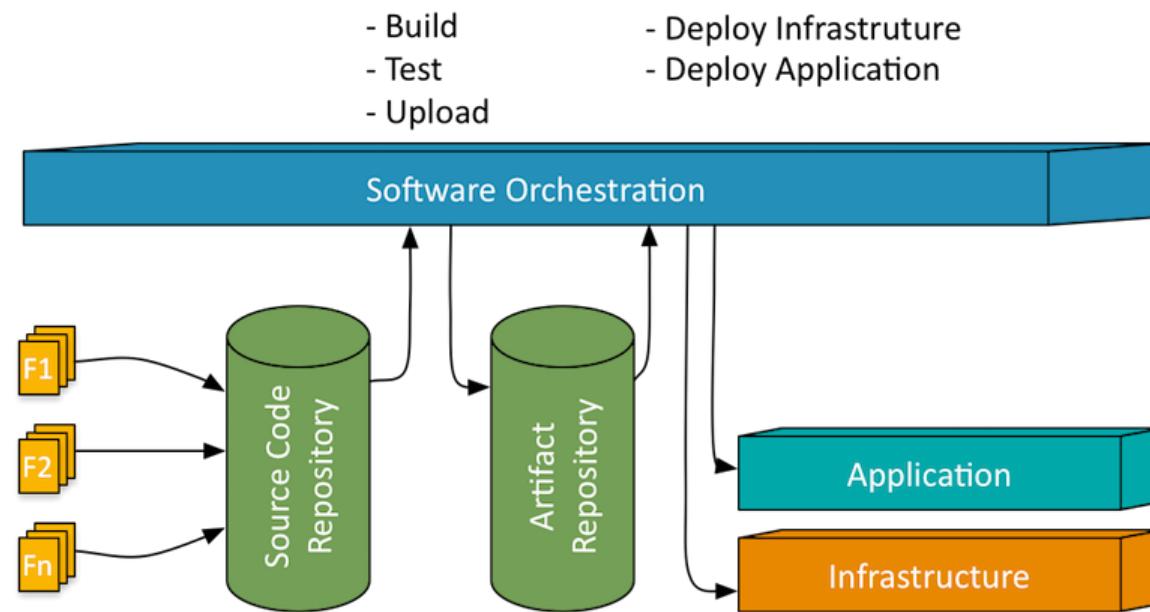
Continuous Integration & Testing



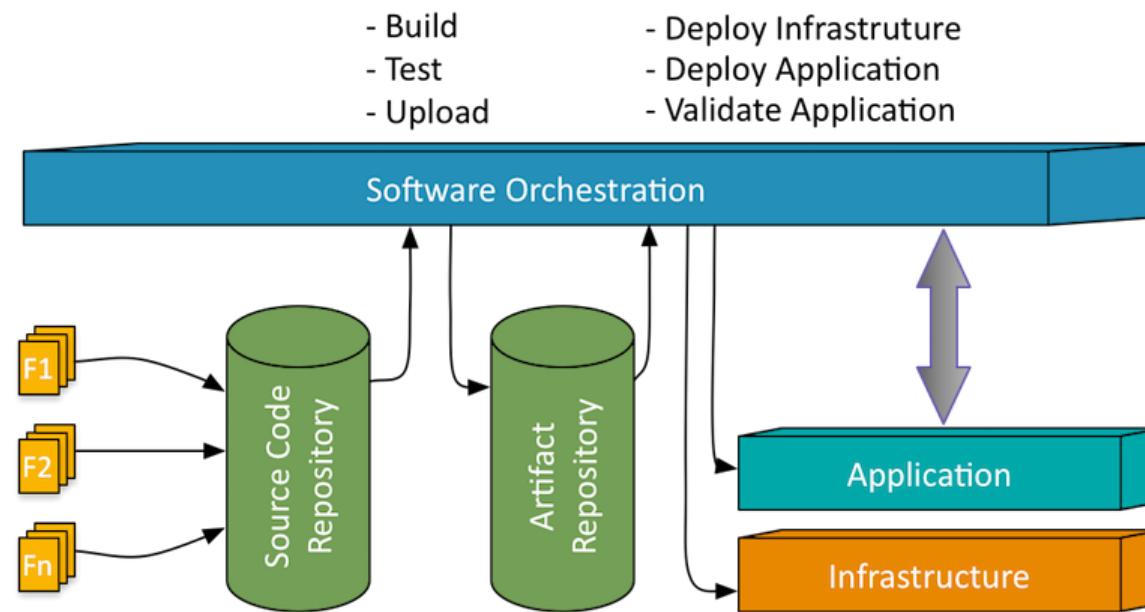
Continuous Release and Deployment



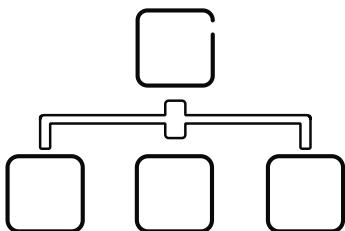
Continuous Release and Deployment



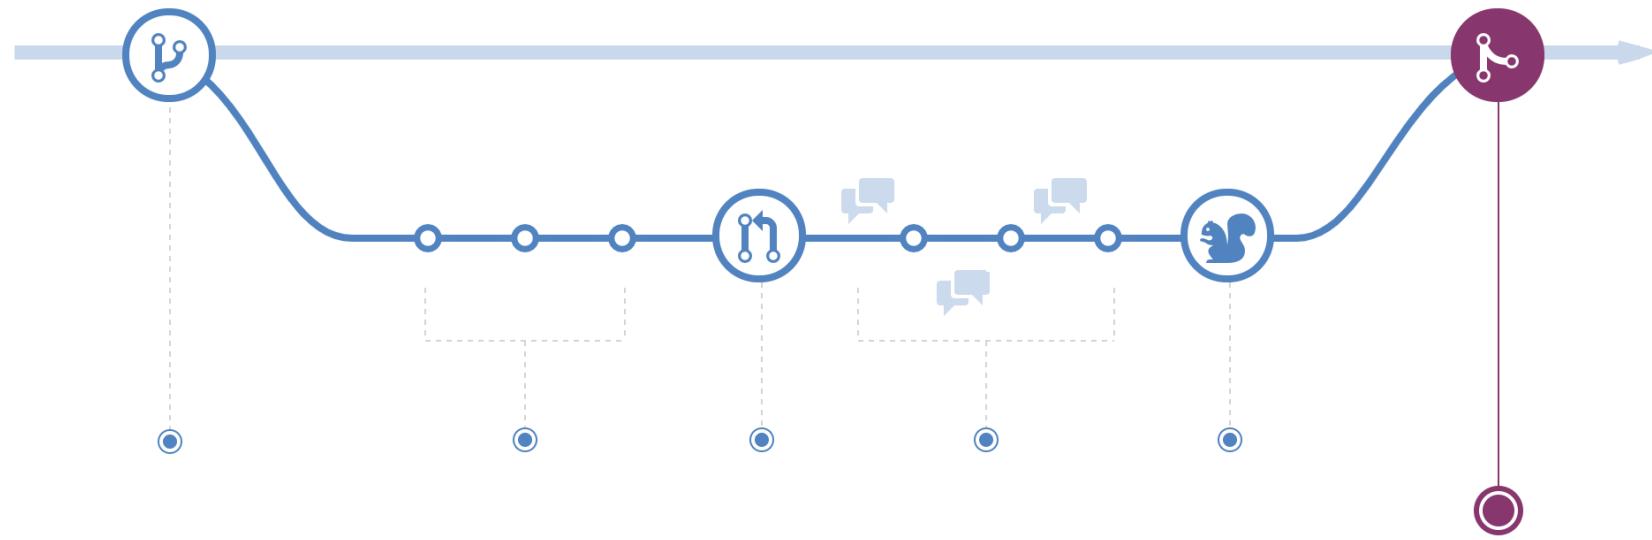
Continuous Release and Deployment



GitHub flow



1. Create a branch
 2. Checkout branch and make changes
 3. Open pull request
 4. Peer review
 5. Deploy
 6. Merge into Master branch



Configuration Management



Configuration Drift



POTHOLE

- Your infrastructure requirements change
- Configuration of a server falls out of policy
- Manage with Infrastructure as Code (IAC)
 - Terraform
 - Ansible
 - Chef
 - Puppet

Manual



POP QUIZ: Manual tasks



What manual tasks do you deal with repeatedly?

Can they be automated?

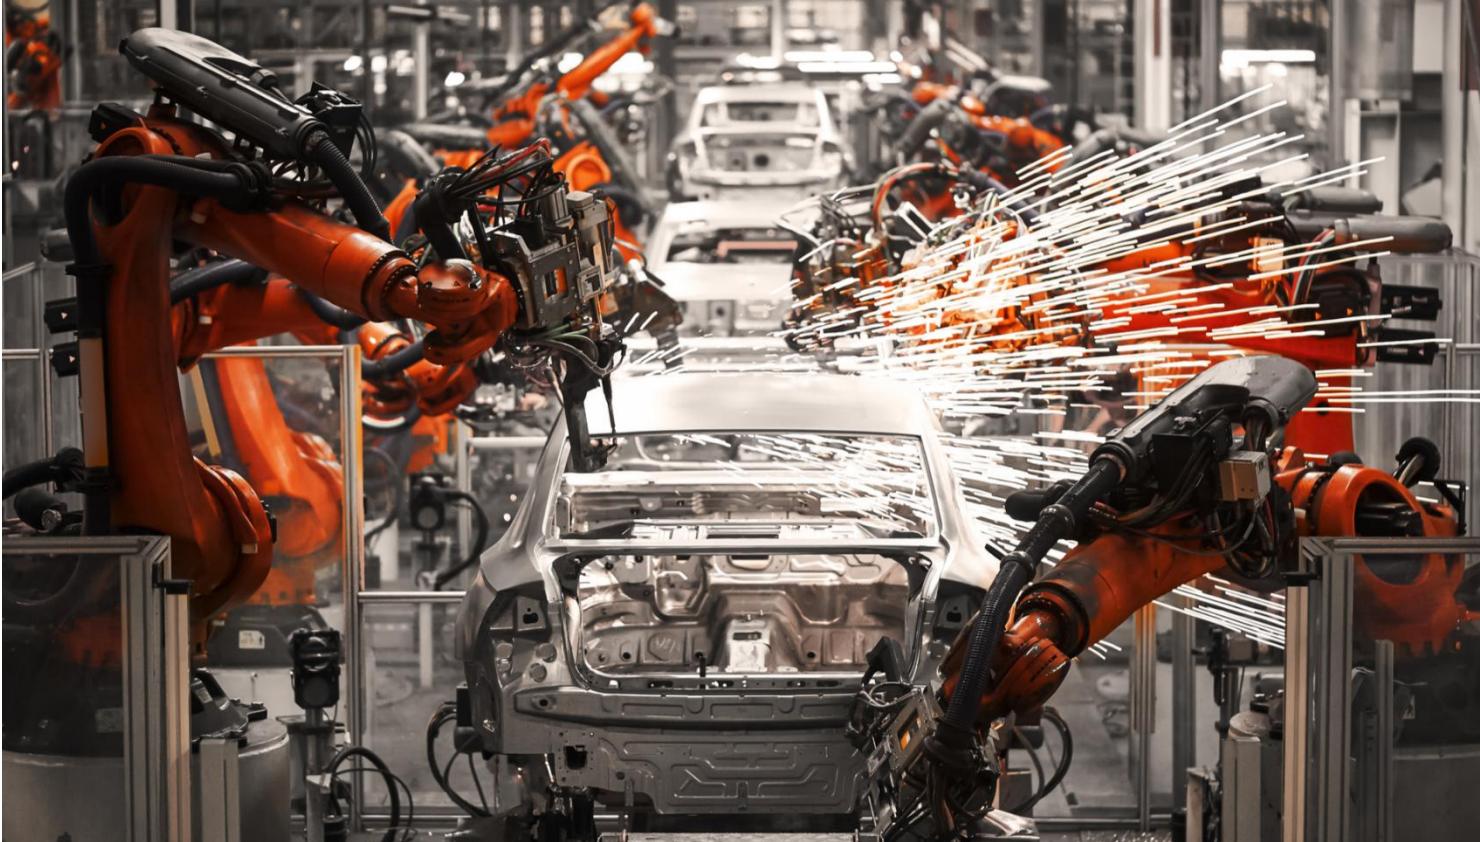
Common manual tasks



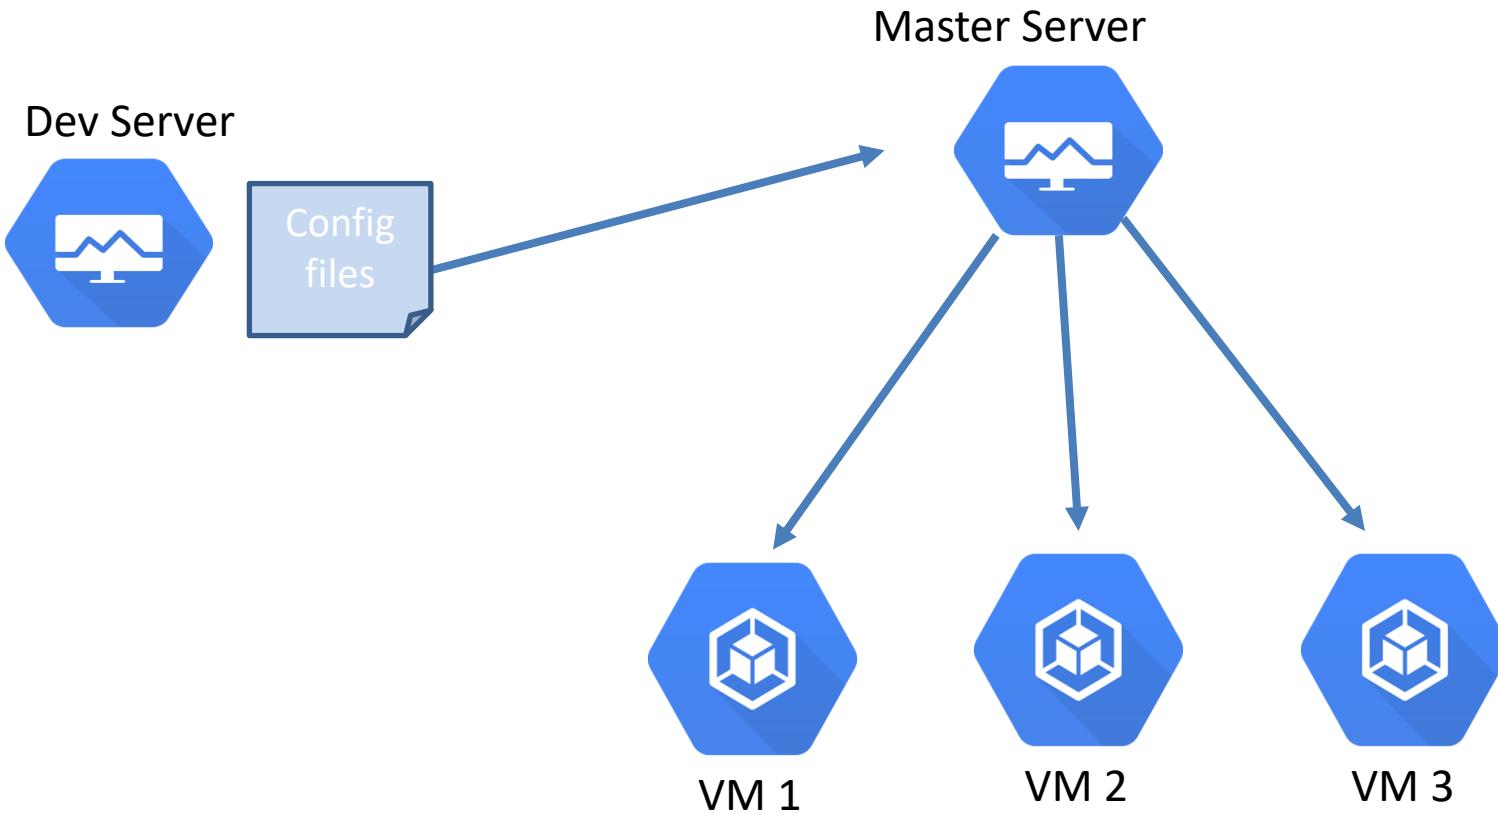
Installation and configuration:

- Operating System
 - Runtime
 - Libraries
 - Utilities
 - Configuration files
-
- Build application
 - Test application
 - Deploy application

Automated



Automated



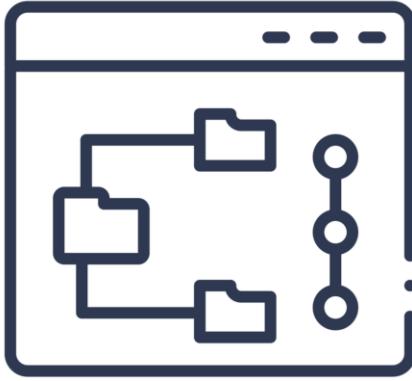
Terraform



Automating Infrastructure



Provisioning resources



Version Control

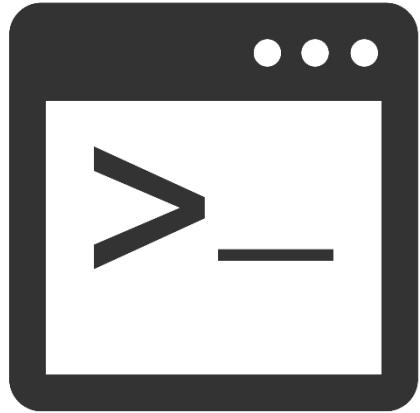


Plan Updates

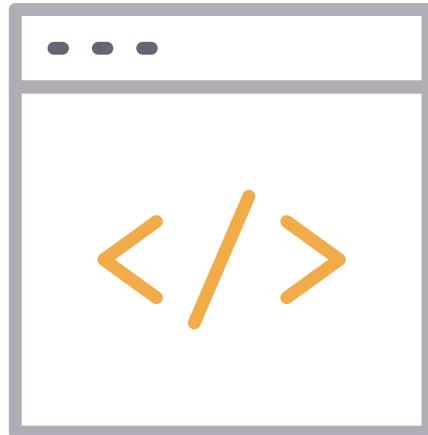


Reusable
Templates

Terraform components



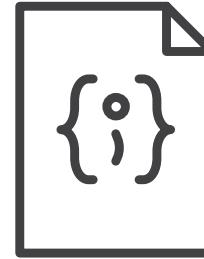
Terraform executable



Terraform files

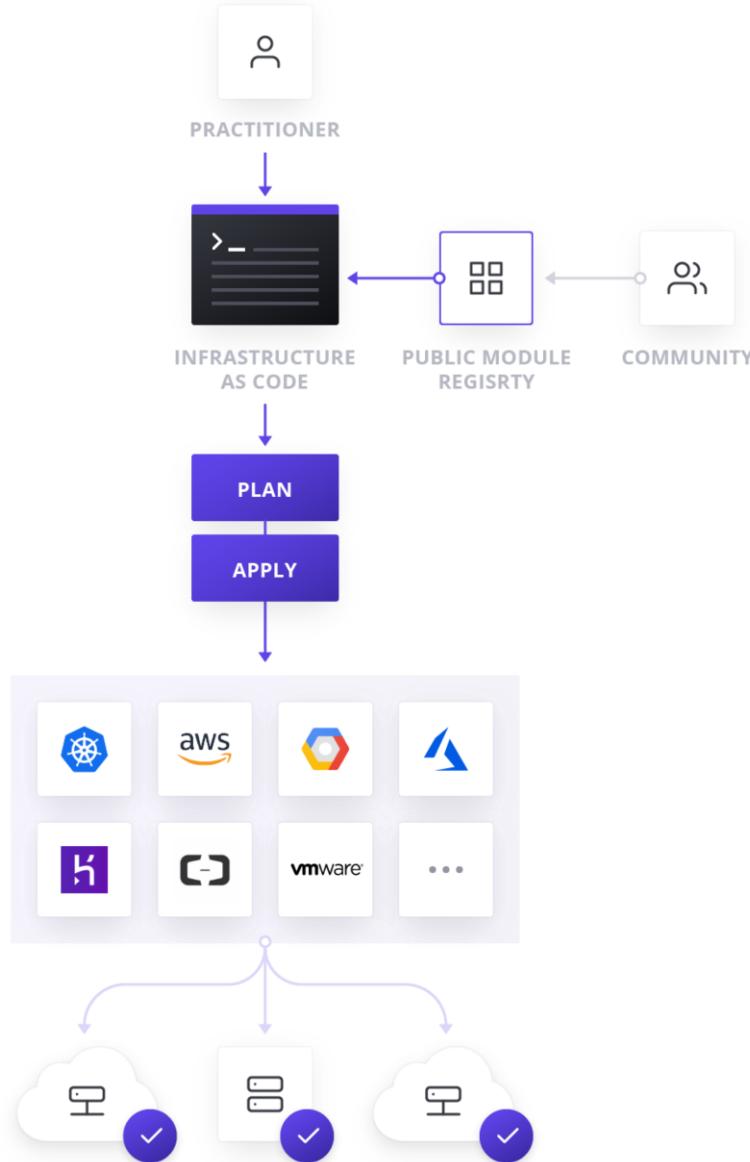


Terraform
plugins

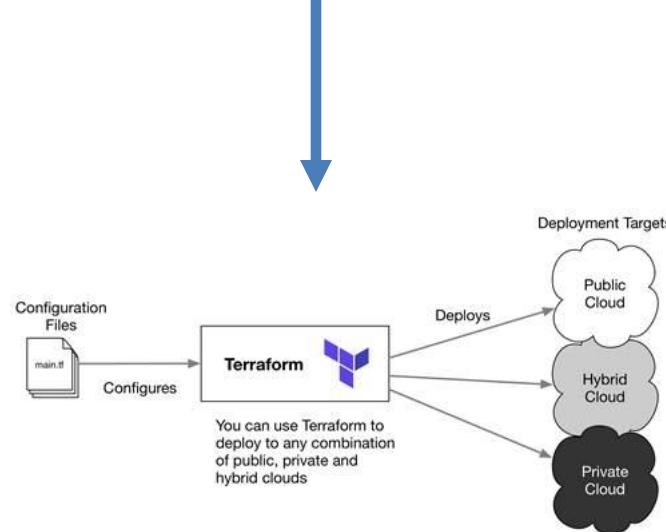
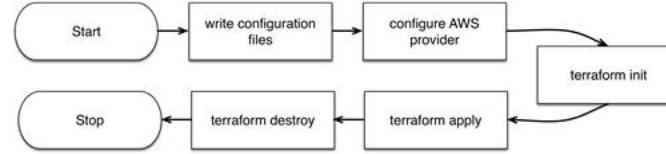


Terraform
state

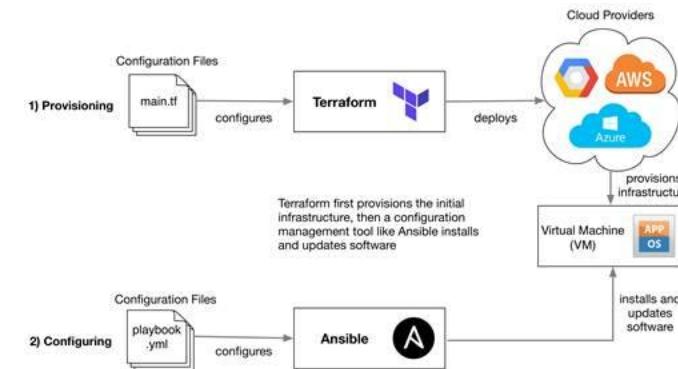
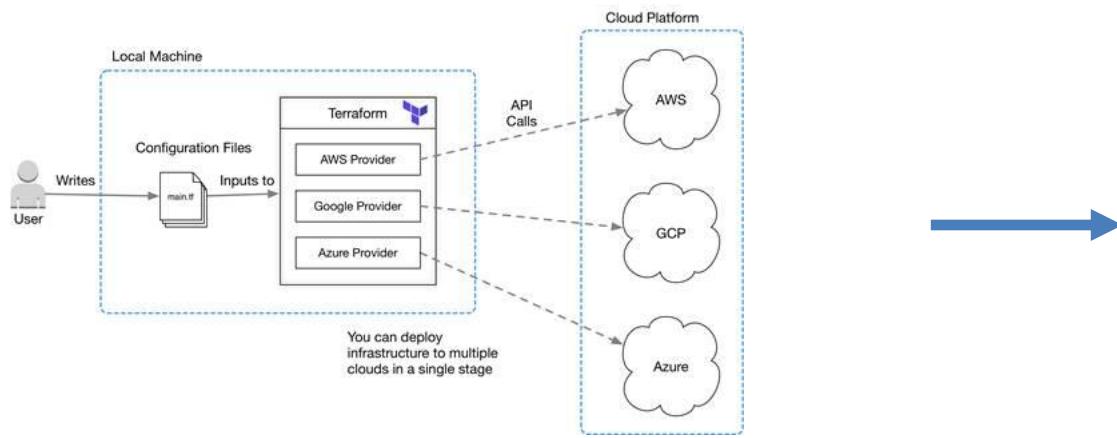
Terraform architecture



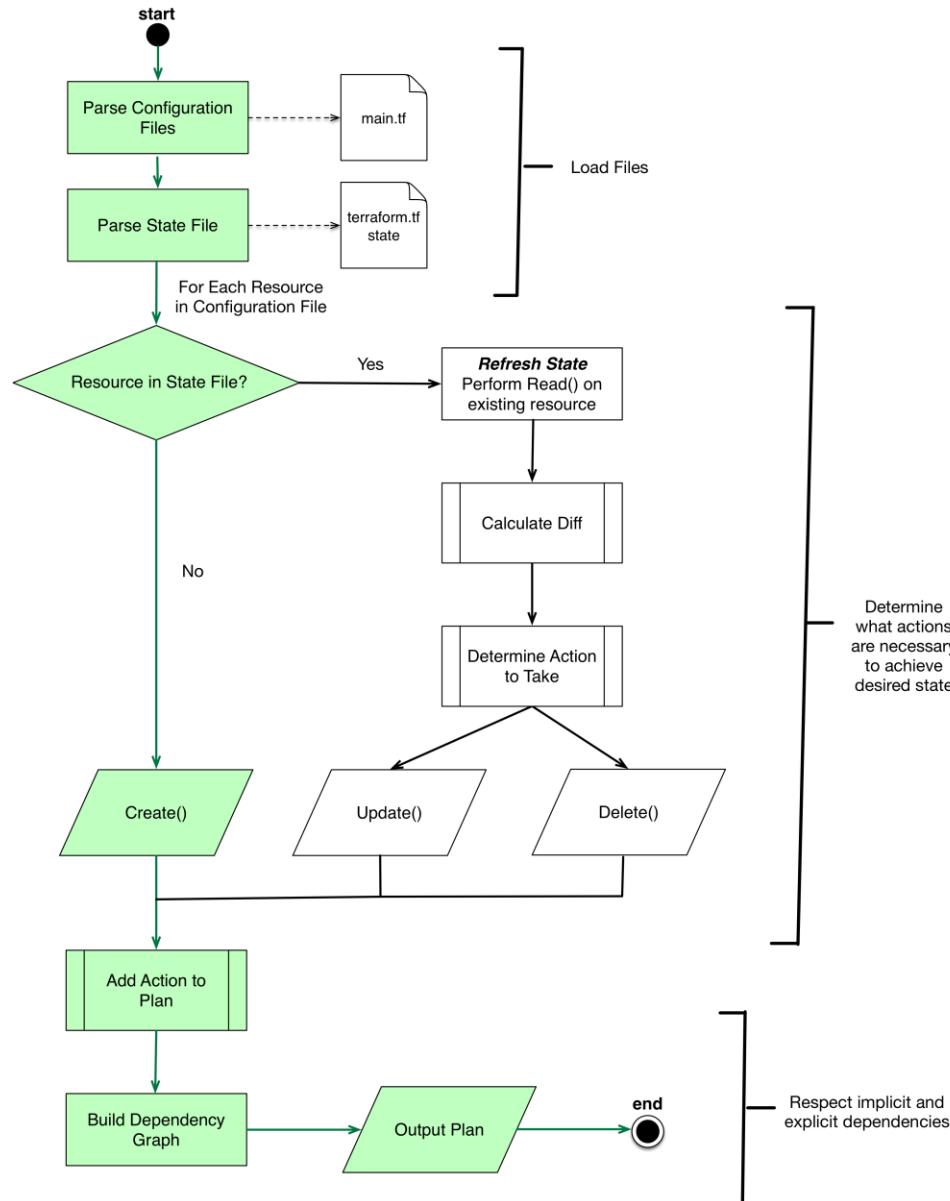
Terraform architecture



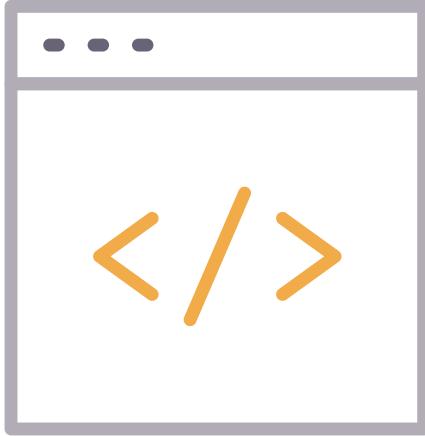
Terraform architecture



Terraform architecture



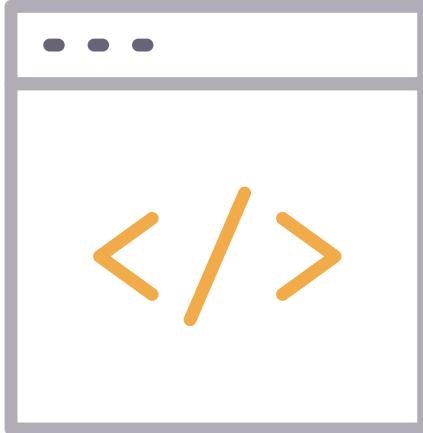
Terraform files



```
##Amazon Infrastructure
provider "aws" {
    region = "${var.aws_region}"
}
```

- Provider defines what infrastructure Terraform will be managing
 - Pass variables to provider
 - Region, Flavor, Credentials

Terraform files

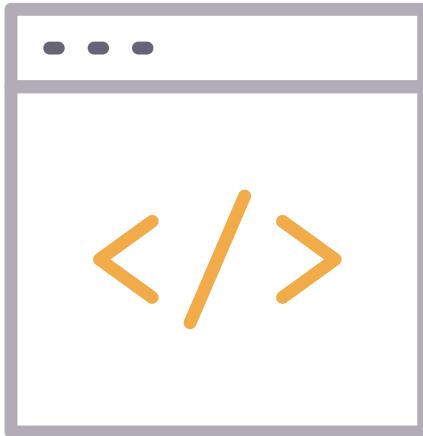


```
data "aws_ami" "ubuntu" {
    most_recent = true

    filter {
        name    = "name"
        values  = ["ubuntu/images/hvm-ssd/ubuntu-xenial-16.04-amd64-server-*"]
    }
    filter {
        name    = "virtualization-type"
        values  = ["hvm"]
    }
    owners   = ["099720109477"] # Canonical
}
```

- Data sources
 - Queries AWS API for latest Ubuntu 16.04 image.
 - Stores results in a variable which is used later.

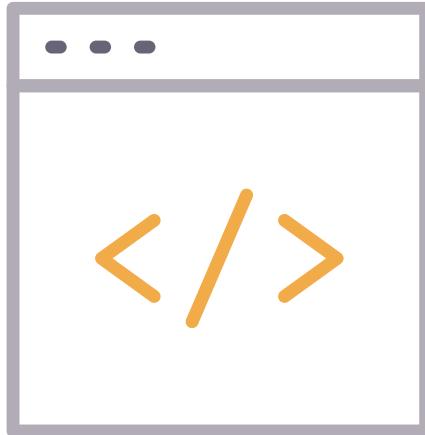
Terraform files



```
resource "aws_instance" "aws-k8s-master" {
    subnet_id          = "${var.aws_subnet_id}"
    depends_on         = [ "aws_security_group.k8s_sg" ]
    ami                = "${data.aws_ami.ubuntu.id}"
    instance_type      = "${var.aws_instance_size}"
    vpc_security_group_ids = ["${aws_security_group.k8s_sg.id}"]
    key_name           = "${var.aws_key_name}"
    count              = "${var.aws_master_count}"
    root_block_device {
        volume_type      = "gp2"
        volume_size       = 20
        delete_on_termination = true
    }
    tags {
        Name = "k8s-master-${count.index}"
        role = "k8s-master"
    }
}
```

- Resource: Defines specifications for creation of infrastructure.

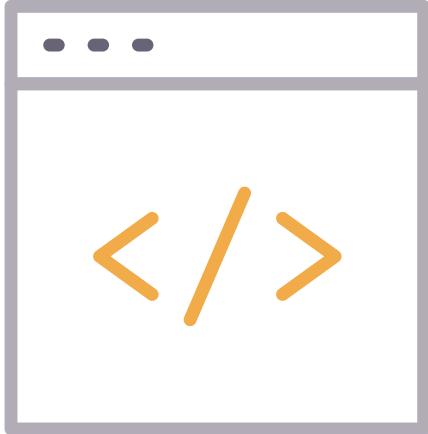
Terraform files



```
resource "aws_instance" "aws-k8s-master" {
  subnet_id          = "${var.aws_subnet_id}"
  depends_on         = ["aws_security_group.k8s_sg"]
  ami                = "${data.aws_ami.ubuntu.id}"
  instance_type      = "${var.aws_instance_size}"
  vpc_security_group_ids = ["${aws_security_group.k8s_sg.id}"]
  key_name           = "${var.aws_key_name}"
  count              = "${var.aws_master_count}"
  root_block_device {
    volume_type      = "gp2"
    volume_size       = 20
    delete_on_termination = true
  }
  tags {
    Name = "k8s-master-${count.index}"
    role = "k8s-master"
  }
}
```

- ami is set by results of previous data source query

Terraform files



```
##AWS Specific Vars
variable "aws_master_count" {
| default = 10
}

variable "aws_worker_count" {
| default = 20
}

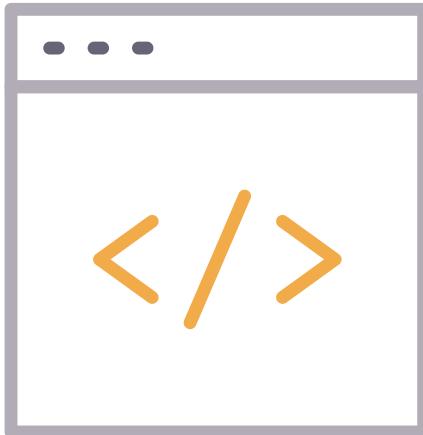
variable "aws_key_name" {
| default = "k8s"
}

variable "aws_instance_size" {
| default = "t2.small"
}

variable "aws_region" {
| default = "us-west-1"
}
```

- Define sane defaults in variables.tf

Terraform files



```
resource "aws_instance" "aws-k8s-master" {
    subnet_id          = "${var.aws_subnet_id}"
    depends_on         = [ "aws_security_group.k8s_sg" ]
    ami                = "${data.aws_ami.ubuntu.id}"
    instance_type      = "${var.aws_instance_size}"
    vpc_security_group_ids = ["${aws_security_group.k8s_sg.id}"]
    key_name           = "${var.aws_key_name}"
    count              = "${var.aws_master_count}"

    root_block_device {
        volume_type      = "gp2"
        volume_size       = 20
        delete_on_termination = true
    }

    tags {
        Name = "k8s-master-${count.index}"
        role = "k8s-master"
    }
}
```

- Value is 'count' is setup by variable in variables.tf

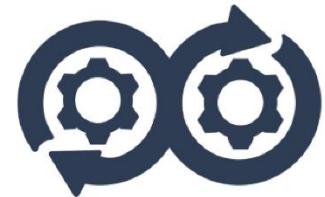
Demo: Terraform



Lab: Terraform Introduction



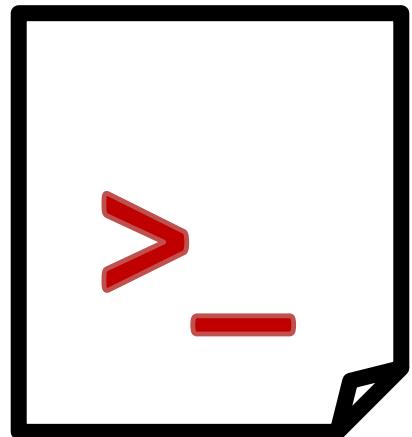
Ansible



Scenario Progression



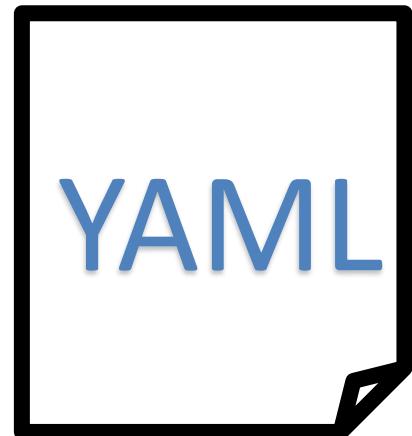
1. Manual
Commands



2. Reusable
Scripts

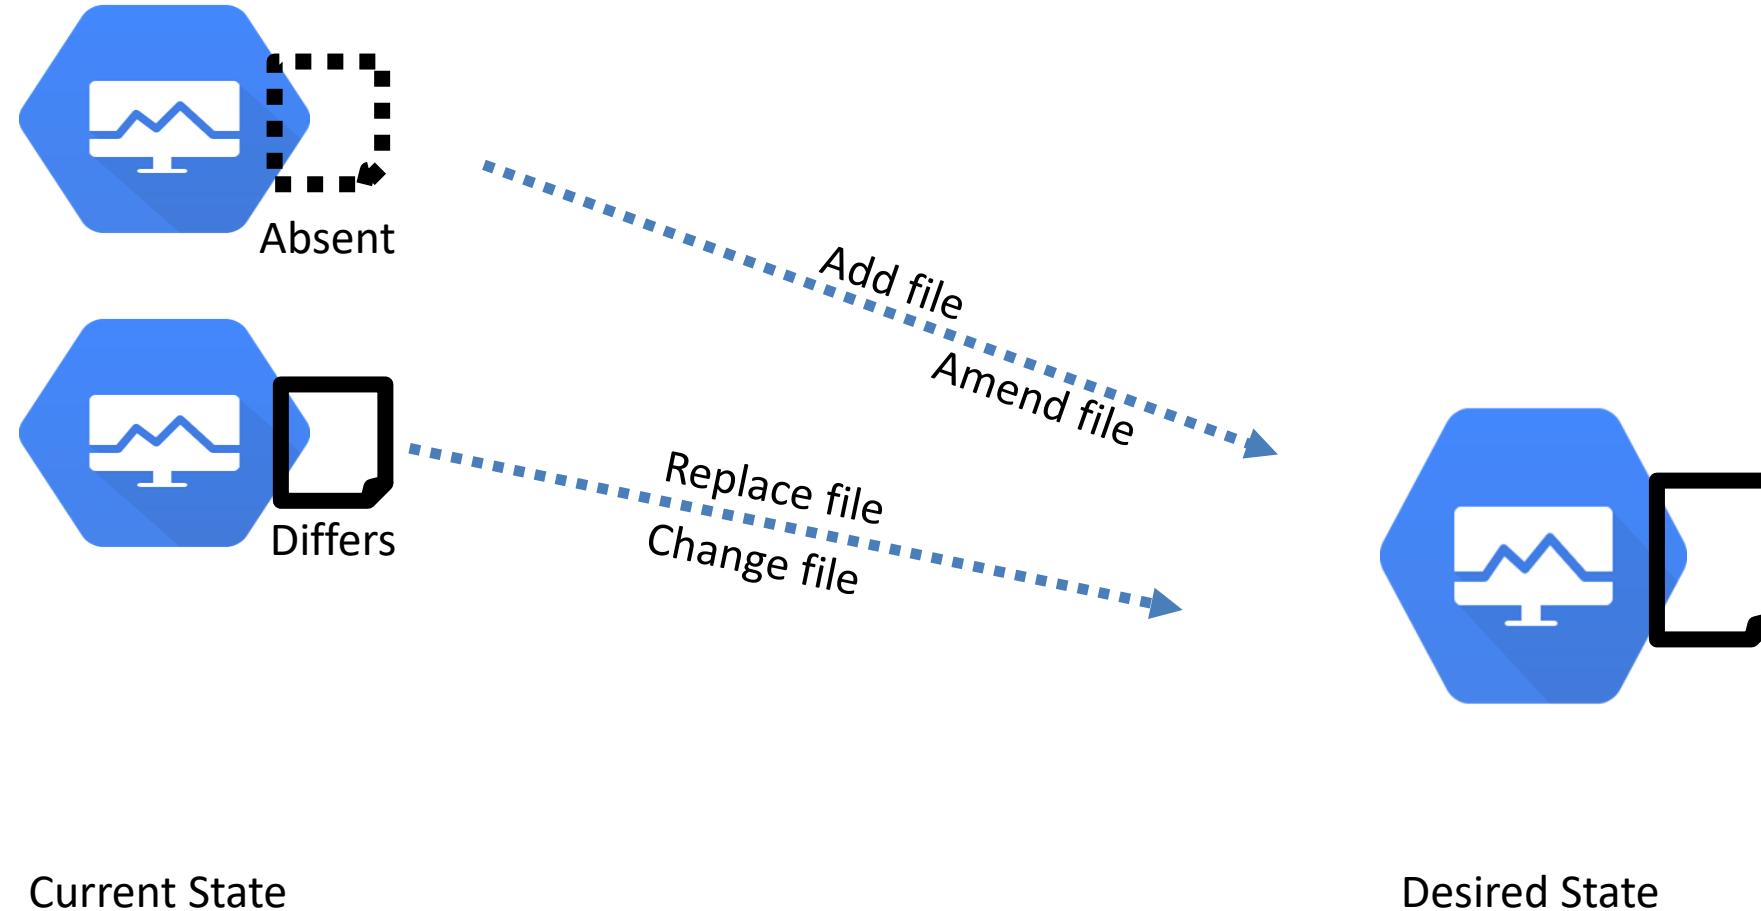


3. Ansible
Ad-hoc

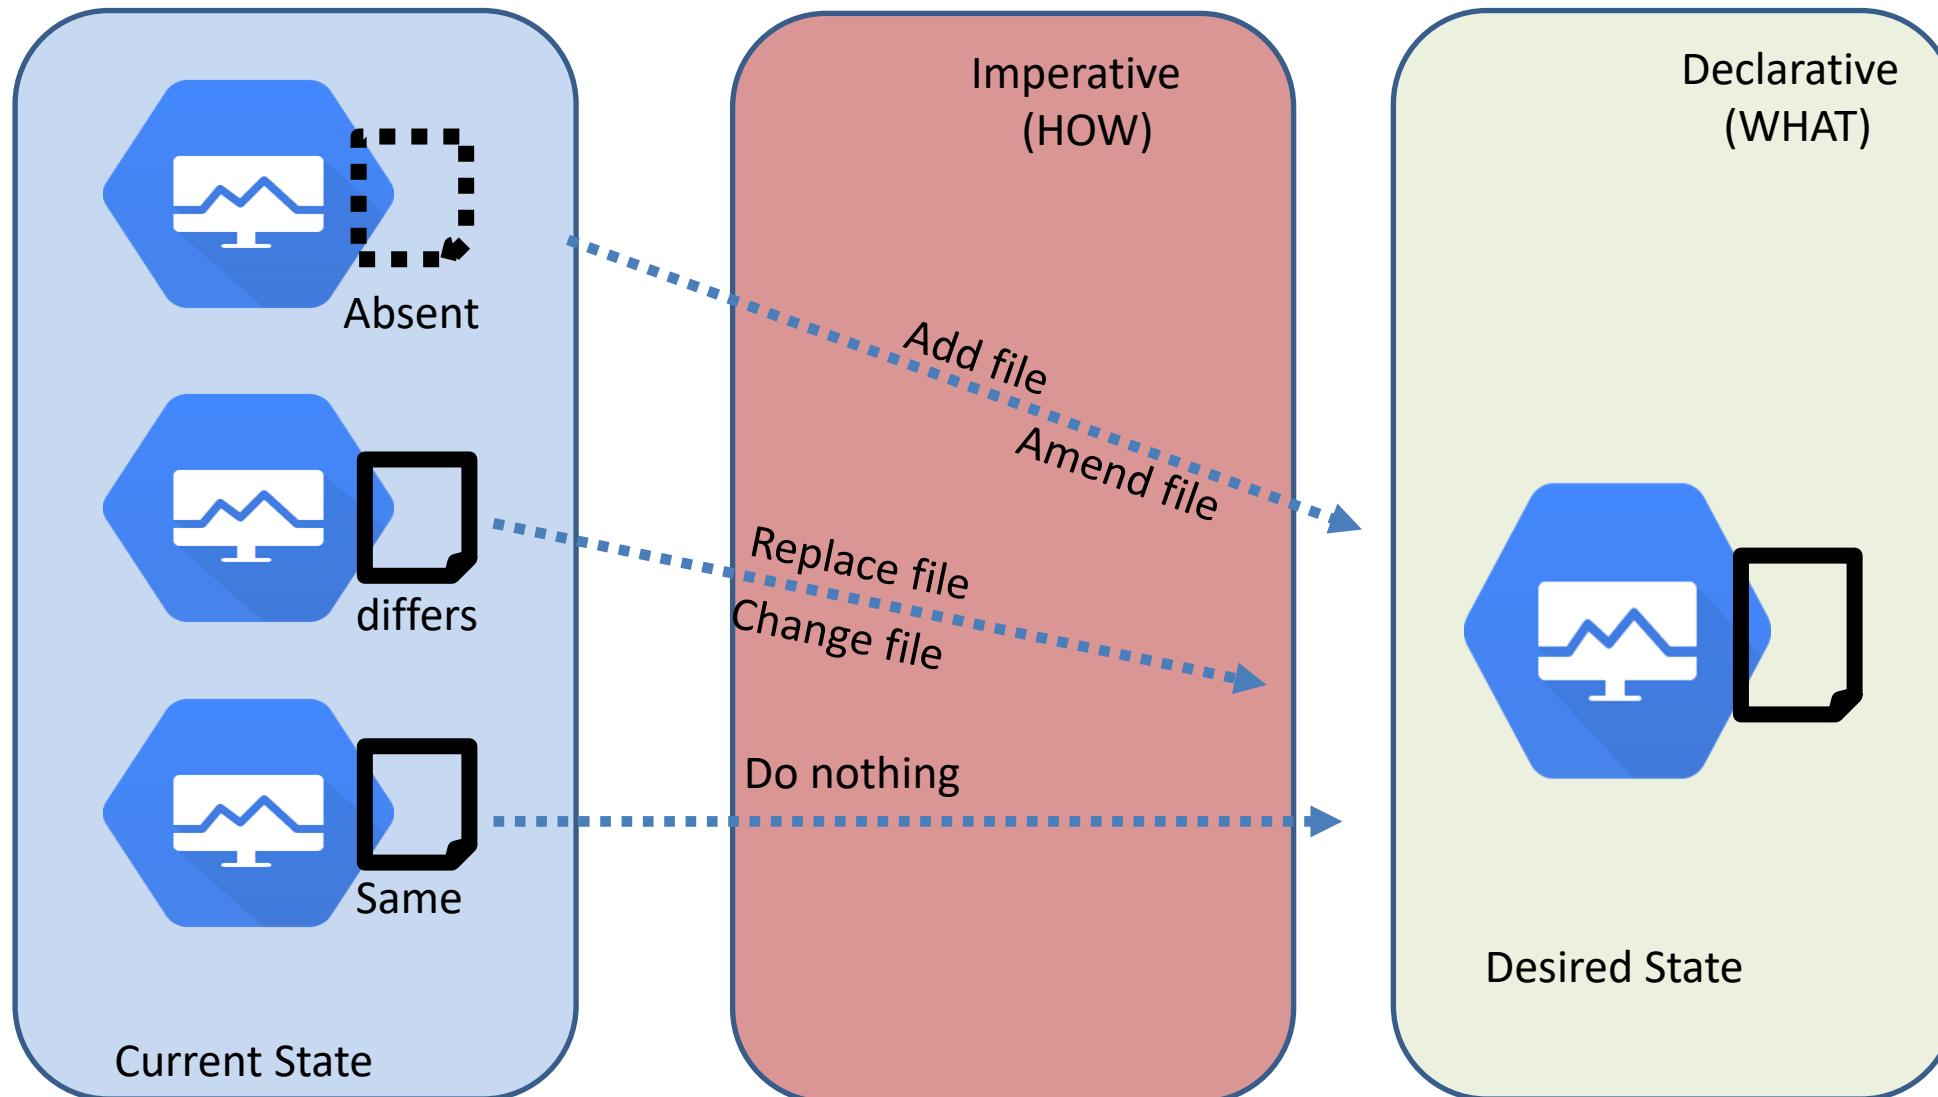


4. Playbooks

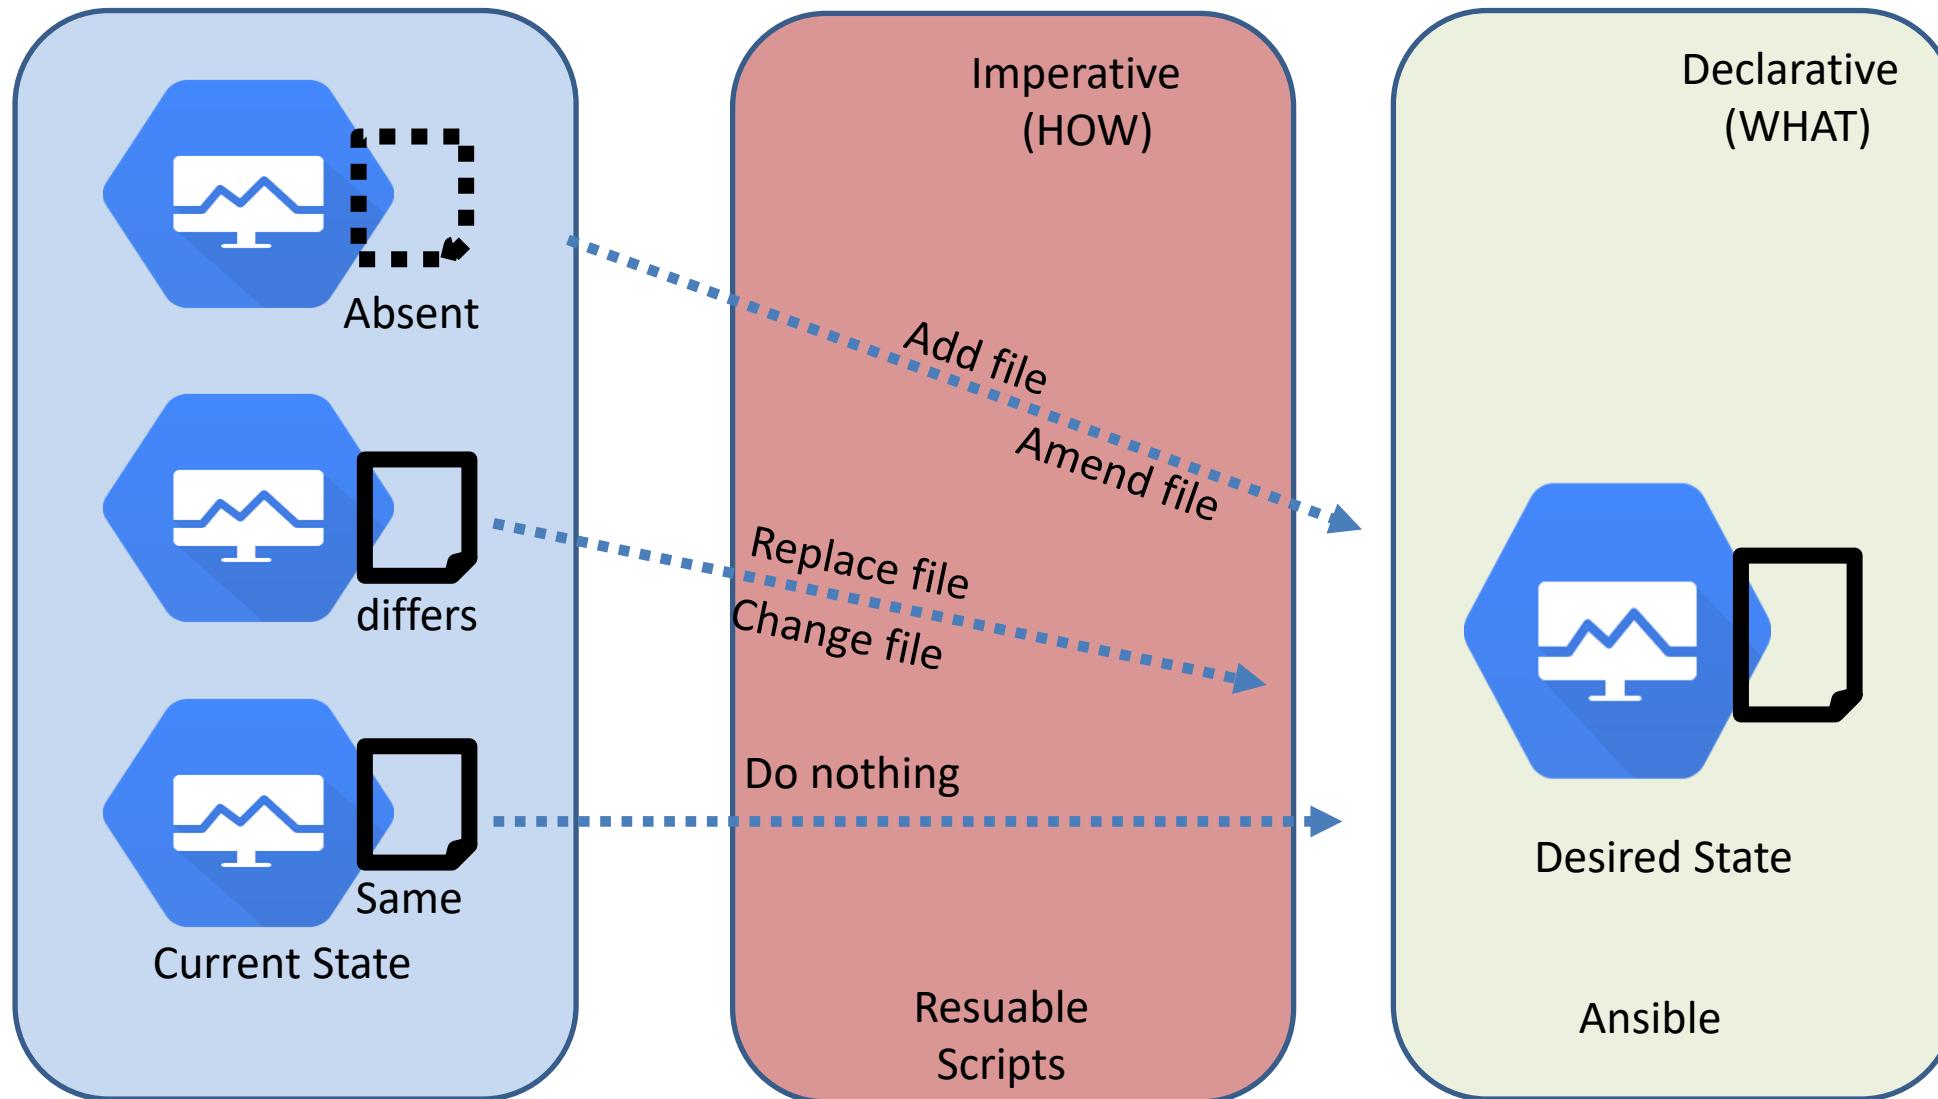
Reconciliation



Reconciliation



Reconciliation

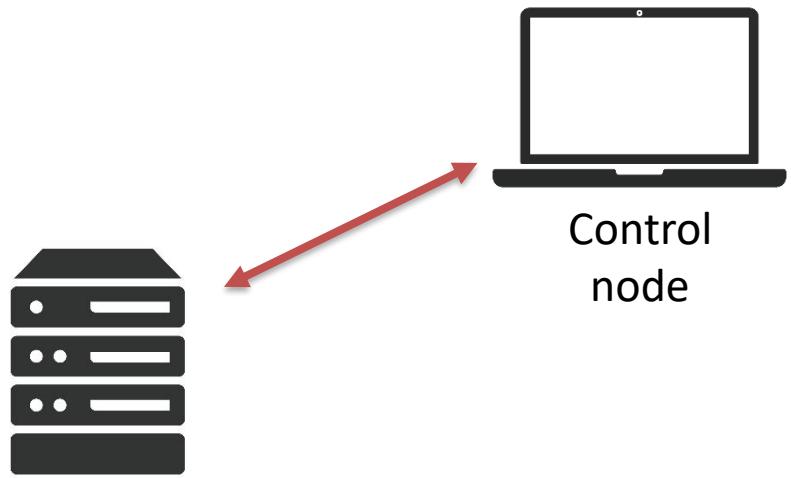


Architecture

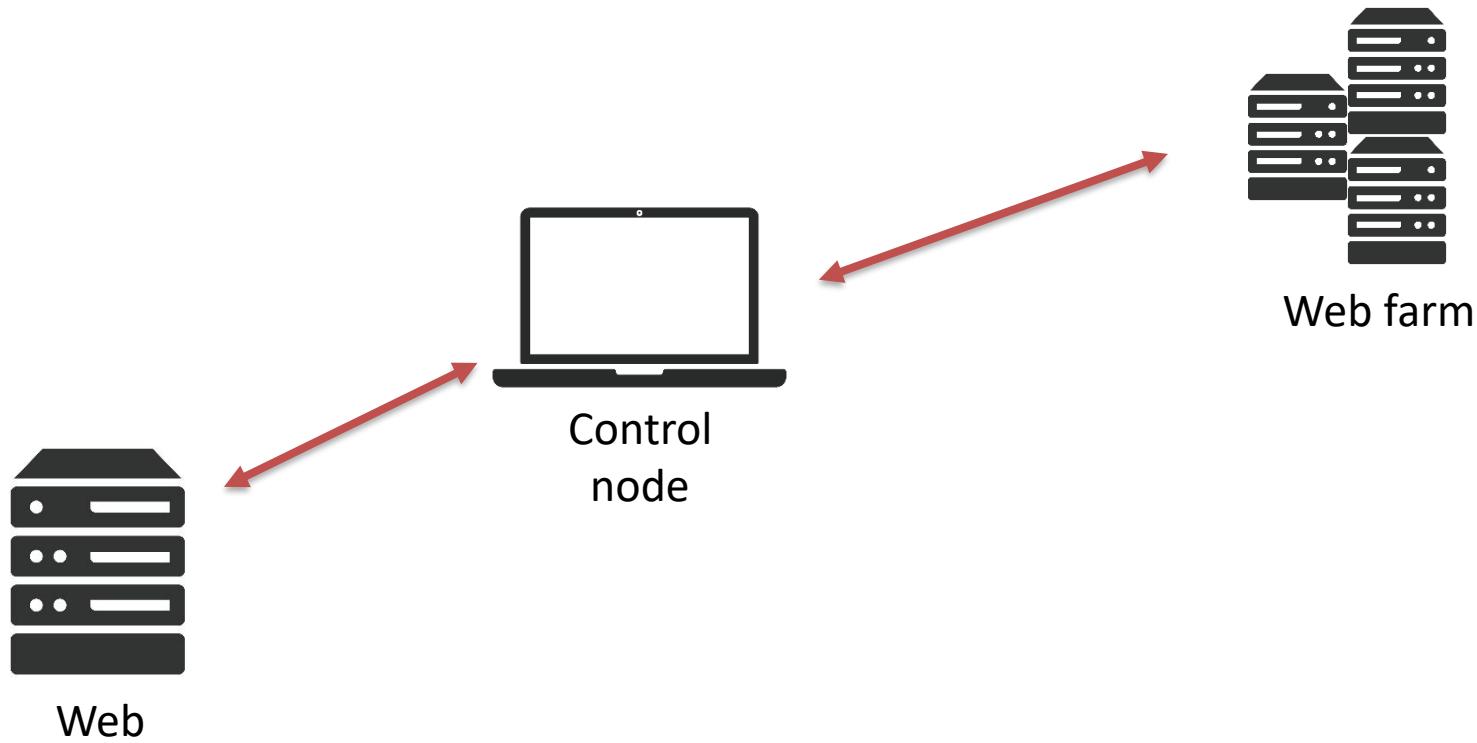


Control
node

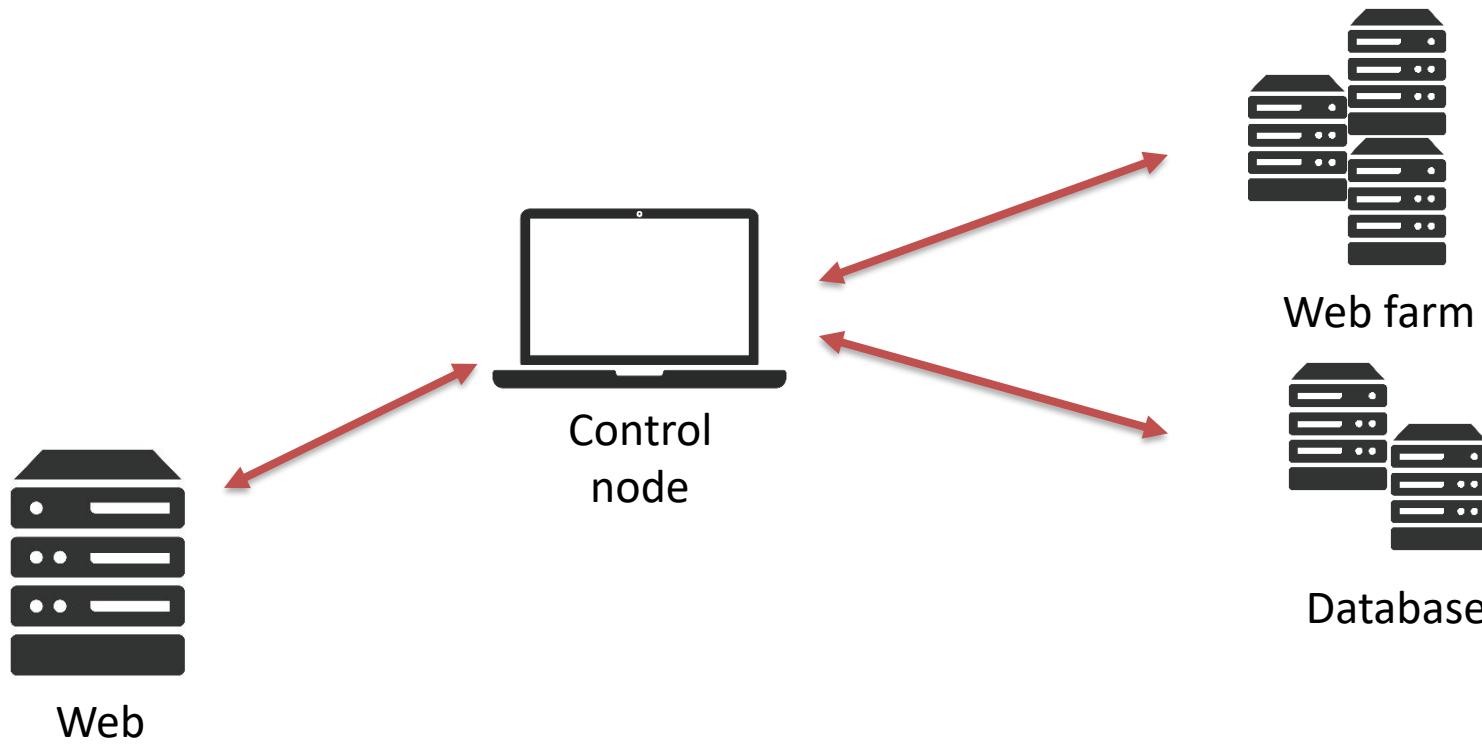
Architecture



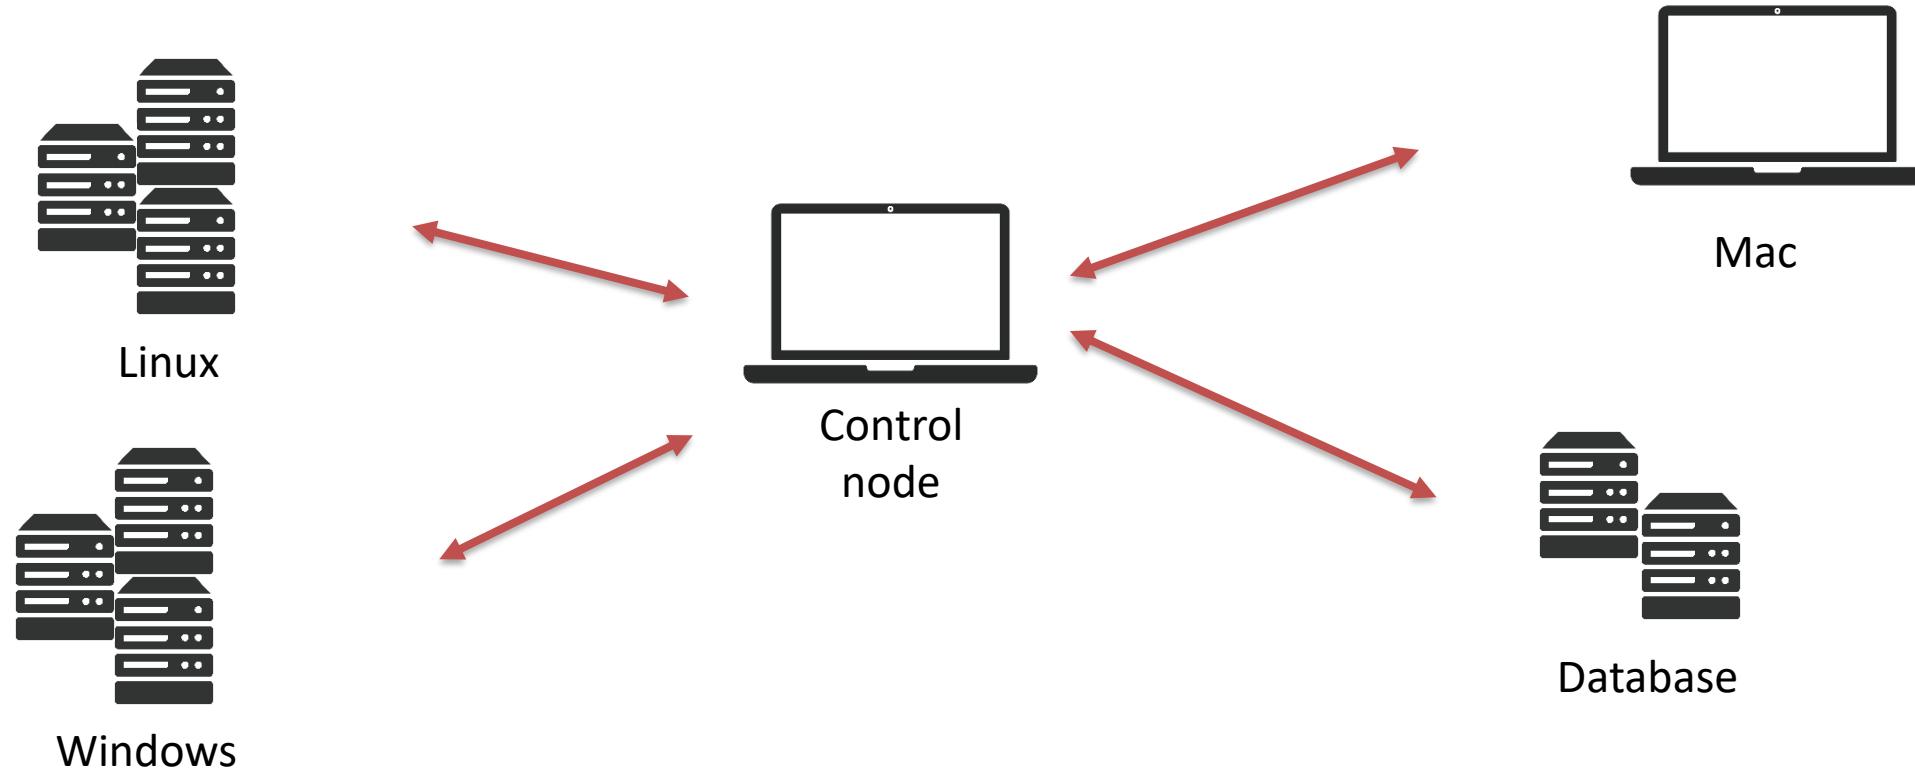
Architecture



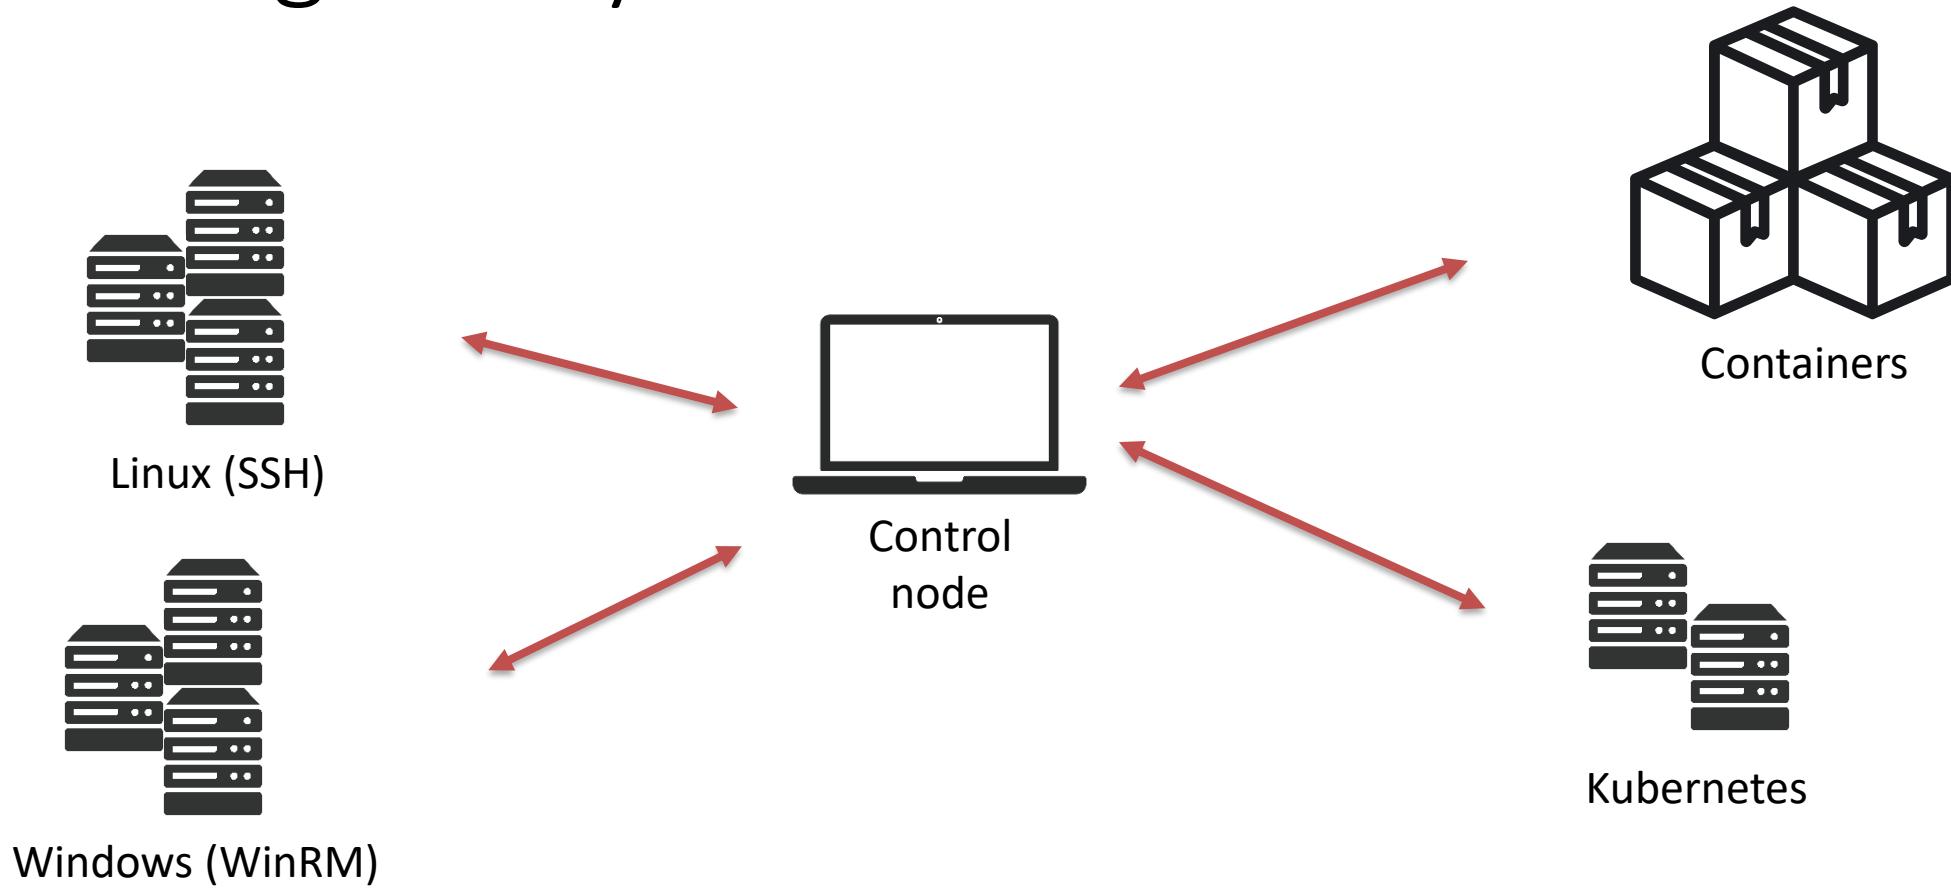
Architecture



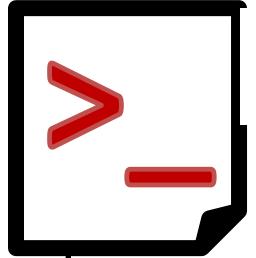
Connecting to anywhere



Connecting to anywhere



Ansible Ad-hoc



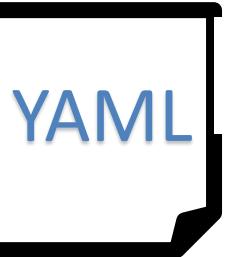
Scripted Ad-hoc

```
ansible -m copy -a "src=master.gitconfig dest=~/gitconfig" localhost
```

```
ansible -m homebrew -a "name=bat state=latest" localhost
```

```
ansible -i hosts all -m ping -u ubuntu
```

Playbook



playbook.yml

```
hosts: localhost
tasks :
  - copy :
      src: "master.gitconfig"
      dest: "~/.gitconfig"
```

Playbook

YAML

playbook.yml

```
hosts: localhost
tasks :
  - copy :
      src: "master.gitconfig"
      dest: "~/.gitconfig"
```

V

```
ansible-playbook playbook.yml
```

More complex playbook

```
- name: Install Docker
  hosts: tag_role_k8s_master:tag_role_k8s_member
  remote_user: ubuntu
  gather_facts: no
  handlers:
    - include: roles/handlers/main.yml
  tasks:
    - package:
        name: "docker-ce"
        state: latest
    - user:
        name: "{{ ansible_ssh_user }}"
        groups: docker
        append: yes
        notify: restart docker
```

Name of Play

More complex playbook

```
- name: Install Docker
  hosts: tag_role_k8s_master:tag_role_k8s_member
  remote_user: ubuntu
  gather_facts: no
  handlers:
    - include: roles/handlers/main.yml
  tasks:
    - package:
        name: "docker-ce"
        state: latest
    - user:
        name: "{{ ansible_ssh_user }}"
        groups: docker
        append: yes
        notify: restart docker
```

Hosts to run Play on

More complex playbook

```
- name: Install Docker
  hosts: tag_role_k8s_master:tag_role_k8s_member
  remote_user: ubuntu
  gather_facts: no
  handlers:
    - include: roles/handlers/main.yml
  tasks:
    - package:
        name: "docker-ce"
        state: latest
    - user:
        name: "{{ ansible_ssh_user }}"
        groups: docker
        append: yes
        notify: restart docker
```

User to run Play as

More complex playbook

```
- name: Install Docker
  hosts: tag_role_k8s_master:tag_role_k8s_member
  remote_user: ubuntu
  gather_facts: no
  handlers:
    - include: roles/handlers/main.yml
  tasks:
    - package:
        name: "docker-ce"
        state: latest
    - user:
        name: "{{ ansible_ssh_user }}"
        groups: docker
        append: yes
        notify: restart docker
```

Define what to do after package is installed.

More complex playbook

```
- name: Install Docker
  hosts: tag_role_k8s_master:tag_role_k8s_member
  remote_user: ubuntu
  gather_facts: no
  handlers:
    - include: roles/handlers/main.yml
  tasks:
    - package:
        name: "docker-ce"
        state: latest
    - user:
        name: "{{ ansible_ssh_user }}"
        groups: docker
        append: yes
        notify: restart docker
          Call the defined handler
```

More complex playbook

```
- name: Install Docker
  hosts: tag_role_k8s_master:tag_role_k8s_member
  remote_user: ubuntu
  gather_facts: no
  handlers:
    - include: roles/handlers/main.yml
  tasks:
    - package:
        name: "docker-ce"
        state: latest
    - user:
        name: "{{ ansible_ssh_user }}"
        groups: docker
        append: yes
        notify: restart docker
```

Tasks to run
- install Docker

More complex playbook

```
- name: Install Docker
  hosts: tag_role_k8s_master:tag_role_k8s_member
  remote_user: ubuntu
  gather_facts: no
  handlers:
    - include: roles/handlers/main.yml
  tasks:
    - package:
        name: "docker-ce"
        state: latest
    - user:
        name: "{{ ansible_ssh_user }}"
        groups: docker
        append: yes
    notify: restart docker
  Create user account
```

Lab: Ansible Introduction

